



BITS Pilani
Pilani Campus

BITS Pilani presentation

Dr. Vivek V. Jog
Dept. of Computer Engineering





BITS Pilani
Pilani Campus



Big Data Systems (S1-24_CCZG522)

Lecture No.9



- Entry of Apache Pig
- Pig vs MapReduce
- Twitter Case Study on Apache Pig
- Apache Pig Architecture
- Pig Components
- Pig Data Model & Operators
- Running Pig Commands and Pig Scripts (Log Analysis)



No Need to be a Programmer



In MapReduce, you need to write a program in Java/Python to process the data.

Developed by YAHOO



- An open-source **high-level dataflow system**
- Introduced by **Yahoo**
- Provides abstraction over MapReduce
- Two main components – the *Pig Latin* language and the *Pig Execution*

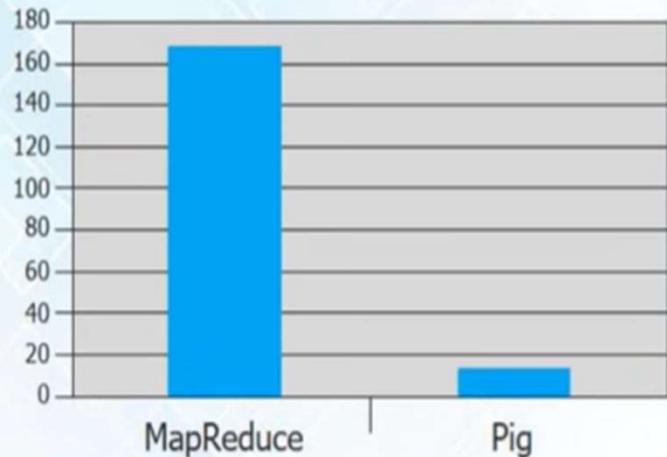
Fun Fact:

✓ *10 lines of pig latin= approx. 200 lines of Map-Reduce Java Program*

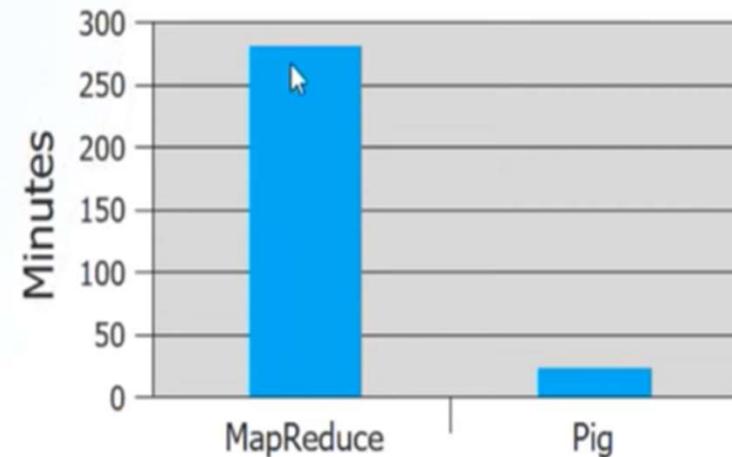


Clear Winner

1/20 the lines of Code



1/16 the development Time



Pig v/s MapReduce



- High-level data flow tool
- No need to write complex programs
- Built-in support for data operations like joins, filters, ordering, sorting etc.
- Provides nested data types like tuples, bags, and maps



- Low-level data processing paradigm
- You need write programs in Java/Python etc.
- Performing data operations in MapReduce is a humongous task
- Nested data types are not there in MapReduce

Pig Specific

- Can take any data
 - Structured data
 - Semi-Structured data
 - Unstructured data
 - Easy to learn, Easy to write and Easy to read
 - Data Flow Language
 - Reads like a series of steps
 - Extensible by UDF (User Defined Functions)
 - Java
 - Python
 - JavaScript
 - Ruby
- Provides common data operations **filters**, **joins**, **ordering**, etc. and nested data types **tuples**, **bags**, and **maps** missing from MapReduce.
- An **ad-hoc** way of creating and executing map-reduce jobs on very large data sets
- Open source and actively supported by a community of developers.

Pig – Hello World

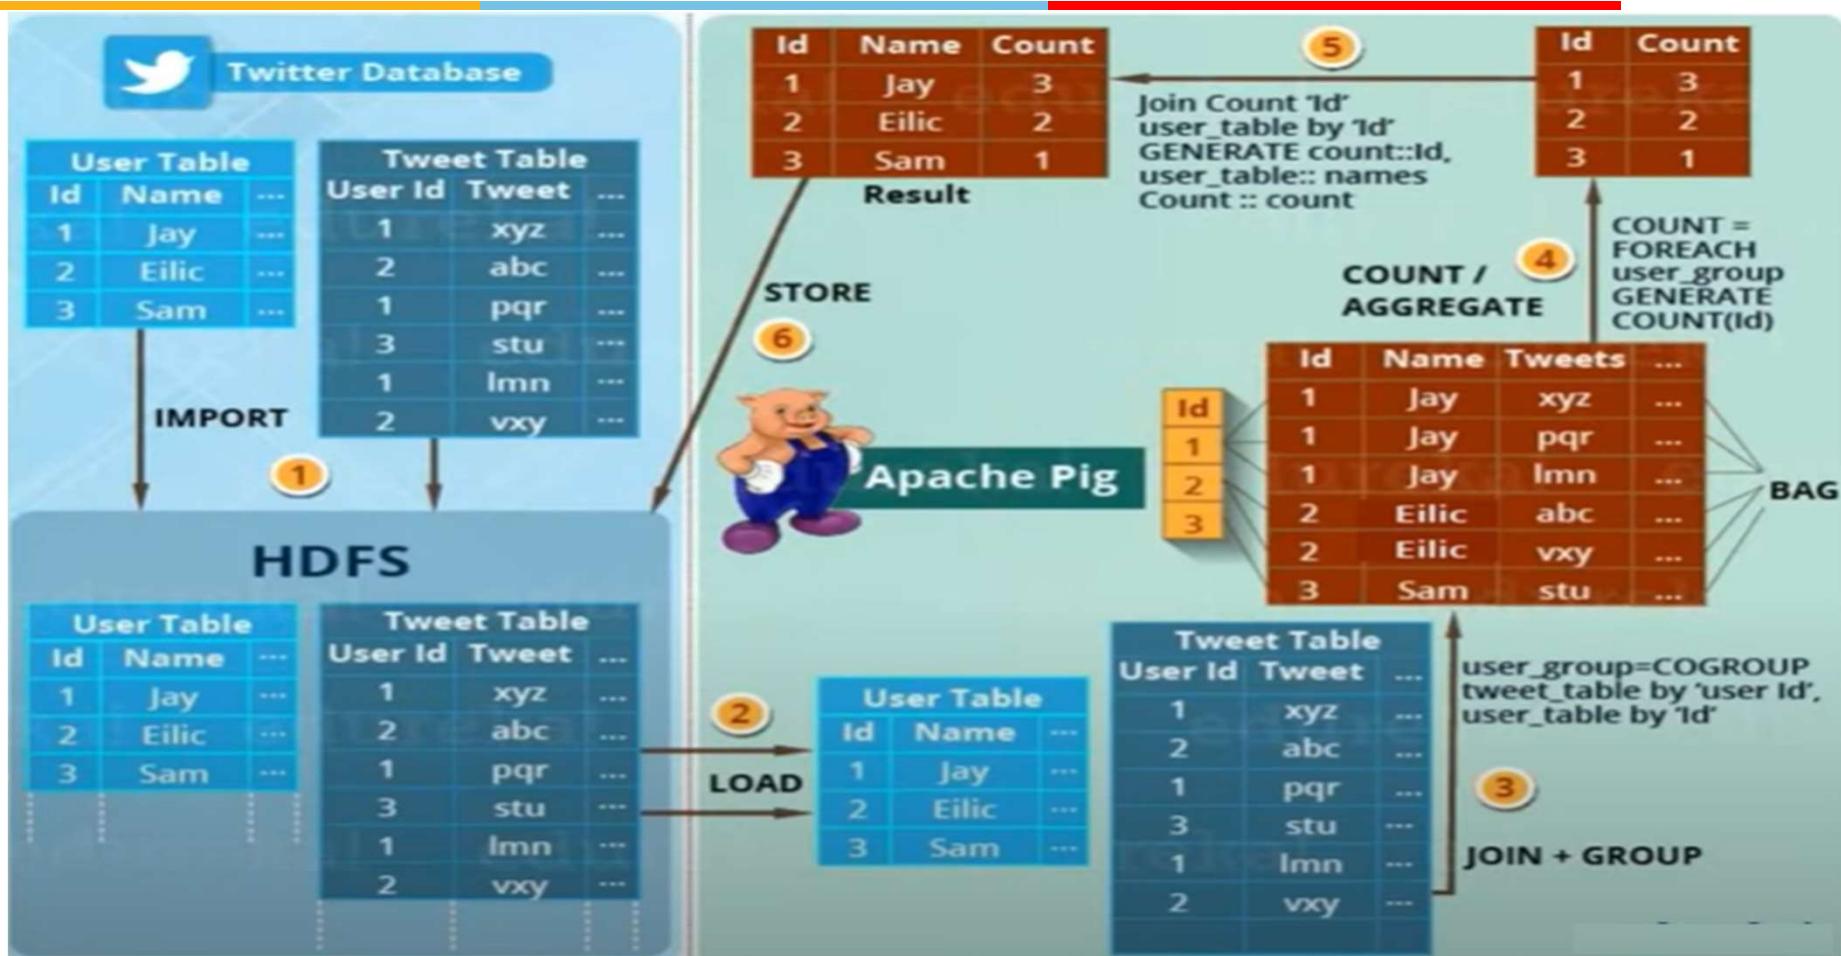


- Twitter's data was growing at an accelerating rate (i.e. 10 TB/day).
- Thus, Twitter decided to move the archived data to HDFS and adopt Hadoop for extracting the business values out of it.
- Their major aim was to analyse data stored in Hadoop to come up with the multiple insights on a daily, weekly or monthly basis.

Let me talk about one of the insight they wanted to know.

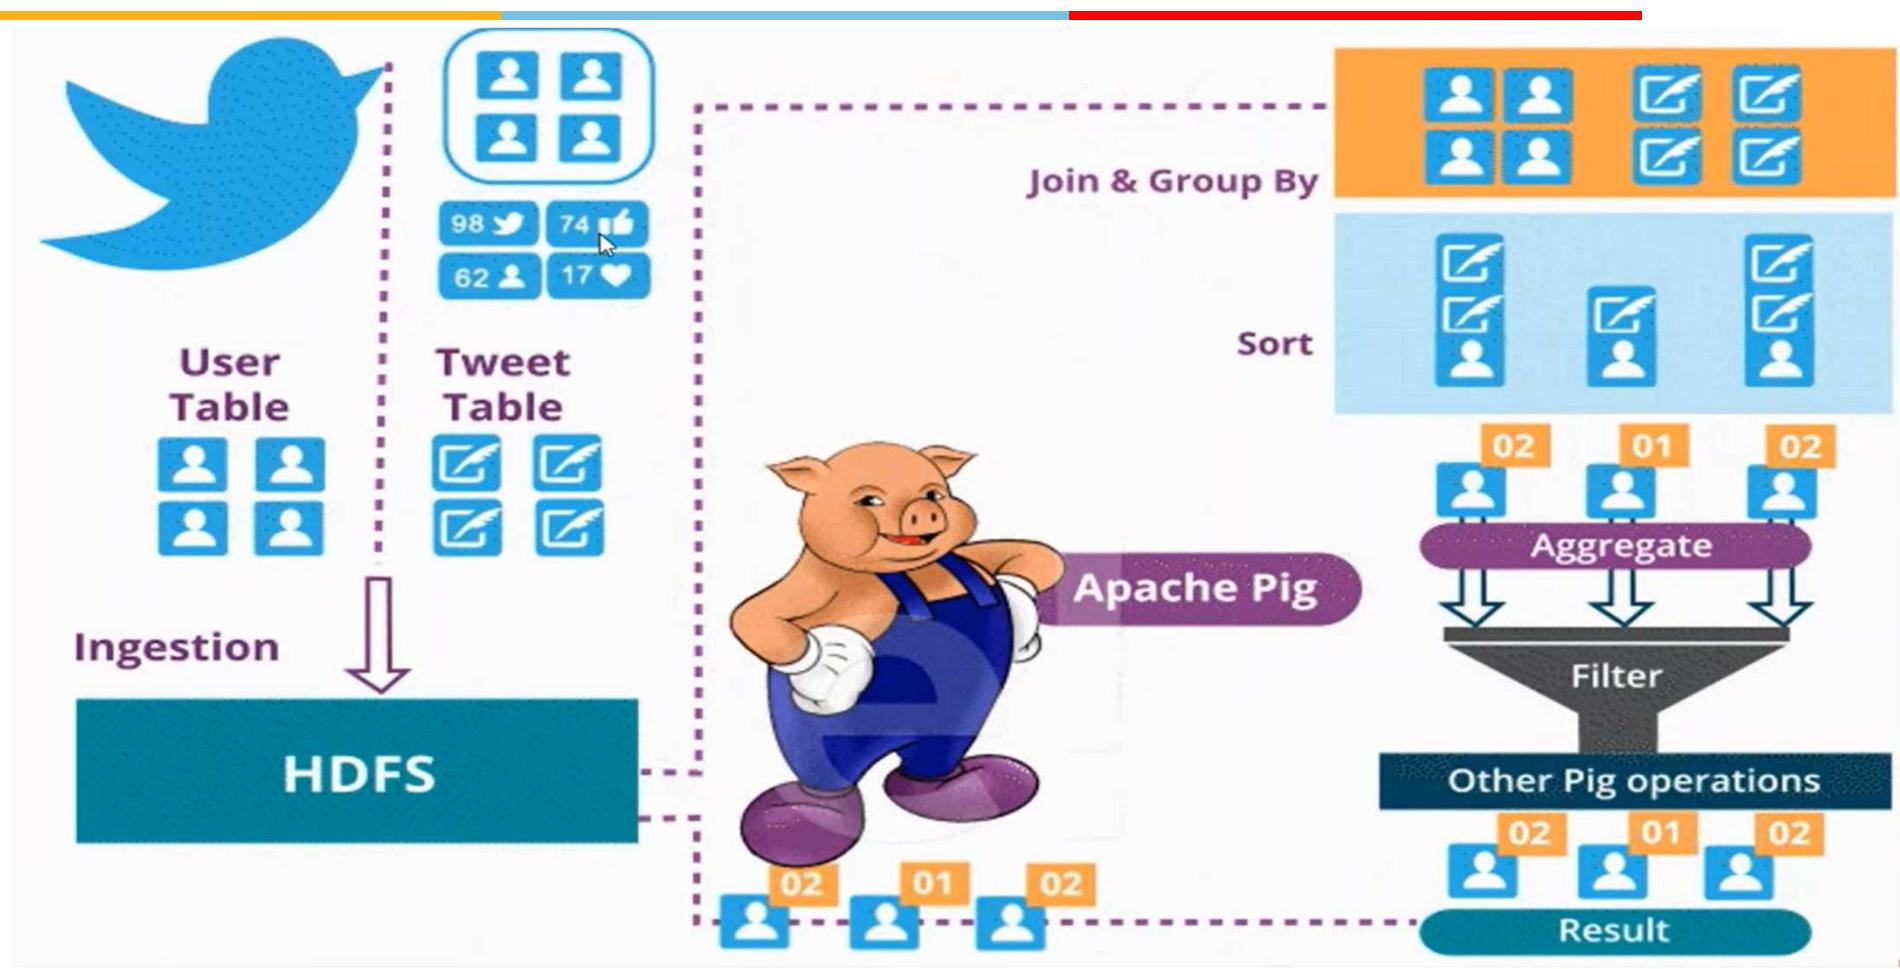
Analyzing how many tweets are stored per user, in the given tweet tables?

Working

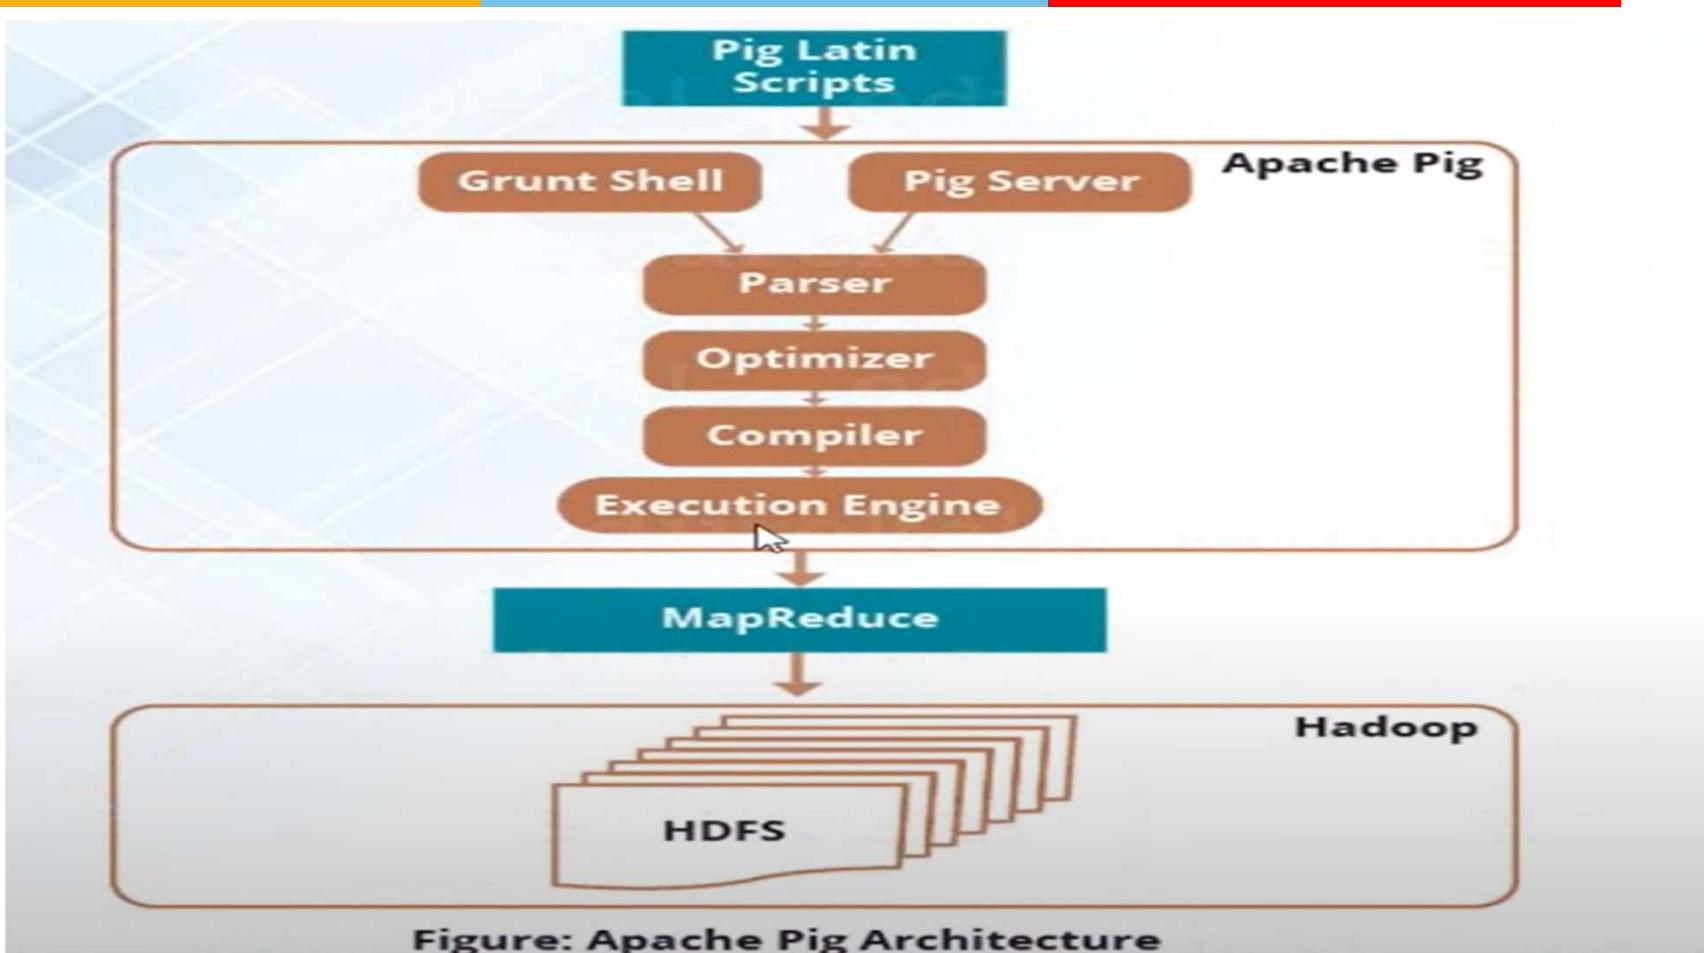




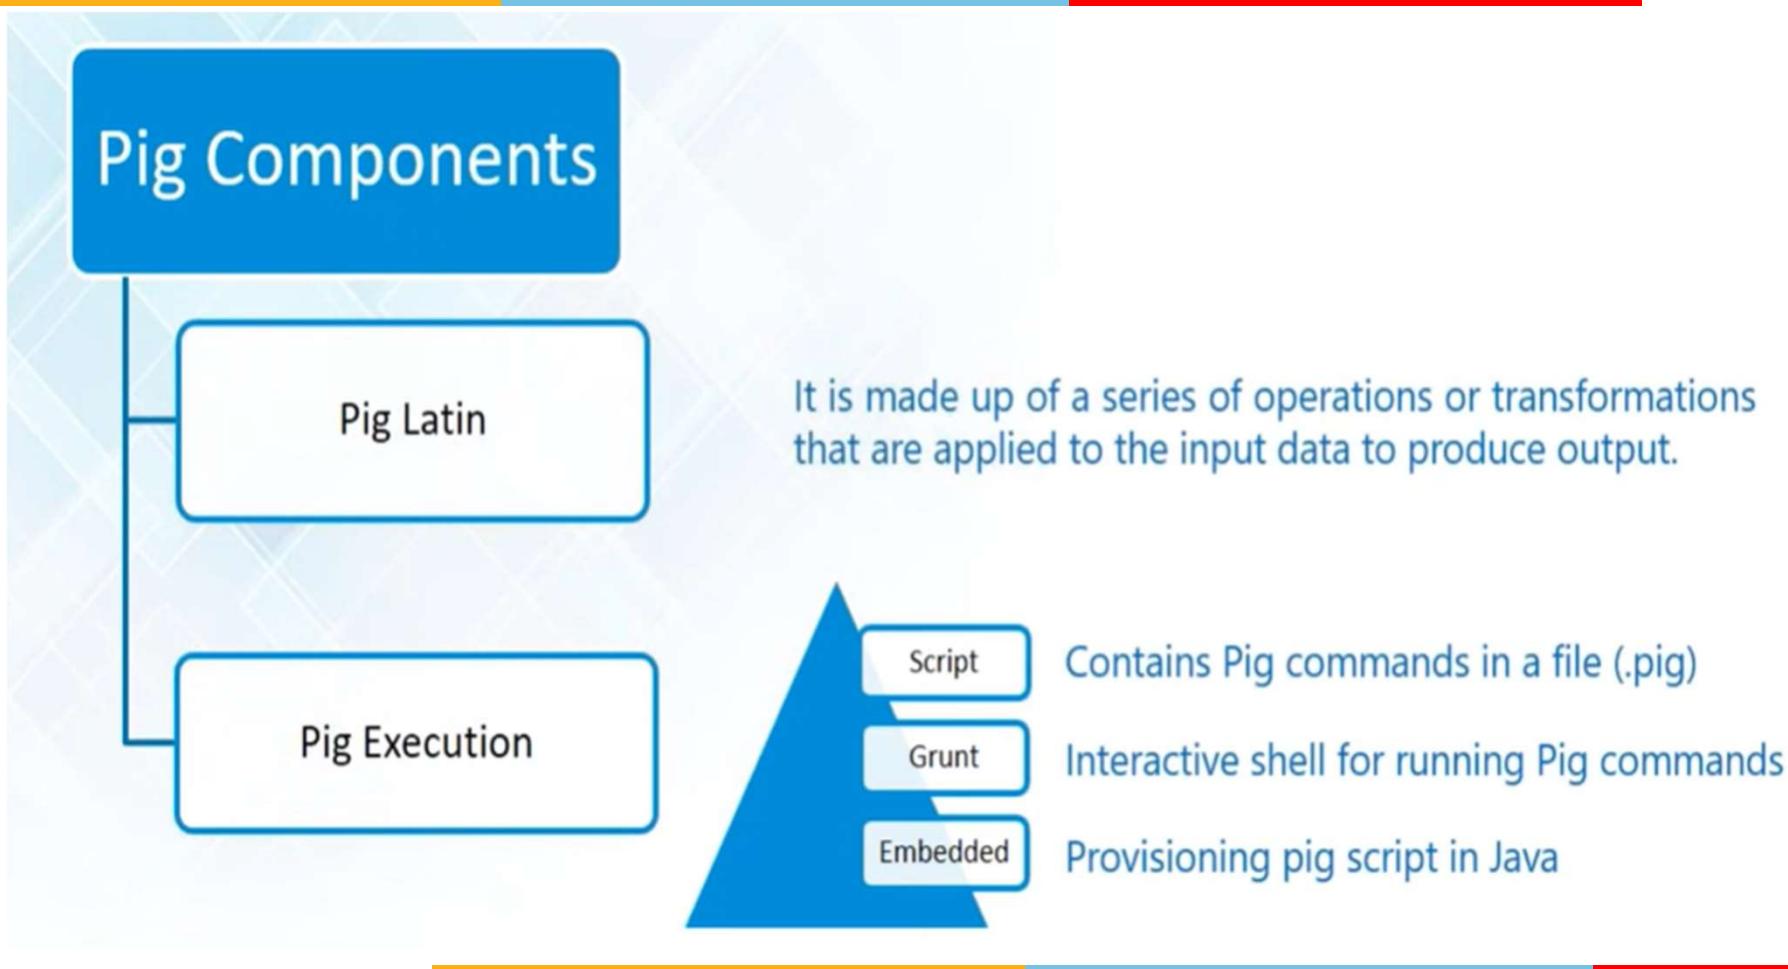
Overview



Architecture



Components



Modes of Execution

You can run
Apache Pig
in 2 modes:

MapReduce Mode – This is the default mode, which requires access to a Hadoop cluster and HDFS installation. The input and output in this mode are present on HDFS.

Command: pig

Local Mode – With access to a single machine, all files are installed and run using a local host and file system. Here the local mode is specified using '-x flag' (pig -x local). The input and output in this mode are present on local file system.

Command: pig -x local

MapReduce & Local Mode

```
[edureka@localhost ~]$ pig  
16/12/23 12:55:06 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL  
16/12/23 12:55:06 INFO pig.ExecTypeProvider: Trying ExecType : MAPREDUCE  
16/12/23 12:55:06 INFO pig.ExecTypeProvider: Picked MAPREDUCE as the ExecType  
2016-12-23 12:55:07,414 [main] INFO org.apache.pig.Main - Apache Pig version 0.16.0 (r1746530) compi  
2016, 23:10:49  
2016-12-23 12:55:07,416 [main] INFO org.apache.pig.Main - Logging error messages to: /home/edureka/  
07395.log  
2016-12-23 12:55:07,776 [main] INFO org.apache.pig.impl.util.Utils - Default bootup file /home/edu  
up not found
```

```
[edureka@localhost ~]$ pig -x local  
SLF4J: Class path contains multiple SLF4J bindings.  
SLF4J: Found binding in [jar:file:/usr/lib/hadoop-2.7.3/share/hadoop/common/lib/slf4j-log4j12-1.7  
f4j-impl/StaticLoggerBinder.class]  
SLF4J: Found binding in [jar:file:/usr/lib/hbase-0.96.2-hadoop2/lib/slf4j-log4j12-1.6.4.jar!/org/  
icLoggerBinder.class]  
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.  
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]  
16/12/23 12:56:22 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL  
16/12/23 12:56:22 INFO pig.ExecTypeProvider: Picked LOCAL as the ExecType  
2016-12-23 12:56:23,163 [main] INFO org.apache.pig.Main - Apache Pig version 0.16.0 (r1746530) c  
2016, 23:10:49  
2016-12-23 12:56:23,165 [main] INFO org.apache.pig.Main - Logging error messages to: /home/edure  
83136.log  
2016-12-23 12:56:23,445 [main] INFO org.apache.pig.impl.util.Utils - Default bootup file /home/e  
up not found  
2016-12-23 12:56:27,264 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.jo  
precated. Instead, use mapreduce.jobtracker.address  
2016-12-23 12:56:27,267 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.defaul
```

Data Models

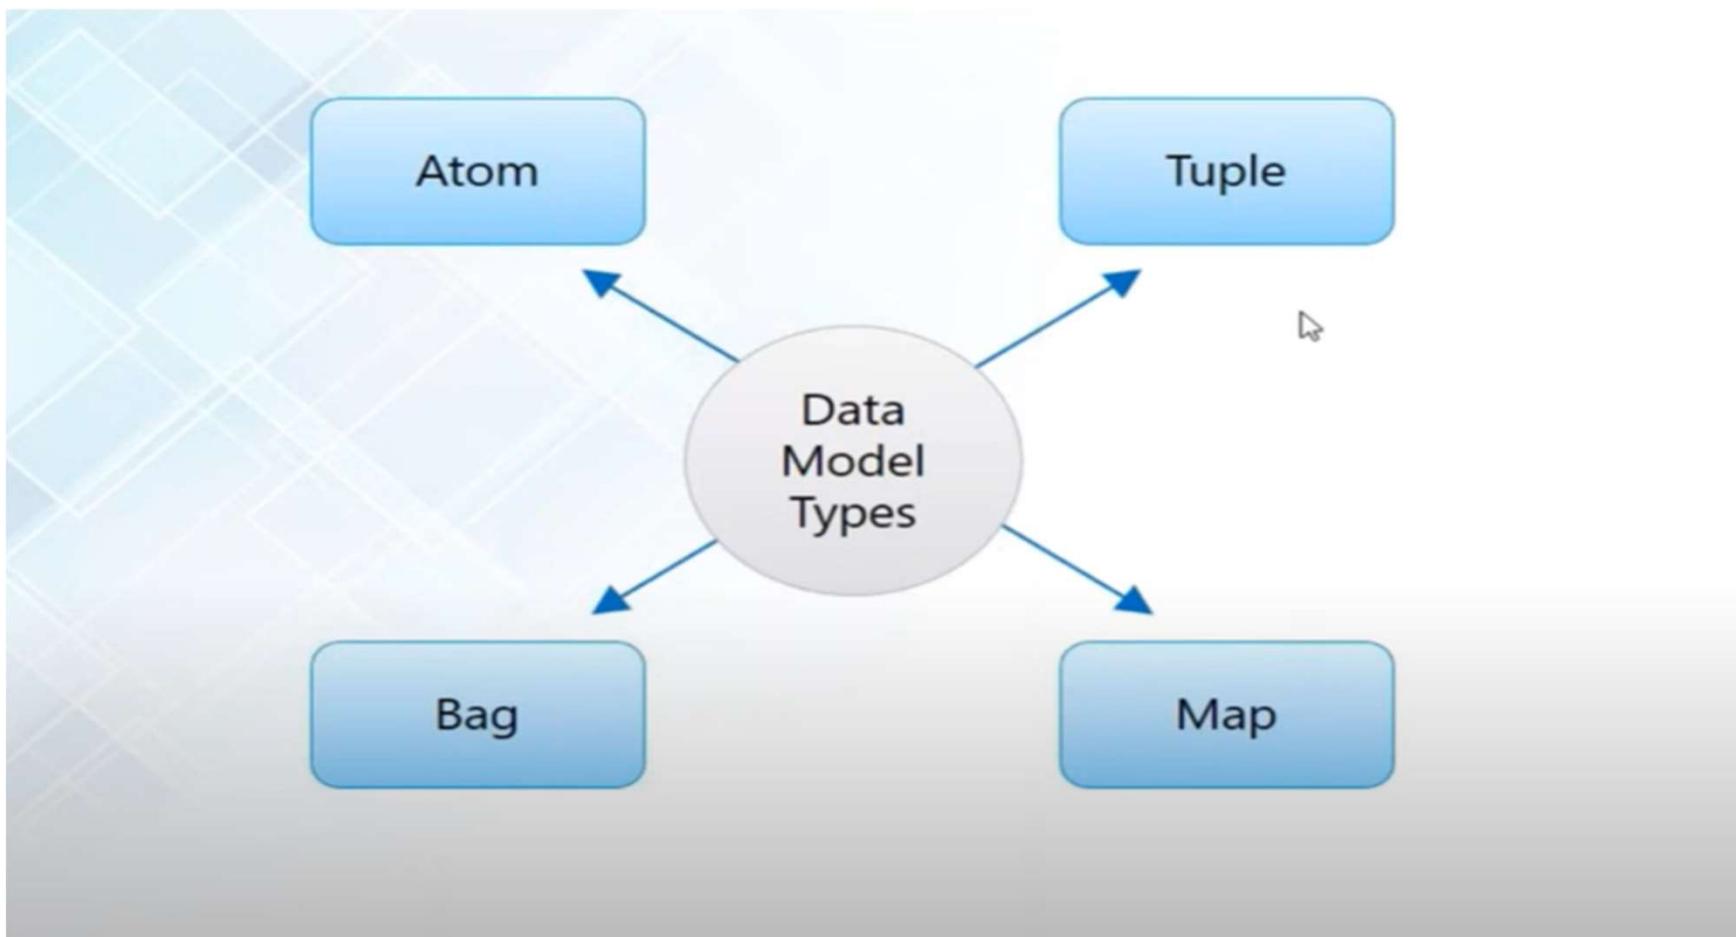




Figure: Apache Pig Data Model

edureka!

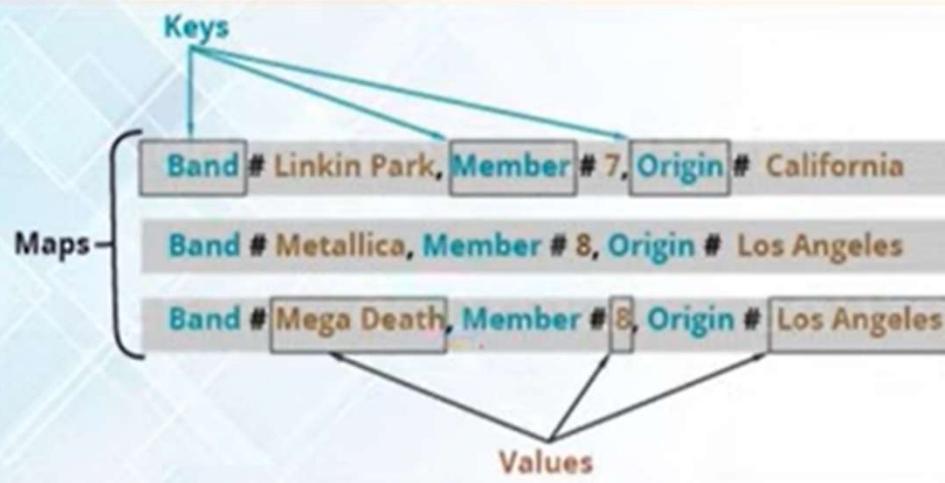
- **Tuple** is an ordered set of fields which may contain different data types for each field.

Example of tuple – (1, Linkin Park, 7, California)

- A **Bag** is a collection of a set of tuples and these tuples are subset of rows or entire rows of a table.

*Example of a bag – { (Linkin Park, 7, California),
(Metallica, 8), (Mega Death, Los Angeles) }*

Map & Atom



- A **Map** is key-value pairs used to represent data elements.

Example of maps- [band#Linkin Park, members#7], [band#Metallica, members#8]

- **Atoms** are basic data types which are used in all the languages like string, int, float, long, double, char[], byte[]

Operators for DB Operations

Operator	Description
LOAD	Load data from the local file system or HDFS storage into Pig
FOREACH	Generates data transformations based on columns of data
FILTER	Selects tuples from a relation based on a condition
JOIN	Join the relations based on the column
ORDER BY	Sort a relation based on one or more fields
STORE	Save results to the local file system or HDFS
DISTINCT	Removes duplicate tuples in a relation
GROUP	Groups together the tuples with the same group key (key field)
COGROUP	It is same as GROUP. But COGROUP is used when multiple relations are involved

move data to HDFS

```
[edureka@localhost ~]$ hadoop dfs -put Desktop/Datasets/input /pigInput  
DEPRECATED: Use of this script to execute hdfs command is deprecated.  
Instead use the hdfs command for it.  
  
16/12/23 13:16:14 WARN util.NativeCodeLoader: Unable to load native-hadoop library for  
native classes where applicable  
[edureka@localhost ~]$
```

Sample Data

```
input  
111,John,Sales,Austin  
222,Alex,Marketing,New York  
333,Philip,Operation,Sacramento  
444,Terry,Sales,New York  
555,Jessi,Development,Boston|
```

Data loaded in HDFS

```
native classes where applicable  
[edureka@localhost ~]$ hadoop dfs -cat /pigInput  
DEPRECATED: Use of this script to execute hdfs command is deprecated.  
Instead use the hdfs command for it.  
  
16/12/23 13:17:23 WARN util.NativeCodeLoader: Unable to load native-hadoop  
native classes where applicable  
111,John,Sales,Austin  
222,Alex,Marketing,New York  
333,Philip,Operation,Sacramento  
444,Terry,Sales,New York  
555,Jessi,Development,Boston|  
[edureka@localhost ~]$
```

Execute Pig in MapReduce mode



```
[edureka@localhost ~]$ pig  
16/12/23 13:18:34 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL  
16/12/23 13:18:34 INFO pig.ExecTypeProvider: Trying ExecType : MAPREDUCE  
16/12/23 13:18:34 INFO pig.ExecTypeProvider: Picked MAPREDUCE as the ExecType  
2016-12-23 13:18:35,694 [main] INFO org.apache.pig.Main - Apache Pig version 0.16.0 (r1746530)  
2016, 23:10:49  
2016-12-23 13:18:35,698 [main] INFO org.apache.pig.Main - Logging error messages to: /home/edureka/15655.log  
2016-12-23 13:18:36,178 [main] INFO org.apache.pig.impl.util.Utils - Default bootup file /home/edureka/.pigrc not found
```

Load data from HDFS to Pig (you can also load data without specifying schema)

```
grunt>  
grunt> employee = LOAD '/pigInput' using PigStorage(',') AS ( ssn:chararray, name:chararray, department:chararray,  
, city:chararray);
```

```
2016-12-23 13:23:30,934 [main] INFO org.apache.hadoop.conf  
cated. Instead, use fs.defaultFc  
grunt> dump employee;
```

(Dump Data on screen)

```
2016-12-23 13:28:04,667 [main] INFO org.apache.hadoop.hdfs  
paths to process : 1  
(111,John,Sales,Austin)  
(222,Alex,Marketing,New York)  
(333,Philip,Operation,Sacramento)  
(444,Terry,Sales,New York)  
(555,Jessi,Development,Boston)  
grunt>
```

Use of “foreach”

```
(222,Alex,Marketing,New York)
(333,Philip,Operation,Sacramento)
(444,Terry,Sales,New York)
(555,Jessi,Development,Boston)
grunt> emp foreach = foreach employee generate name,department;
```

Selecting two column from every ROW

```
o process : 1
2016-12-23 13:35:23,637 [main] INFO org.apache.pig.
paths to process : 1
(John,Sales)
(Alex,Marketing)
(Philip,Operation)
(Terry,Sales)
(Jessi,Development)
grunt>
```

Filter Employee from Austin -- Filter operation

```
(Terry,Sales)
(Jessi,Development)
grunt> emp filter = filter employee by city == 'Austin';
```

```
.. will not generate code.
2016-12-23 13:40:18,827 [main] INFO org.apache.
o process : 1
2016-12-23 13:40:18,827 [main] INFO org.apache.
paths to process : 1
(111,John,Sales,Austin) I
grunt>
```

```

o process : 1
2016-12-23 13:40:18,827 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher
paths to process : 1
(111,John,Sales,Austin)
grunt> emp_order = order employee by ssn desc; ORDER Operation

```

[Browsing HDFS - Mozilla Firefox edureka@localhost:~ [Datasets - File Brow... queries (~/Desktop) - ...]

```

.. will not generate code.
2016-12-23 13:55:31,748 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat
o process : 1
2016-12-23 13:55:31,748 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat
paths to process : 1
(555,Jessi,Development,Boston)
(444,Terry,Sales,New York)
(333,Philip,Operation,Sacramento)
(222,Alex,Marketing,New York)
(111,John,Sales,Austin)
grunt>

```

[Browsing HDFS - Mozilla Firefox edureka@localhost:~ [Datasets - File Brow...]

(111,John,Sales,Austin)

Store Data in HDFS

```

grunt>
grunt>
grunt> STORE emp_order into '/pigresult';

```

[Browsing HDFS - Mozilla Firefox edureka@localhost:~ [Datasets - File Brow... queries (~/Desktop) - ...]

drwxr-xr-x	edureka	supergroup	0 B	12/14/2016, 3:29:26 PM	0	0 B	out_card
drwxr-xr-x	edureka	supergroup	0 B	12/14/2016, 3:34:22 PM	0	0 B	out_card1
drwxr-xr-x	edureka	supergroup	0 B	12/22/2016, 3:24:52 PM	0	0 B	output
-rw-r--r--	edureka	supergroup	136 B	12/23/2016, 1:16:35 PM	1	128 MB	pigInput
drwxr-xr-x	edureka	supergroup	0 B	12/23/2016, 2:01:43 PM	0	0 B	pigresult
-rw-r--r--	edureka	supergroup	96.94 KB	12/20/2016, 9:16:07 PM	1	128 MB	sample.log
drwx-----	edureka	supergroup	0 B	12/23/2016, 1:24:11 PM	0	0 B	tmp
drwxr-xr-x	edureka	supergroup	0 B	12/20/2016, 5:02:55 PM	0	0 B	user
drwxr-xr-x	edureka	supergroup	0 B	12/16/2016, 1:53:05 PM	0	0 B	weather_out

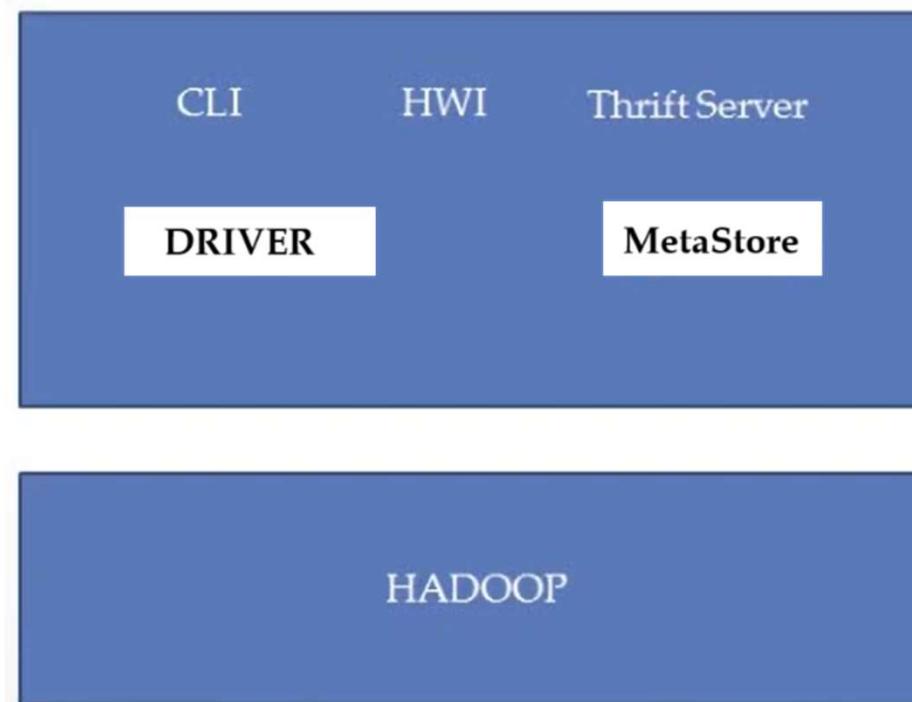


HIVE

- Hive Architecture
- Hive Components
- Metastore
- Limitations of Hive
- Where to use Hive
- Differences with traditional RDBMS
- Type System
- Hive Data Model



Architecture





Components

Shell - Environment

MetaStore - Metadata

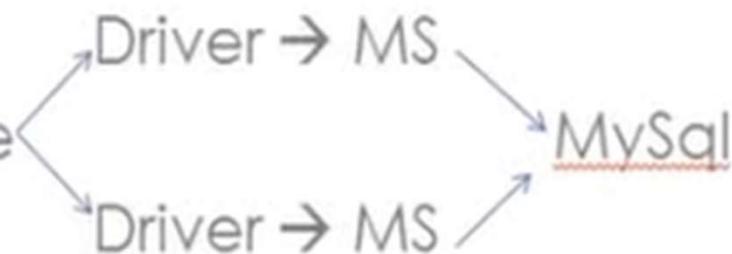
Driver- Management of Life Cycle of Hql Statements

Compiler- Convert Hql to MR

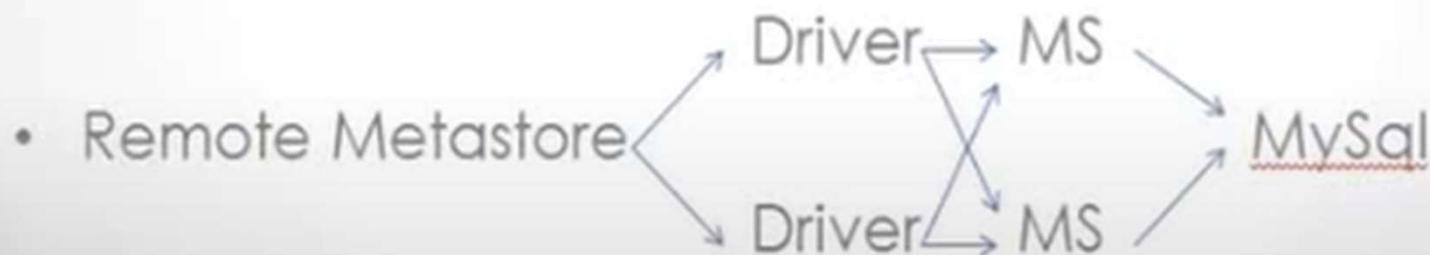
Execution Engine- Execution of Job

Metastore

- Embedded metastore → Driver → MS → Derby



- Local Metastore



- Remote Metastore



Limitations

- Not recommended for row level updates .
- Latency for Hive Queries is High
- Not designed for OLTP



Difference w.r.t RDBMS

Schema on Read V/s Schema on Write



Data Types

- Boolean :- true/false
- Integers :- TinyInt - 1 byte integer
SmallInt - 2 byte integer
INT - 4 byte integer
BigInt - 8 byte integer
- Floating point :- Float / Double
- String



Data Types

NUMERIC TYPES	DESCRIPTION	DATE/TIME TYPES	DESCRIPTION
TINYINT	1-byte signed integer, from -128 to 127	TIMESTAMP	Accepts Both Date and Time
SMALLINT	2-byte signed integer, from -32,768 to 32,767	DATE	Accepts just Date
INT/INTEGER	4-byte signed integer, from -2,147,483,648 to 2,147,483,647	INTERVAL	Interval
BIGINT	8-byte signed integer, from -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807		
FLOAT	4-byte single precision floating point number		
DOUBLE	8-byte double precision floating point number		
DOUBLE PRECISIO N	Alias for DOUBLE, only available starting with Hive 2.2.0		
DECIMAL	It accepts a precision of 38 digits.		
NUMERIC	Same as DECIMAL type.		



Data Types

STRING TYPES	DESCRIPTION
STRING	The string is an unbounded type. Not required to specify the lenght. It can accept max up to 32,767 bytes.
VARCHAR	Variable length of characters. It is bounded meaning you still need to specify the length like VARCHAR(10).
CHAR	Fixed length of Characters. if you define char(10) and assigning 5 chars, the remaining 5 characters space will be wasted.



Complex Data Types

Struts

Maps

Array



ACID -limitations

- To support ACID, Hive tables should be created with **TRANSACTIONAL** table property.
- Currently, Hive supports ACID transactions on tables that store **ORC** file format.
- Enable ACID support by setting transaction manager to **DbTxnManager**
- Transaction tables cannot be accessed from the non-ACID Transaction Manager (**DummyTxnManager**) session.
- **External tables** cannot be created to support ACID since the changes on external tables are beyond Hive control.
- **LOAD** is not supported on ACID transactional Tables. hence use **INSERT INTO**.
- On the Transactional session, all operations are auto commit as **BEGIN, COMMIT,** and **ROLLBACK** are not yet supported.



Data Models

- Database
- Table
- Partition
- Buckets & Clusters

Partitioning - Static

```
hive> create database part ;          Create database
OK
Time taken: 0.439 seconds
hive> set hive.cli.print.current.db=true;
hive (default)> use part ;           Use database already created
OK
Time taken: 0.042 seconds
hive (part)> create table student(name string,rollno int,per float)
                                         Create table in database
```



Partition Info

```
hive (default)> use part ;
OK
Time taken: 0.042 seconds
hive (part)> create table student(name string,rollno int,per float)
    > partitioned by (state string,city string) Partition column
    > row format delimited
    > fields terminated by ',';
OK
Time taken: 0.513 seconds
hive (part)> describe student ;
OK
name          string
rollno        int
per           float
state         string
city          string
Partition Information Partition Column
# col_name      data_type            comment
state          string
city           string
Time taken: 0.327 seconds, Fetched: 11 row(s)
hive (part)>
```

localhost:50070/explorer.html#/user/hive/warehouse

Hadoop Overview Datanodes Snapshot Startup Progress Utilities

Browse Directory

Brows

/user/hive/warehouse/part.db

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 01:37:14 PM IST	0	0 B	student

Hadoop 2016

BITS Pilani, Pilani Campus

Inserting Data in Partitioned Table

Partitioning.txt x partstud mah x

```
sdp,12,75.2
xcd,45,45.5
dss,47,74.2
fss75,90
qdz,5,32
dds,55,78.2
gff,66,85.4
jhh,22,65.02
sd,45,14.04
sf,15,78.87
vdf,22,65.21
fdf,24,75.25
sdsd,2,95.23
dsd,20,71.05
cFDF 20 60 AA
```

Trying to insert this data from "partstud.mah"
(Name, Rollno, Percentage)

```
hive (part)> load data local inpath '/home/hadoop/Desktop/HadoopData/partstudmah' into table student
      > partition(state="maharashtra",city="aurangabad");
Loading data to table part.student partition (state=maharashtra, city=aurangabad)
OK
Time taken: 2.259 seconds
```

hive (part)> load data local inpath '/home/hadoop/Desktop/HadoopData/partstudmah' into table student
 > partition(state="maharashtra",city="aurangabad");
Loading data to table part.student partition (state=maharashtra, city=aurangabad)
OK
Time taken: 2.259 seconds

hive (part)> load data local inpath '/home/hadoop/Desktop/HadoopData/partstudgujarat' into table student
 > partition(state="gujrat",city="surat");
Loading data to table part.student partition (state=gujrat, city=surat)
OK
Time taken: 0.674 seconds

second partition on partstudgujarat

/user/hive/warehouse/part.db/student							
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 01:46:27 PM IST	0	0 B	state=gujrat
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 01:43:47 PM IST	0	0 B	state=maharashtra

/user/hive/warehouse/part.db/student/state=gujrat/city=surat							
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rwxr-xr-x	hadoop	supergroup	2.93 KB	Thu 22 Jun 2017 01:46:27 PM IST	1	122.07 MB	partstudgujarat

BITS Pilani, Pilani Campus

Partitioning Dynamic

```
Loading data to table part.student partition (state=gujrat, city=surat)
OK
Time taken: 0.674 seconds
hive (part)> select * from student limit 10 ;
OK
sdp    12      75.2    gujrat  surat
xcd    45      45.5    gujrat  surat
dss    47      74.2    gujrat  surat
fss75  90      NULL     gujrat  surat
qdz    5       32.0    gujrat  surat
dds    55      78.2    gujrat  surat
gff    66      85.4    gujrat  surat
jhh    22      65.02   gujrat  surat
sd     45      14.04   gujrat  surat
sf     15      78.87   gujrat  surat
Time taken: 2.61 seconds, Fetched: 10 row(s)
```

Data in student table

```
hive (part)> set hive.exec.dynamic.partition=true ;
hive (part)> set hive.exec.dynamic.partition.mode=nonstrict
Enable Dynamic Partition with these two statements
```

Load file data (3 files) –Static mode

```

hive (part)> set hive.exec.dynamic.partition=true ;
hive (part)> set hive.exec.dynamic.partition.mode=nonstrict;
hive (part)> create table student1(name string,rollno int,per float,state string,city string)
    > row format delimited
    > fields terminated by ',';
OK
Time taken: 0.13 seconds
hive (part)> load data local inpath '/home/hadoop/Desktop/HadoopData/partstud1' into table student
    > ;
FAILED: SemanticException [Error 10062]: Need to specify partition columns because the destination table is partitioned
hive (part)> load data local inpath '/home/hadoop/Desktop/HadoopData/partstud1' into table student1;
Loading data to table part.student1
OK
loading data from student1 , student2 and student3
Time taken: 0.399 seconds
hive (part)> load data local inpath '/home/hadoop/Desktop/HadoopData/partstud2' into table student1;
Loading data to table part.student1
OK
Time taken: 0.355 seconds
hive (part)> load data local inpath '/home/hadoop/Desktop/HadoopData/partstud3' into table student1;

```

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rwxr-xr-x	hadoop	supergroup	8.84 KB	Thu 22 Jun 2017 01:59:44 PM IST	1	122.07 MB	partstud1
-rwxr-xr-x	hadoop	supergroup	8.84 KB	Thu 22 Jun 2017 01:59:53 PM IST	1	122.07 MB	partstud2
-rwxr-xr-x	hadoop	supergroup	8.84 KB	Thu 22 Jun 2017 02:00:05 PM IST	1	122.07 MB	partstud3

Time taken: 0.418 seconds
 hive (part)> select count(*) from student1;
WARNING: Hive on MR is deprecated in Hive 2 and may not be available in park, tez) or using Hive 1.X releases.
 Query ID = hadoop_20170622140043_366dacff-7545-4d9d-98e6-5ace8bde5a8a
 Total jobs = 1
 Launching Job 1 out of 1
 Number of reduce tasks determined at compile time: 1
 In order to change the average load for a reducer (in bytes):
 set hive.exec.reducers.bytes.per.reducer=<number>
 In order to limit the maximum number of reducers:
 set hive.exec.reducers.max=<number>
 In order to set a constant number of reducers:
 set mapreduce.job.reduces=<number>
 Job running in-process (local Hadoop)
 2017-06-22 14:00:48,527 Stage-1 map = 0%, reduce = 0%
 2017-06-22 14:00:49,557 Stage-1 map = 100%, reduce = 100%
 Ended Job = job_local1879476739_0001
 MapReduce Jobs Launched:
 Stage-Stage-1: HDFS Read: 63236 HDFS Write: 63236 SUCCESS
 Total MapReduce CPU Time Spent: 0 msec
 OK
 936
Totally 936 records inserted
 Time taken: 5.725 seconds, Fetched: 1 row(s)

Insert data from table to partitioned table

```
hive (part)> insert into stud_part
>   partition(state,city)
>   select name,rollno,per
> ,state,
> city           Insert data from stuент1 to "stud_part" table
>   from student1; having partition on state and city
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available
park, tez) or using Hive 1.X releases.
Query ID = hadoop_20170622140910_9f43aa75-98c0-448b-84b3-f36b6c0
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operation
Job running in-process (local Hadoop)
2017-06-22 14:09:13,096 Stage-1 map = 0%,  reduce = 0%
2017-06-22 14:09:14,111 Stage-1 map = 100%,  reduce = 0%
Ended Job = job_local1776471965_0002
```

Data Presentation in dynamic mode

/user/hive/warehouse/part.db							Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 02:09:15 PM IST	0	0 B	stud_part
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 01:46:27 PM IST	0	0 B	student
/user/hive/warehouse/part.db/stud_part							Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 02:09:14 PM IST	0	0 B	state=gujrat
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 02:09:14 PM IST	0	0 B	state=maharashtra
/user/hive/warehouse/part.db/stud_part/state=gujrat							Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 02:09:13 PM IST	0	0 B	city=surat
/user/hive/warehouse/part.db/stud_part/state=gujrat/city=surat							Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rwxr-xr-x	hadoop	supergroup	5.94 KB	Thu 22 Jun 2017 02:09:13 PM IST	1	122.07 MB	000000_0

Transactional table

```
CREATE TABLE emp.employee_trans (
    id int,
    name string,
    age int,
    gender string)
STORED AS ORC
TBLPROPERTIES ('transactional'='true');
```

You can run the `DESCRIBE FORMATTED emp.employee` to check if the table is created with the transactional data_type as `TRUE`.

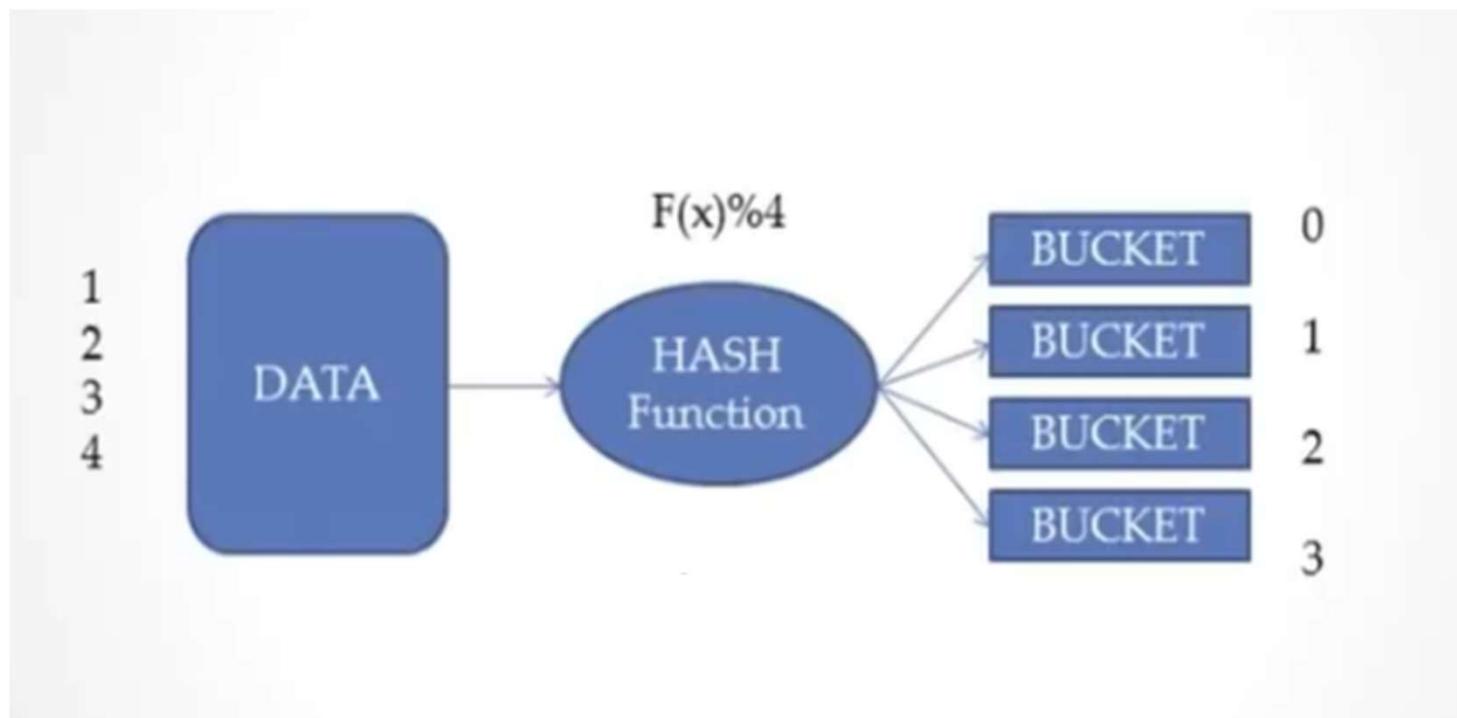
col_name	data_type	comment
# col_name		
id	int	NULL
name	string	NULL
age	int	NULL
gender	string	NULL
# Detailed Table Information		
Database:		
OwnerType:		
Owner:		
CreateTime:		
LastAccessTime:		
Retention:		
Location:		
Table Type:	MANAGED_TABLE	NULL
Table Parameters:		
	NULL	NULL
	bucketing_version	2
	numFiles	4
	totalSize	3356
	transactional	true
	transactional_properties	default
	transient_lastDdlTime	1602914461
# Storage Information		
SerDe Library:		
InputFormat:		
OutputFormat:		



Bucketing

- Data is divided into buckets
- Used for Data Sampling
- Helpful for implementing Map-side Join

Hashing Technique





```
hadoop@Sandeep:~$ jps      Start Hadoop
4448 NodeManager
4135 SecondaryNameNode
4313 ResourceManager
3786 NameNode
3915 DataNode
8523 Jps
hadoop@Sandeep:~$ hive      Start Hive

Logging initialized using configuration in jar:file:/home/hadoop/Downloads/apache-hive-2.1.1-bin/lib/hive
es Async: true
Java HotSpot(TM) Server VM warning: You have loaded library /home/hadoop/Downloads/hadoop-2.7.3/lib/nati
isabled stack guard. The VM will try to fix the stack guard now.
It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noe
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a dif
) or using Hive 1.X releases.
hive> create database bucketeddb;      Create database
OK
Time taken: 1.597 seconds
hive> use bucketeddb;
OK
Time taken: 0.045 seconds
hive> create table sada_table1(id int,firstname string,lastname string)      Create simple table
    > row format delimited
    > fields terminated by ','
    > stored as textfile ;
OK
Time taken: 0.575 seconds
hive> load data local inpath '/home/hadoop/Desktop/HadoopData/bd1' into table sada_table ;
FAILED: SemanticException [Error 10001]: Line 1:72 Table not found 'sada_table'
hive> load data local inpath '/home/hadoop/Desktop/HadoopData/bd1' into table sada_table1 ;
Loading data to table bucketeddb.sada_table1      load data in sada_table from bd1
OK
Time taken: 1.163 seconds
hive> █
```



```
hive> create table bucketwala_table1(id int,firstname string,lastname string)
  > clustered by (id) into 5 buckets
  > row format delimited
  > fields terminated by ','
  > stored as textfile;
          Create bucketed table
OK
Time taken: 0.191 seconds
hive> insert overwrite table bucketwala_table1 select * from sada_table1;      insert data in bucketed table from sada_table
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (spark, tez) or using Hive 1.X releases.
Query ID = hadoop_20170823134810_55a07640-70f3-4d33-9e0a-afab8220cb96
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2017-08-23 13:48:14,765 Stage-1 map = 100%,  reduce = 0%
2017-08-23 13:48:15,797 Stage-1 map = 100%,  reduce = 100%
Ended Job = job_local1899406628_0001
>Loading data to table bucketeddb.bucketwala_table1
MapReduce Jobs Launched:
  Stage-Stage-1: HDFS Read: 2064 HDFS Write: 4374 SUCCESS
  Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 6.241 seconds
```



/user/hive/warehouse/bucketeddb.db

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Wed 23 Aug 2017 01:48:16 PM IST	0	0 B	bucketwala_table1
drwxr-xr-x	hadoop	supergroup	0 B	Wed 23 Aug 2017 01:43:34 PM IST	0	0 B	sada_table1

ble1/000000_0

Data in Bucket ZERO

Java Hotspot(TM) Server VM warning: You have
ds/hadoop-2.7.3/lib/native/libhadoop.so.1.
ard. The VM will try to fix the stack guard
It's highly recommended that you fix the
or link it with '-z noexecstack'.
17/08/23 13:51:20 WARN util.NativeCodeLoad
ry for your platform... using builtin-jav

/user/hive/warehouse/bucketeddb.db/sada_table1

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rwxr-xr-x	hadoop	supergroup	344 B	Wed 23 Aug 2017 01:43:34 PM IST	1	122.07 MB	bd1

/user/hive/warehouse/bucketeddb.db/bucketwala_table1

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rwxr-xr-x	hadoop	supergroup	70 B	Wed 23 Aug 2017 01:48:14 PM IST	1	122.07 MB	000000_0
-rwxr-xr-x	hadoop	supergroup	75 B	Wed 23 Aug 2017 01:48:15 PM IST	1	122.07 MB	000001_0
-rwxr-xr-x	hadoop	supergroup	69 B	Wed 23 Aug 2017 01:48:15 PM IST	1	122.07 MB	000002_0
-rwxr-xr-x	hadoop	supergroup	63 B	Wed 23 Aug 2017 01:48:15 PM IST	1	122.07 MB	000003_0
-rwxr-xr-x	hadoop	supergroup	67 B	Wed 23 Aug 2017 01:48:15 PM IST	1	122.07 MB	000004_0

20,Sandeep,Patil
15,Ashwini,Vispute
10,Bhausaheb,Khamat
5,Vikas,Joshi
hadoop@Sandeep:~\$

Partitioning V/s Bucketing

/user/hive/warehouse

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Wed 23 Aug 2017 01:46:32 PM IST	0	0 B	bucketeddb.db
drwxr-xr-x	hadoop	supergroup	0 B	Wed 23 Aug 2017 01:31:09 PM IST	0	0 B	bucketwala_table1
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 02:04:03 PM IST	0	0 B	part.db
drwxr-xr-x	hadoop	supergroup	0 B	Wed 23 Aug 2017 01:26:49 PM IST	0	0 B	sada_table1

Bucket

Partition

Replication	Block Size	Name	Replication	Block Size	Name
0	0 B	stud_part	0	0 B	state=gujrat
0	0 B	student	0	0 B	state=maharashtra
0	0 B	student1			

/user/hive/warehouse/part.db/stud_part/state=gujrat

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Thu 22 Jun 2017 02:09:13 PM IST	0	0 B	city=surat

city=surat

/user/hive/warehouse/part.db/stud_part/state=gujrat/city=surat

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rwxr-xr-x	hadoop	supergroup	5.94 KB	Thu 22 Jun 2017 02:09:13 PM IST	1	122.07 MB	000000_0



/user/hive/warehouse/bucketeddb.db	Go!
------------------------------------	-----

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	hadoop	supergroup	0 B	Wed 23 Aug 2017 01:48:16 PM IST	0	0 B	bucketwala_table1
drwxr-xr-x	hadoop	supergroup	0 B	Wed 23 Aug 2017 01:43:34 PM IST	0	0 B	sada_table1

/user/hive/warehouse/bucketeddb.db/bucketwala_table1	Go!
--	-----

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rwxr-xr-x	hadoop	supergroup	70 B	Wed 23 Aug 2017 01:48:14 PM IST	1	122.07 MB	000000_0
-rwxr-xr-x	hadoop	supergroup	75 B	Wed 23 Aug 2017 01:48:15 PM IST	1	122.07 MB	000001_0
-rwxr-xr-x	hadoop	supergroup	69 B	Wed 23 Aug 2017 01:48:15 PM IST	1	122.07 MB	000002_0
-rwxr-xr-x	hadoop	supergroup	63 B	Wed 23 Aug 2017 01:48:15 PM IST	1	122.07 MB	000003_0
-rwxr-xr-x	hadoop	supergroup	67 B	Wed 23 Aug 2017 01:48:15 PM IST	1	122.07 MB	000004_0



Table Types

- Internal Table
- External Table



Internal/Managed & External Table

Internal tables are also known as Managed tables that are owned and managed by Hive. By default, Hive creates a table as an Internal table and owned the table structure and the files.

In other words, Hive completely manages the lifecycle of the table (metadata & data) similar to tables in RDBMS.

For Internal tables, Hive by default stores the files at the [data warehouse location](#) which is located at
[/user/hive/warehouse](#)

When you drop an internal table, it drops the data and also drops the metadata of the table.

Below is an example of creating internal table.



Internal / Managed Table

```
CREATE TABLE IF NOT EXISTS emp.employee (
  id int,
  name string,
  age int,
  gender string )
COMMENT 'Employee Table'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';
```

Use `DESCRIBE FORMATTED emp.employee;` to get the description of the table and you should see Table Type as **MANAGED TABLE**.

col_name	data_type	comment
# col_name		
id	int	NULL
name	string	NULL
age	int	NULL
gender	string	NULL
# Detailed Table Information		
Database:	NULL	NULL
OwnerType:	emp	NULL
Owner:	USER	NULL
CreateTime:	prabha	NULL
LastAccessTime:	Thu Oct 15 04:58:36 UTC 2020	NULL
Retention:	UNKNOWN	NULL
Location:	0	NULL
Table Type:	hdfs://192.168.1.148:9000/user/hive/warehouse/emp.db/employee MANAGED_TABLE	NULL
Table Parameters:	NULL	NULL



External Table

Data in External tables are not owned or managed by Hive. To create an External table you need to use `EXTERNAL` clause.

Hive default stores external table files also at [Hive managed data warehouse location](#) but recommends to use external location using `LOCATION` clause.

Dropping an external table just drops the metadata but not the actual data. The actual data is still accessible outside of Hive.

Below is an example of creating an external table in Hive. If you noticed we use `EXTERNAL` and `LOCATION` options.

External Table

```
CREATE EXTERNAL TABLE emp.employee_external (
    id int,
    name string,
    age int,
    gender string)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/user/hive/data/employee_external';
```

Use `DESCRIBE FORMATTED emp.employee_external;` to get the description of the table and you should see Table Type as `EXTERNAL_TABLE`.

col_name	<th>comment</th>	comment
# col_name	data_type	
id	int	NULL
name	string	NULL
age	int	NULL
gender	string	NULL
# Detailed Table Information		
Database:		
OwnerType:	emp	NULL
Owner:	USER	NULL
CreateTime:	prabha	NULL
LastAccessTime:	Thu Oct 15 05:08:10 UTC 2020	NULL
Retention:	UNKNOWN	NULL
Location:	0	NULL
Table Type:	hdfs://192.168.1.148:9000/user/hive/warehouse/emp.db/employee_external EXTERNAL_TABLE	NULL
Table Parameters:	COLUMN_STATS_ACCURATE	{"BASIC_STATS": "true", "CO



Dropping table

Regardless of the Internal and external table, Hive manages the table definition and its partition information in Hive Metastore. Dropping an internal table deletes the table metadata from Metastore and also removes all its data/files from HDFS.

Dropping an external table, just drop the metadata of the table from Metastore and keeps the actual data as-is on HDFS location.



Key differences – Summary

INTERNAL OR MANAGED TABLE	EXTERNAL TABLE
By default, Hive creates an Internal or Managed Table.	Use EXTERNAL option/clause to create an external table
Hive owns the metadata, table data by managing the lifecycle of the table	Hive manages the table metadata but not the underlying file.
Dropping an Internal table drops metadata from Hive Metastore and files from HDFS	Dropping an external table drops just metadata from Metastore without touching actual file on HDFS.
Hive supports ARCHIVE, UNARCHIVE, TRUNCATE, MERGE, CONCATENATE operations	Not supported
Supports ACID/Transactional	Not supported
Supports result caching	Not supported



Hadoop...Hbase....Hive....Pig

Hadoop:

- Hadoop is basically 2 things: a **Distributed FileSystem (HDFS)** + **Computation or Processing framework (MapReduce)**.
- Like all other FS, HDFS also provides us storage, but in a fault tolerant manner with high throughput and lower risk of data loss (because of the replication). But, being a FS, HDFS lacks **random read and write access**.
- This is where HBase comes into picture. It's a **distributed, scalable, big data store**, modelled after Google's BigTable. It stores data as key/value pairs.

HBase:

- Apache HBase is an open source NoSQL database that provides real-time read/write access to those large datasets.
- HBase scales linearly to handle huge data sets with billions of rows and millions of columns, and it easily combines data sources that use a wide variety of different structures and schemas.
- HBase is natively integrated with Hadoop and works seamlessly alongside other data access engines through YARN.

Hive:

- It provides us data warehousing facilities on top of an existing Hadoop cluster. Along with that it provides an SQL like interface which makes work easier.
- While Pig is basically a **dataflow language** that allows us to process enormous amounts of data very easily and quickly.
- Pig basically has 2 parts: the Pig **Interpreter** and the language, **PigLatin**. We can write Pig script in PigLatin and using Pig interpreter process them.
- Pig makes our life a lot easier, otherwise writing MapReduce is always not easy.



Pig V/s Hive

S.No.	Pig	Hive
1.	Pig operates on the client side of a cluster.	Hive operates on the server side of a cluster.
2.	Pig uses pig-latin language.	Hive uses HiveQL language.
3.	Pig is a Procedural Data Flow Language.	Hive is a Declarative SQLish Language.
4.	It was developed by Yahoo.	It was developed by Facebook.
5.	It is used by Researchers and Programmers.	It is mainly used by Data Analysts.
6.	It is used to handle structured and semi-structured data.	It is mainly used to handle structured data.
7.	It is used for programming.	It is used for creating reports.
8.	Pig scripts end with .pig extension.	In Hive, all extensions are supported.
9.	It does not support partitioning.	It supports partitioning.
10.	It loads data quickly.	It loads data slowly.
11.	It does not support JDBC.	It supports <u>JDBC</u>.
12.	It does not support ODBC.	It supports <u>ODBC</u>.
13.	Pig does not have a dedicated metadata database.	Hive makes use of the exact variation of dedicated SQL-DDL language by defining tables beforehand.



Hive V/s HBase

S. No.	Parameters	Hive	HBase
1.	Basics	Hive is a query engine that uses queries that are mostly similar to SQL queries.	It is Data storage, particularly for unstructured data.
2.	Used for	It is mainly used for batch processing (that means OLAP-based).	It is extensively used for transactional processing (that means OLTP).
3.	Processing	It cannot be used for real-time processing since immediate analysis results are unable to obtain. In other words, the operations in Hive require batch processing, they normally take a long time to complete.	It can be used to process data in real-time. Transactional operations are faster than non-transactional operations (since HBase stores data in the form of key-value pairs).
4.	Queries	It is used only for analytical queries. It is mostly used to analyze Big Data.	It is used for real-time querying. It is mostly used to query Big Data.
5.	Runs on	Hive runs on the top of Hadoop.	HBase runs on the top of HDFS (Hadoop Distributed File System).
6.	Database	Apache Hive is not a database.	It supports the NoSQL database.
7.	Schema	It has a schema model.	It is free from the schema model.



Contd...

8.	Latency	Made for high latency operations as batch processing takes time.	Made for low-level latency operations.
9.	Cost	It is expensive as compared to HBase.	It is cost-effective as compared to Hive.
10.	Query Language	Hive uses HQL (Hive Query Language).	To conduct CRUD (Create, Read, Update, and Delete) activities, HBase does not have a specialized query language. HBase includes a Ruby-based shell where you can use Get, Put, and Scan functions to edit your data.
11.	Level of Consistency	Eventual consistency	Immediate consistency
12.	Secondary Indexes	It does not support Secondary Indexes.	It supports Secondary Indexes.
13.	Example	Hubspot	Facebook

Clear Difference

Pig Latin

```

countrys = load '/user/gharris/PIG_COUNTRIES' AS
    (country_id, country_name , country_subregion , region);

customers= load '/user/gharris/PIG_CUSTOMERS' AS
    (cust_id,first_name, last_name, gender, yob, marital, postcode,city,country_id);

asianCountryrs = filter countrys by region matches 'Asia';

joined = join customers by country_id, asianCountryrs by country_id;

grouped = group joined by country_name;

agged = foreach grouped generate group, COUNT(joined.customers::cust_id);

morethan500cust = filter agged by $1 > 500;

ordered=order morethan500cust by $1 desc;

dump ordered;

```

SQL or Hive QL

```

SELECT country_name,COUNT(cust_id) AS cust_count
FROM countries co
JOIN customers cu
ON (co.country_id=cu.country_id)
WHERE country_region='Asia'
GROUP BY country_name
HAVING COUNT(cust_id)>500
ORDER BY cust_count DESC

```



THAT IS ALL