



BITS Pilani presentation

BITS Pilani
Pilani Campus

Dr. Vivek V. Jog
Dept. of Computer Engineering

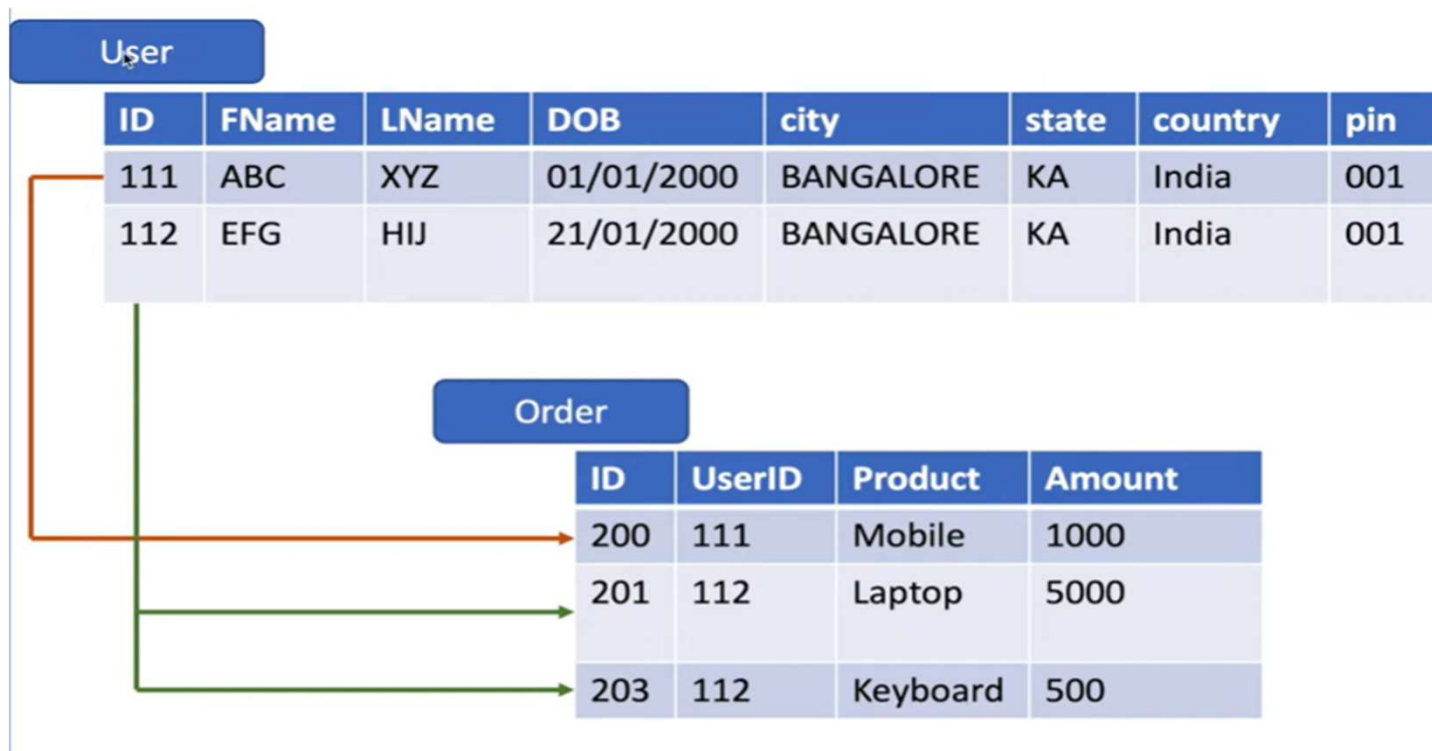


BITS Pilani
Pilani Campus

Big Data Systems (S1-24_CCZG522)

Lecture No.11

Traditional RDBMS



SQL Joins

SQL JOIN operations are used to combine rows from two or more tables, and are of two types:

- Conditional Join: Combine rows based on a condition over one or more common columns (typically primary or foreign keys).
- There are 4 types of conditional joins: inner, left, right, and full.
- Cross Join: Cartesian Product of two tables.

Inner Join

An Inner Join produces a row combining a row from both tables only if the key is present in *both* tables.

For example, if you had run two marketing campaigns on different channels that captured potential customer info in two tables:

- Table A: Name and phone number
- Table B: Name and email

Table A:

Name	Phone
Warren Buffett	+1-123-456-7890
Bill Gates	+1-987-654-3210

Table B:

Name	Email
Bill Gates	bgates@live.com
Larry Page	lpage@gmail.com

Inner Join Cont....

You can find all customer leads for which you have both their phone number and email address using an Inner Join:

SQL

```
1 SELECT A.Name, A.Phone, B.Email
2 FROM A
3 INNER JOIN B
4 ON A.name = B.name;
```

The word **INNER** is optional and you can omit it if you want. The result will be:

Code



Name	Phone	Email
Bill Gates	+1-987-654-3210	bgates@live.com

Left Outer Join

A Left Join returns all rows from the left table and the matched rows from the right table. Say, your sales team decides to reach out to all customer leads over the phone, but, if possible, wants to follow up by email too.

With a Left Join, you can create a list of customers where you have their phone numbers for sure but may also have their email:

SQL

```
1 SELECT A.Name, A.Phone, B.Email
2 FROM A
3 LEFT OUTER JOIN B
4 ON A.name = B.name;
```

The word **OUTER** is optional and you can omit it if you want. The result will be:

Code

Name	Phone	Email
Warren Buffett	+1-123-456-7890	null
Bill Gates	+1-987-654-3210	bgates@live.com

Right Outer Join

A Right Join returns all rows from the right table and the matched rows from the left table.

Say, you want the opposite: all customers with an email address and, if possible, their phone numbers too:

SQL

```
1 SELECT B.Name, A.Phone, B.Email
2 FROM A
3 RIGHT OUTER JOIN B
4 ON A.name = B.name;
```

The word **OUTER** is optional and you can omit it if you want. The result will be:

Code

Name	Phone	Email
Bill Gates	+1-987-654-3210	bgates@live.com
Larry Page	null	lpage@gmail.com

Full Outer Join

A Full Join returns all rows from both tables matching them whenever possible.

Say, you want to have a consolidated list of all customers having a phone number or email address or, if possible, both:

SQL

```
1 SELECT
2 CASE
3     WHEN A.Name IS NOT NULL THEN A.Name
4     ELSE B.Name
5 END AS Name,
6 A.Phone, B.Email
7 FROM A
8 FULL OUTER JOIN B
9 ON A.name = B.name;
```


The word **OUTER** is optional and you can omit it if you want. In some databases, a simple **SELECT *** might work and you will not need that **CASE** statement. The result will be:

Code

Name	Phone	Email
Warren Buffett	+1-123-456-7890	null
Bill Gates	+1-987-654-3210	bgates@live.com
Larry Page	null	lpage@gmail.com

Cross Join

- QL Cross Join is the Cartesian Product of two tables. It pairs each row of the left table with each row of the right table.
- Say, you have a table CarType with 7 rows having values: Sedan, Hatchback, Coupe, MPV, SUV, Pickup Truck, and Convertible.
- And you also have a table Color with 5 rows: Red, Black, Blue, Silver Grey, and Green.
- A Cross Join between the two will generate a table with 35 rows having all possible combinations of car types and colors:



Cross Join Example

SQL

```
1 SELECT *  
2 FROM CarType  
3 CROSS JOIN Color;
```

SQL Statement:

```
SELECT Customers.CustomerName, Orders.OrderID  
FROM Customers  
CROSS JOIN Orders;
```


For total 3 customers in customers table and 4 Orders in Orders table total 12 row output Will be generated

SQL Joins - Cheatsheet

SQL Joins Cheatsheet

INNER JOIN

SELECT <columns>
FROM A
[INNER] JOIN B
ON A.Key = B.Key;



Key	X	Y
1	x1	y1
2	x2	y2

LEFT JOIN

SELECT <columns>
FROM A
LEFT [OUTER] JOIN B
ON A.Key = B.Key;



Key	X	Y
1	x1	y1
2	x2	y2
3	x3	NULL

RIGHT JOIN


SELECT <columns>
FROM A
RIGHT [OUTER] JOIN B
ON A.Key = B.Key;



Key	X	Y
1	x1	y1
2	x2	y2
4	NULL	y3

FULL JOIN

SELECT <columns>
FROM A
FULL [OUTER] JOIN B
ON A.Key = B.Key;



Key	X	Y
1	x1	y1
2	x2	y2
3	x3	NULL
4	NULL	y3

POINTER

Data warehouses are Columnar databases, i.e., values in a column are stored together (in RDBMS, the rows are stored together).

Therefore, `SELECT *` is a particularly bad idea.

Real Time use of Row & Columnar DB -OLAP VS OLTP

RDBMS vs. Columnar: OLTP vs. OLAP

OLTP

D1	P1	S1	D2	P2	S2	D3	P3	S3	D4	P4	S4
----	----	----	----	----	----	----	----	----	----	----	----

- Rows are stored together
- Word-aligned heterogeneous values take more space

- Fast row reads and updates
- High ACID transaction throughput

- GBs of data, scale vertically
- Suitable for typical CRUD applications

Logical Table

Date	Product	Sale
D1	P1	S1
D2	P2	S2
D3	P3	S3
D4	P4	S4

RDBMS: Row-Oriented Datastore

Columnar: Column-Oriented Datastore

OLAP

D1	D2	D3	D4	P1	P2	P3	P4	S1	S2	S3	S4
----	----	----	----	----	----	----	----	----	----	----	----

Encoded Chunk

Encoded Chunk

Encoded Chunk

- Columns are stored together
- Type aware encoding aids compression

- Fast column summarizations
- High aggregation query throughput

- TBs of data, scale horizontally
- Suitable for analytics and machine learning

Why NoSQL

Relational databases are optimized for transaction operations. Transactions often update multiple records in multiple tables. Indexes are optimized for frequent low-latency writes of [ACID Transactions](#).


While transactions are on rows (records), analytics properties are computed on columns (attributes). OLAP applications need an optimized column-read operation on a table. Columnar databases are designed for high-throughput of column aggregations. That's why Columnar DBs are row-oriented databases.

However, the primary RDBMS operation is low-latency high-frequency ACID transactions. That does not scale to the Big Data scale common in analytics applications.


Why NOSQL now?? Ans. Driving Trends

11

New Trends



Big data



Connectivity



P2P Knowledge



Concurrency



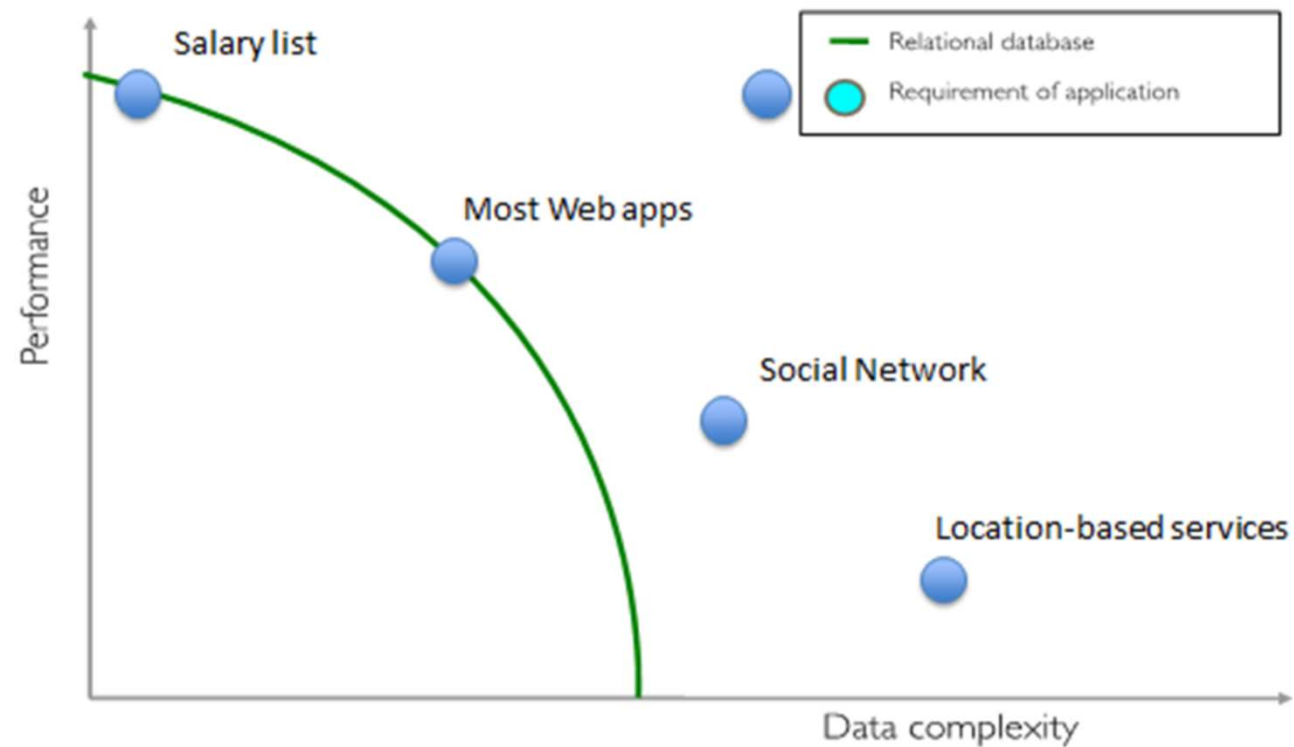
Diversity



Cloud-Grid

Side note: RDBMS performance

12



Contrast

A presentation slide for NoSQL with MongoDB. It features a dark background with a faint image of a plant. The title 'NoSQL' is in a large, bold, white font, with 'MongoDB' in a smaller, white font below it. A list of characteristics is presented in white text with bullet points.

NoSQL
MongoDB

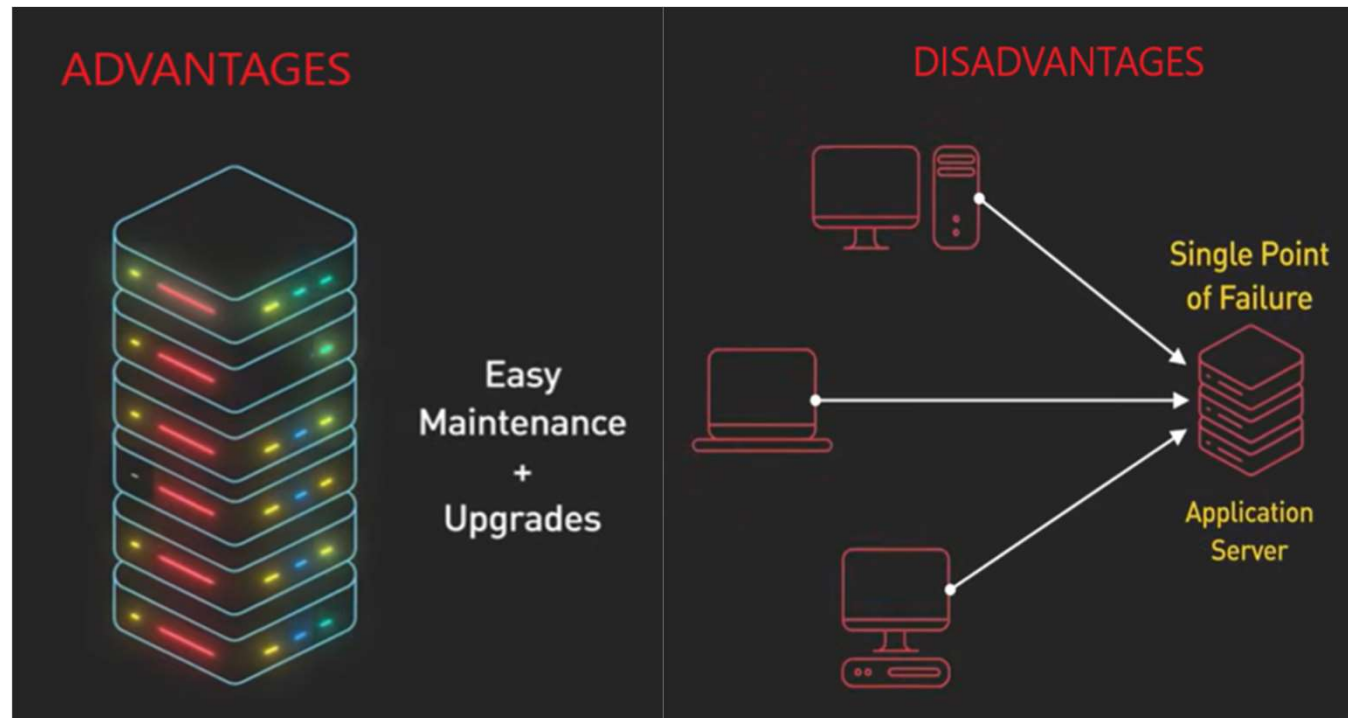

- > **Non - Relational**
- > **Query - JSON/BSON**
- > **Data Stored**
-Key Value & Collection
- > **Dynamic Schema**
- > **Horizontally Scalable**
- > **Follows CAP property**

A presentation slide for SQL. It features a dark background with a faint image of a plant. The title 'SQL' is in a large, bold, white font, with 'Programmer Mitra' in a smaller, white font above it. A list of characteristics is presented in white text with bullet points.




Programmer Mitra
SQL

- > **RDBMS**
- > **Query - SQL**
- > **Data Stored**
-Rows & Columns
- > **Static Schema**
- > **Vertically Scalable**
- > **Follows ACID property**


RDBMS = Vertical Scaling



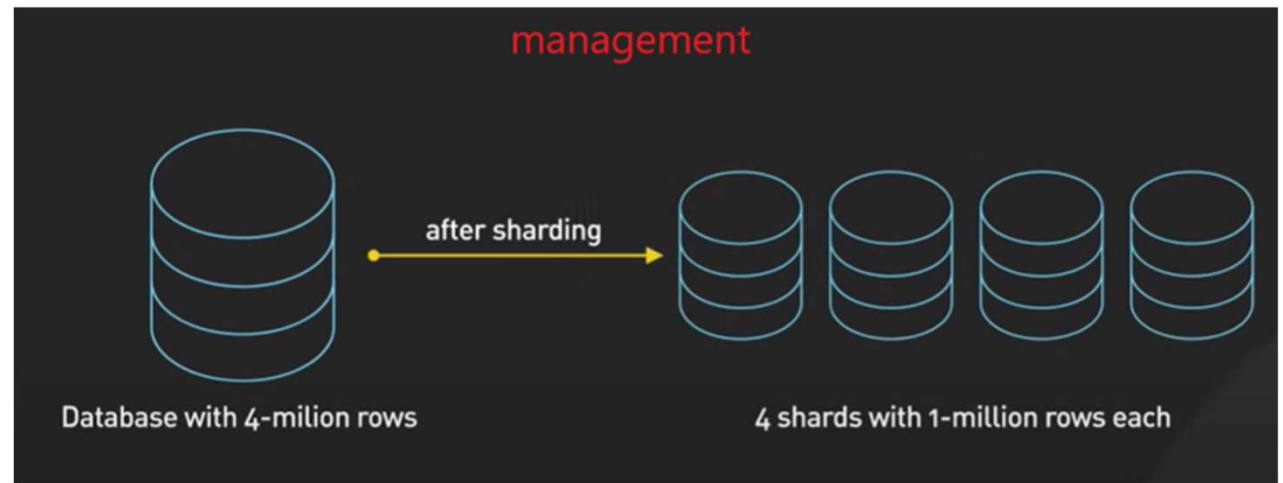

NoSQL = Horizontal Scaling




Cost...



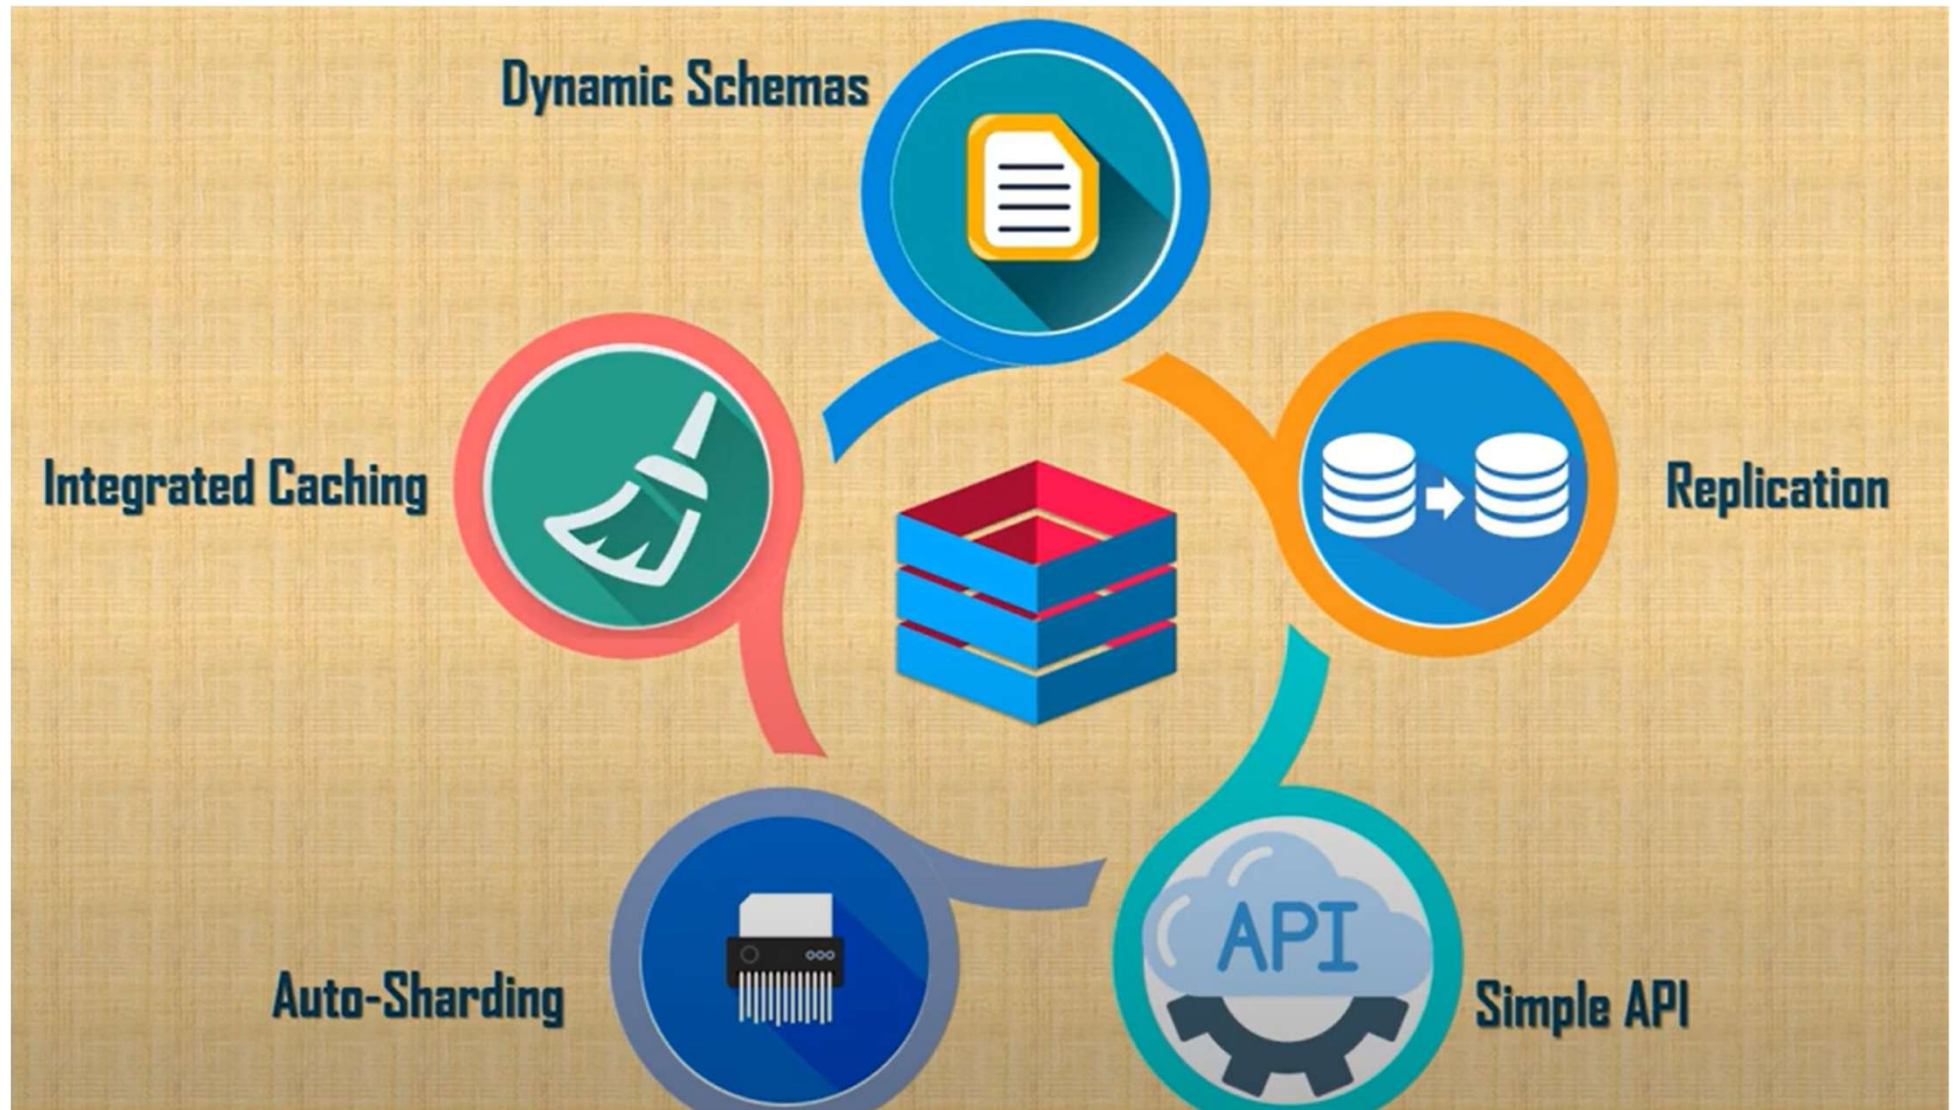
Disadvantages -- Horizontal



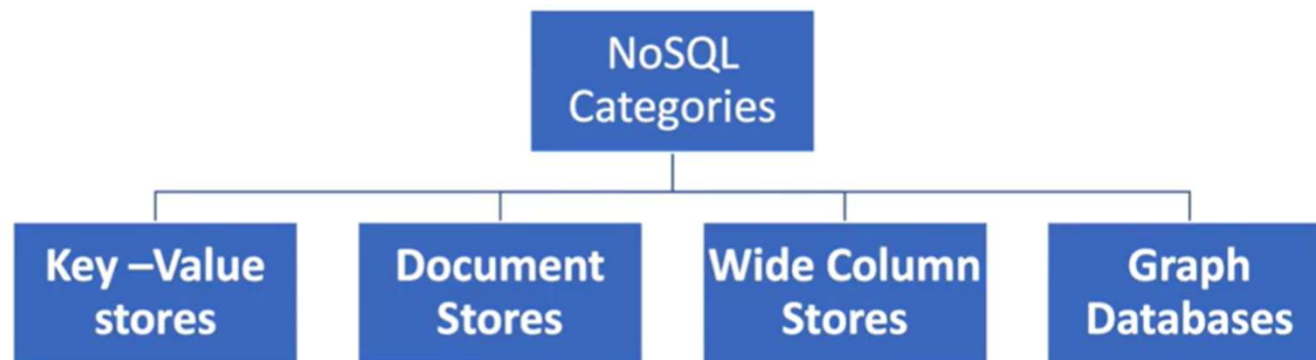
Disadvantages



NoSQL Features

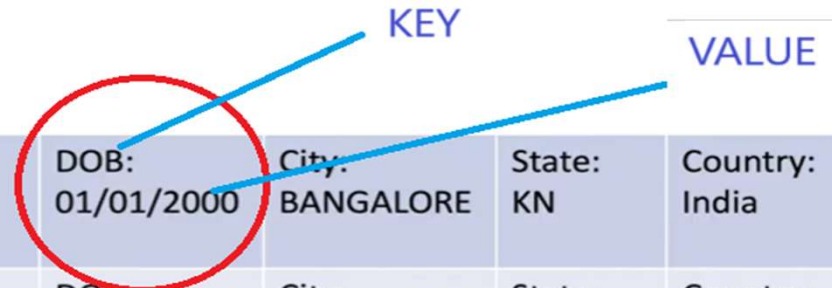


NoSQL DB Types



NoSQL DB

User							
ID: 111	Fname: ABC	Lname: XYZ	DOB: 01/01/2000	City: BANGALORE	State: KN	Country: India	Pin: 001
ID: 112	Fname: EFG	Lname: HIJ	DOB: 01/01/2001	City: NY	State: NY	Country: US	Zip: 002
ID: 113	Fname: EFG	Lname: HIJ	DOB: 01/01/2001	Village: NY	District: NY	Country: US	Postal code: 002



Order						
ID: 2001	UserID: 111	Product: Computer	ST TAX: 10	Amont: 1000	SHIPPING COST: 10	DISCOUNT: 5
ID: 2002	UserID: 112	Product: HEADSET	Amont: 100	SHIPPING COST: 10		
ID: 2003	UserID: 112	Product: MOBILE	Amont: 600	IMPORT TAX: 10		

Key-Value Pair

Database																									
<table><tr><th colspan="2">Table:User</th></tr><tr><td>ID</td><td>111</td></tr><tr><td>FName</td><td>ABC</td></tr><tr><td>LName</td><td>XYZ</td></tr><tr><td>city</td><td>BANGALORE</td></tr><tr><td>state</td><td>KA</td></tr><tr><td>country</td><td>India</td></tr></table>	Table:User		ID	111	FName	ABC	LName	XYZ	city	BANGALORE	state	KA	country	India	<table><tr><th colspan="2">Table:Order</th></tr><tr><td>ID</td><td>200</td></tr><tr><td>UserId</td><td>111</td></tr><tr><td>Product</td><td>Mobile</td></tr><tr><td>Amount</td><td>1000</td></tr></table>	Table:Order		ID	200	UserId	111	Product	Mobile	Amount	1000
Table:User																									
ID	111																								
FName	ABC																								
LName	XYZ																								
city	BANGALORE																								
state	KA																								
country	India																								
Table:Order																									
ID	200																								
UserId	111																								
Product	Mobile																								
Amount	1000																								
<table><tr><td>ID</td><td>111</td></tr><tr><td>FName</td><td>EFG</td></tr><tr><td>city</td><td>BANGALORE</td></tr><tr><td>state</td><td>KA</td></tr><tr><td>country</td><td>India</td></tr></table>	ID	111	FName	EFG	city	BANGALORE	state	KA	country	India	<table><tr><td>ID</td><td>201</td></tr><tr><td>UserId</td><td>112</td></tr><tr><td>Product</td><td>Laptop</td></tr><tr><td>Amount</td><td>5000</td></tr></table>	ID	201	UserId	112	Product	Laptop	Amount	5000						
ID	111																								
FName	EFG																								
city	BANGALORE																								
state	KA																								
country	India																								
ID	201																								
UserId	112																								
Product	Laptop																								
Amount	5000																								
	<table><tr><td>ID</td><td>203</td></tr><tr><td>UserId</td><td>112</td></tr><tr><td>Product</td><td>keyborad</td></tr><tr><td>Amount</td><td>500</td></tr></table>	ID	203	UserId	112	Product	keyborad	Amount	500																
ID	203																								
UserId	112																								
Product	keyborad																								
Amount	500																								

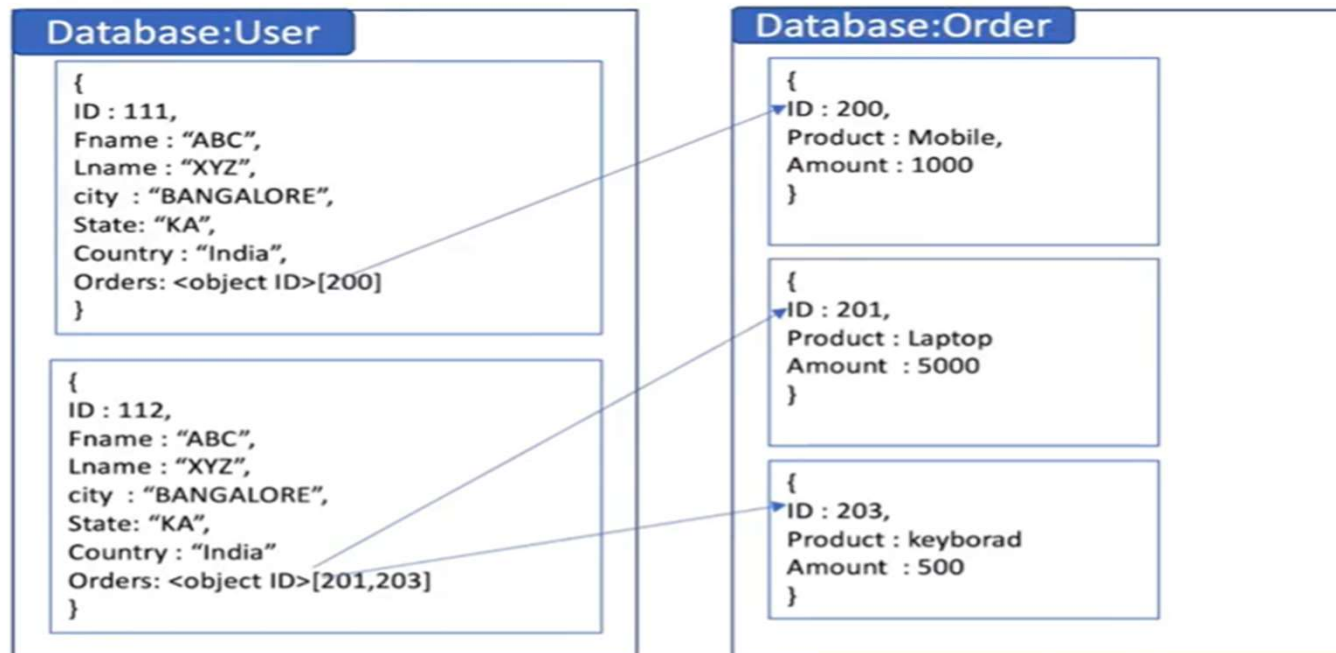
Popular usecases

- Caching
- Session management
- Leadboard

Example

- Redis
- DynamoDB
- Oracle NoSQL

Document DB



Popular usecases

- Blog/ website CMS
- Products catalog
- Big Data
- Analytics

Example

- MongoDB
- Apache CouchDB
- Azure Cosmos DB

MongoDB JSON

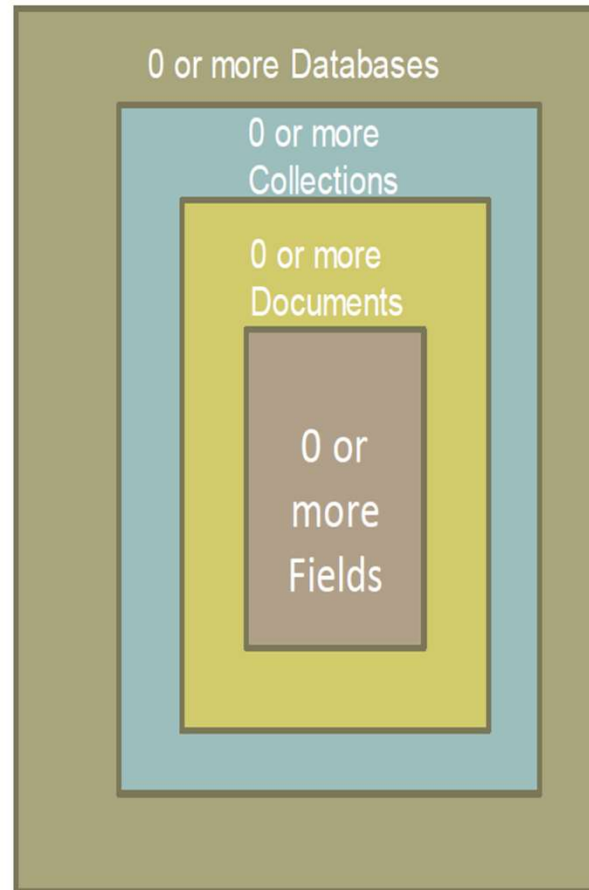
```
{
  "widget": {
    "debug": "on",
    "window": {
      "title": "Sample Konfabulator Widget",
      "name": "main_window",
      "width": 500,
      "height": 500
    },
    "image": {
      "src": "Images/Sun.png",
      "name": "sun1",
      "hOffset": 250,
      "vOffset": 250,
      "alignment": "center"
    },
    "text": {
      "data": "Click Here",
      "size": 36,
      "style": "bold",
      "name": "text1",
      "hOffset": 250,
      "vOffset": 100,
      "alignment": "center",
      "onMouseUp": "sun1.opacity = (sun1.opacity / 100) * 90;"
    }
  }
}
```

Structure



Mapping with RDBMS

- A MongoDB instance may have zero or more 'databases'
- A database may have zero or more 'collections'.
- A collection may have zero or more 'documents'.
- A document may have one or more 'fields'.
- MongoDB 'Indexes' function much like their RDBMS counterparts.



RDBMS		MongoDB
Database	⇒	Database
Table, View	⇒	Collection
Row	⇒	Document (BSON)
Column	⇒	Field
Index	⇒	Index
Join	⇒	Embedded Document
Foreign Key	⇒	Reference
Partition	⇒	Shard




Replication and Partitioning



Replication

Replication (Copying data)— Keeping a copy of same data on multiple servers that are connected via a network.



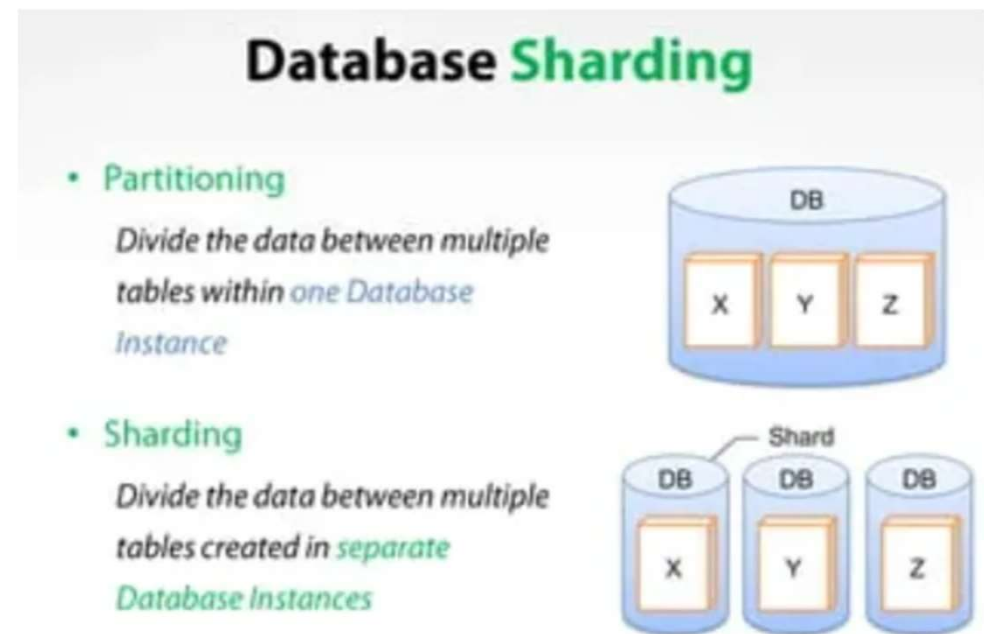
Partitioning

Partitioning — Splitting up a large monolithic database into multiple smaller databases based on data cohesion. e.g. Horizontal (sharding) and Vertical (increase server size) partitioning.

Partitioning and Sharding

All orders placed in January can be stored in one partition, all orders placed in February can be stored in another partition, and so on. Each partition can then be stored on a separate server. This way, when the company needs to retrieve orders from a particular time period, it can query only the relevant partition, which can significantly improve performance.

All orders from customers in North America can be stored on one server, all orders from customers in Europe can be stored on another server, and so on. This way, when the company needs to retrieve orders from a particular region, it can query only the relevant server, which can help improve performance and reduce the load on any one server.



Wide Column DB

Table:User		
ID	111	
Super Column : Name		
Column : FName	ABC	
Column : LName	XYZ	
Super Column : Address		
Column : city	BANGALORE	
Column : state	KA	
Column : country	India	
Super Column : Orders		
Orders	[200, 201]	

ID	111	
Super Column : Name		
Column : FName	ABC	
Column : LName	XYZ	

Table:Order		
ID	200	
Super Column : PRODUCT		
Column : Product	Laptop	
Super Column : Price		
Column : Amount	5000	

ID	201	
Super Column : PRODUCT		
Column : Product	Keyboard	
Super Column : Price		
Column : Amount	500	

Popular usecases

- IOT
- Inventory management
- Big Data

Table:Order

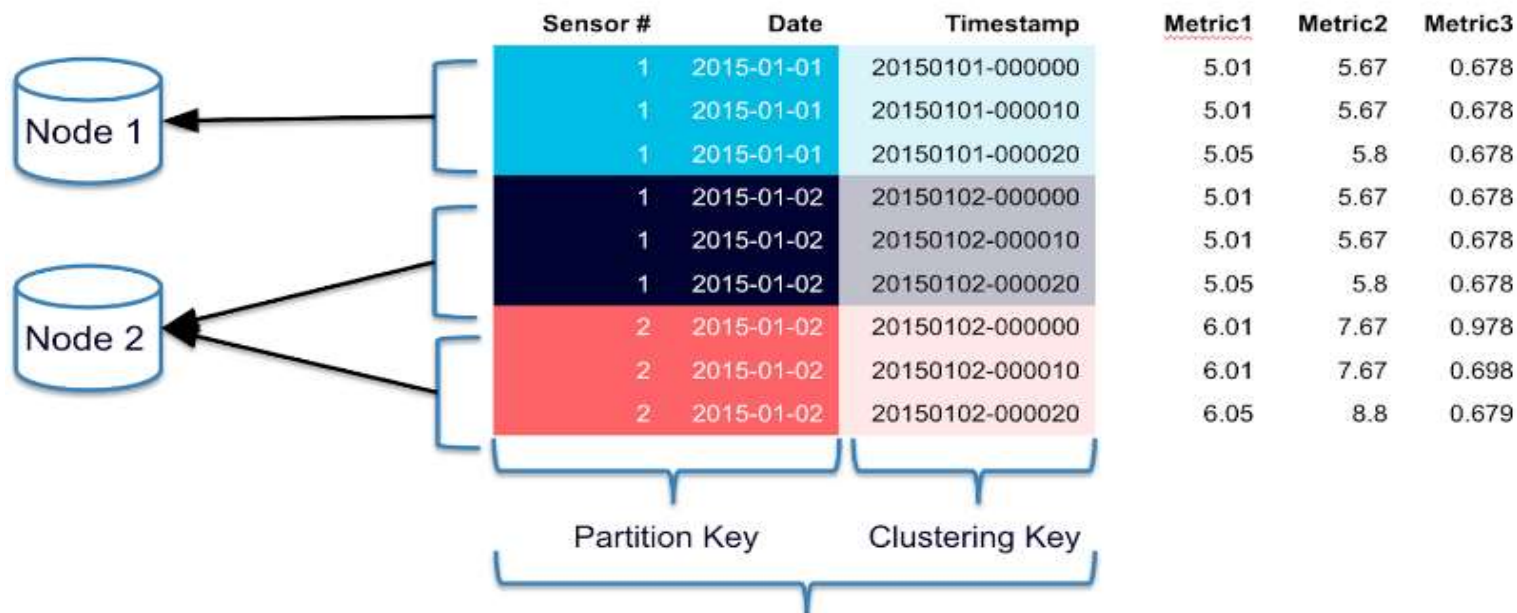
ID	200	
Super Column : PRODUCT		
Column : Product	Laptop	
Super Column : Price		
Column : Amount	5000	

ID	201	
Super Column : PRODUCT		
Column : Product	Keyboard	
Super Column : Price		
Column : Amount	500	

Example

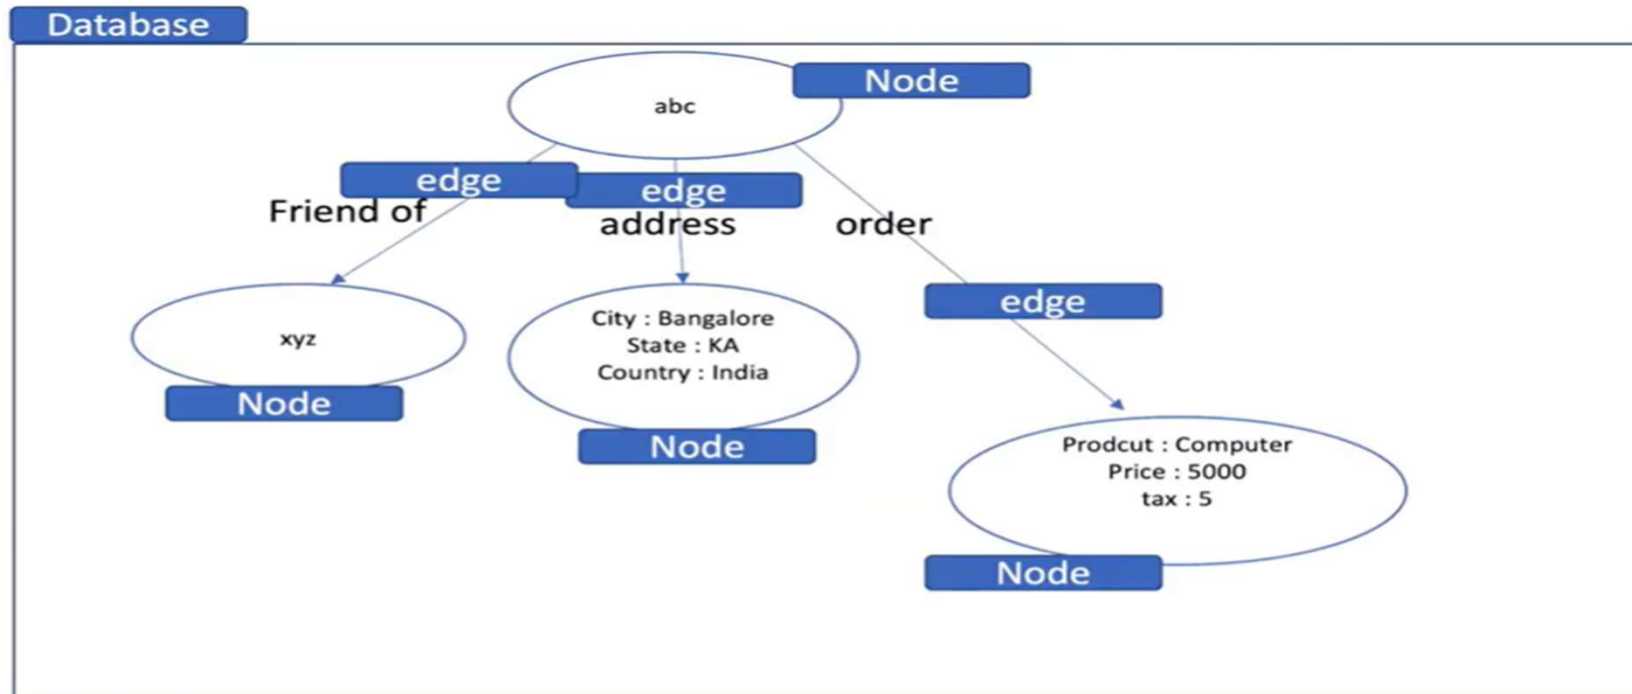
- Cassandra
- HBase

Cassandra Strategy



- Cassandra achieves a compromise between the two partitioning strategies .
- A table in Cassandra can be declared with a *compound primary key* consisting of several columns. Only the first part of that key is hashed to determine the partition, but the other columns are used as a concatenated index for sorting the data in Cassandra's SSTables.

Graph DB



Popular usecases

- Social networking
- IAM
- Recommendation

Example

- Neo4j
- JanusGraph



That is all