



**BITS Pilani**  
Pilani | Dubai | Goa | Hyderabad

# Containers

Aditya Goel

---

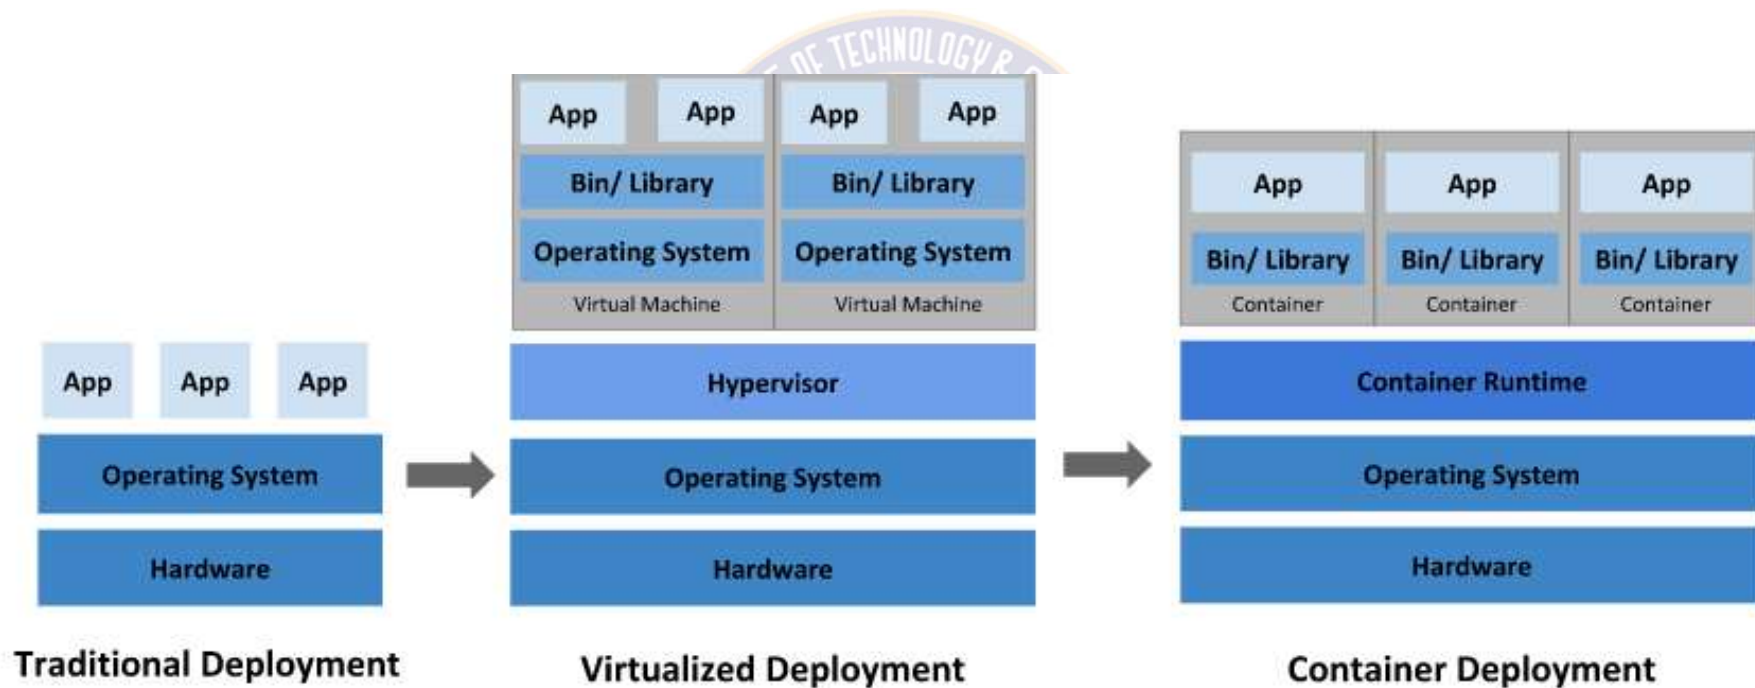
# Agenda



- **Containers**

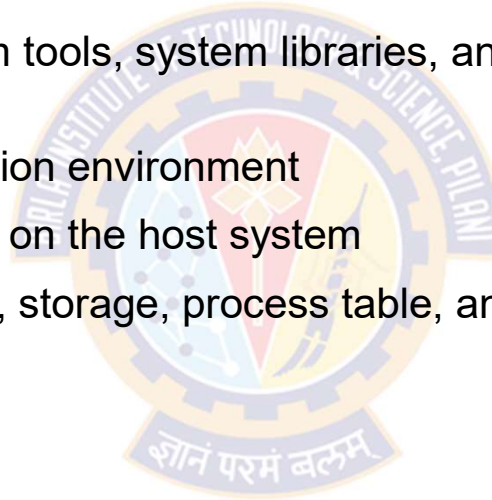
- ✓What are containers?
- ✓Containers – History
- ✓Cgroups
- ✓Namespaces
- ✓Virtual Machine vs Containers
- ✓Containers and Virtual Machines
- ✓Types of Containers
- ✓Dockers

# Going back in time



# What are Containers?

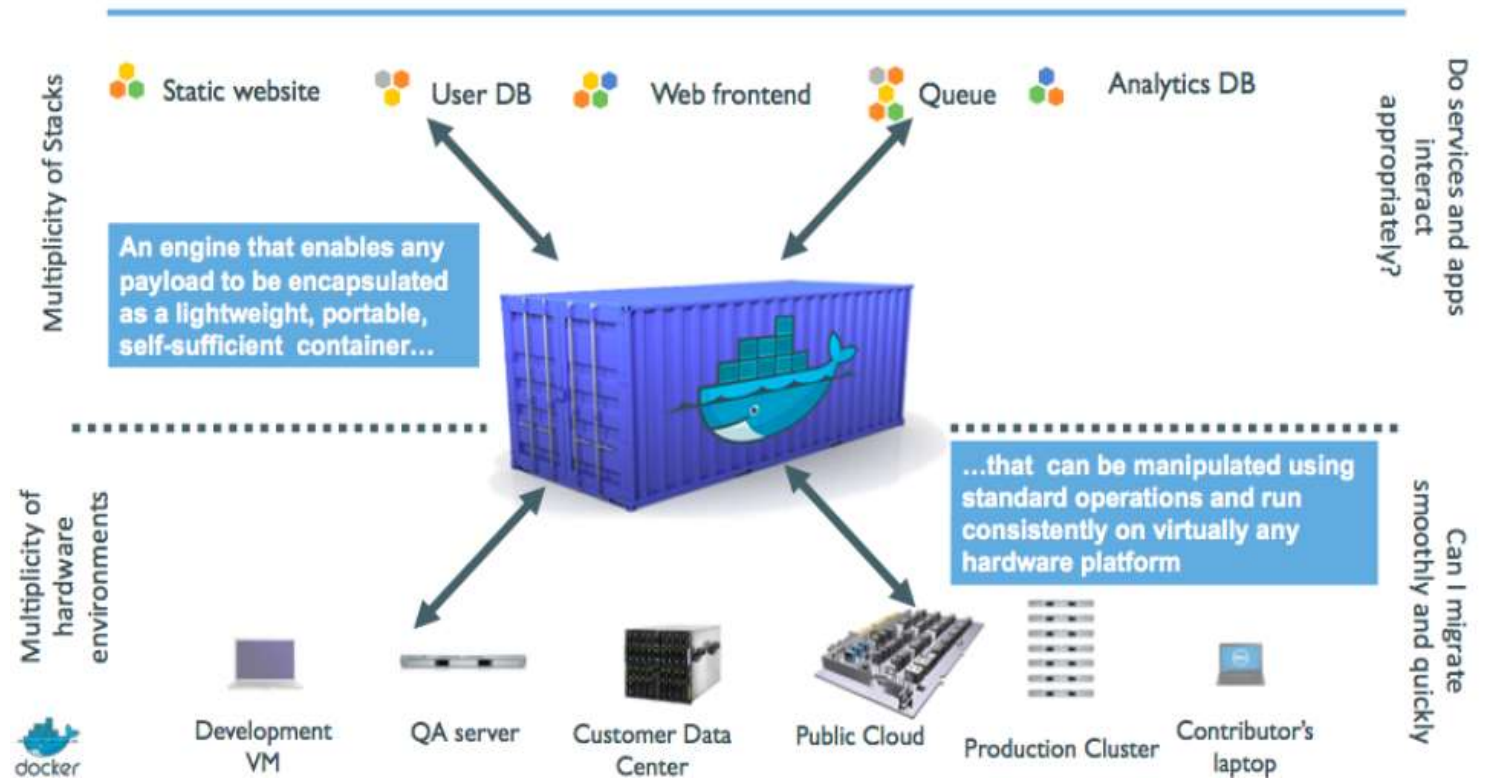
- A software container is a standardized package of software.
- Everything needed for the software to run is inside the container.
- The software code, runtime, system tools, system libraries, and settings are all inside a single container
- Container is a semi-isolated execution environment
- Managed by the OS kernel running on the host system
- Has its own isolated memory, CPU, storage, process table, and networking interfaces






# What are Containers?

## A shipping container system for applications



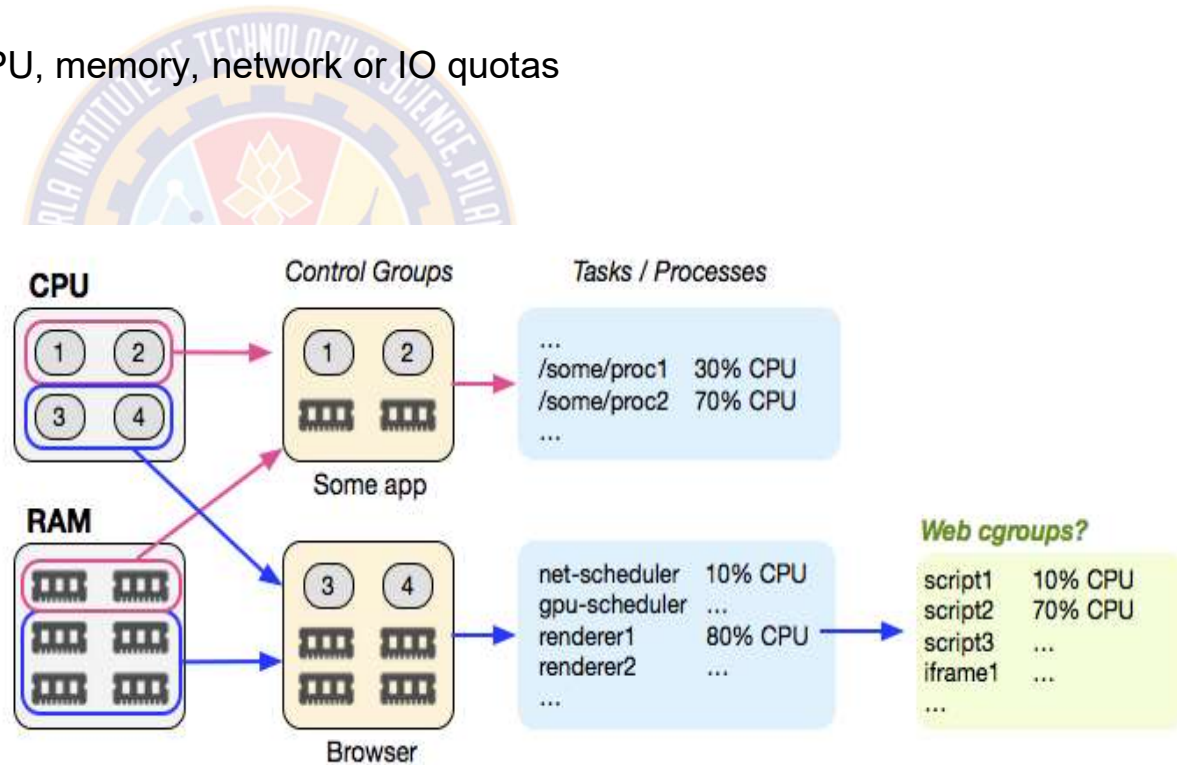
# Containers - History

- Containers are powered by two underlying Linux Kernel technologies:
  - cgroups
  - namespaces



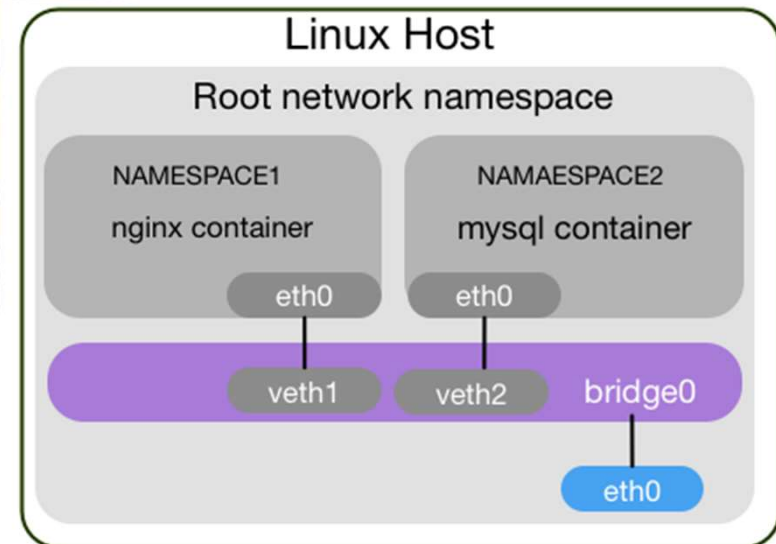
# cgroups

- cgroups – Control groups
  - A kernel mechanism for limiting and measuring the total resources used by a group of processes running on a system
  - Processes can be applied with CPU, memory, network or IO quotas
- Cgroup merged into Linux 2.6.24
- Cgroups provides:
  - **Resource limiting**
  - **Prioritization**
  - **Accounting**
  - **Control**



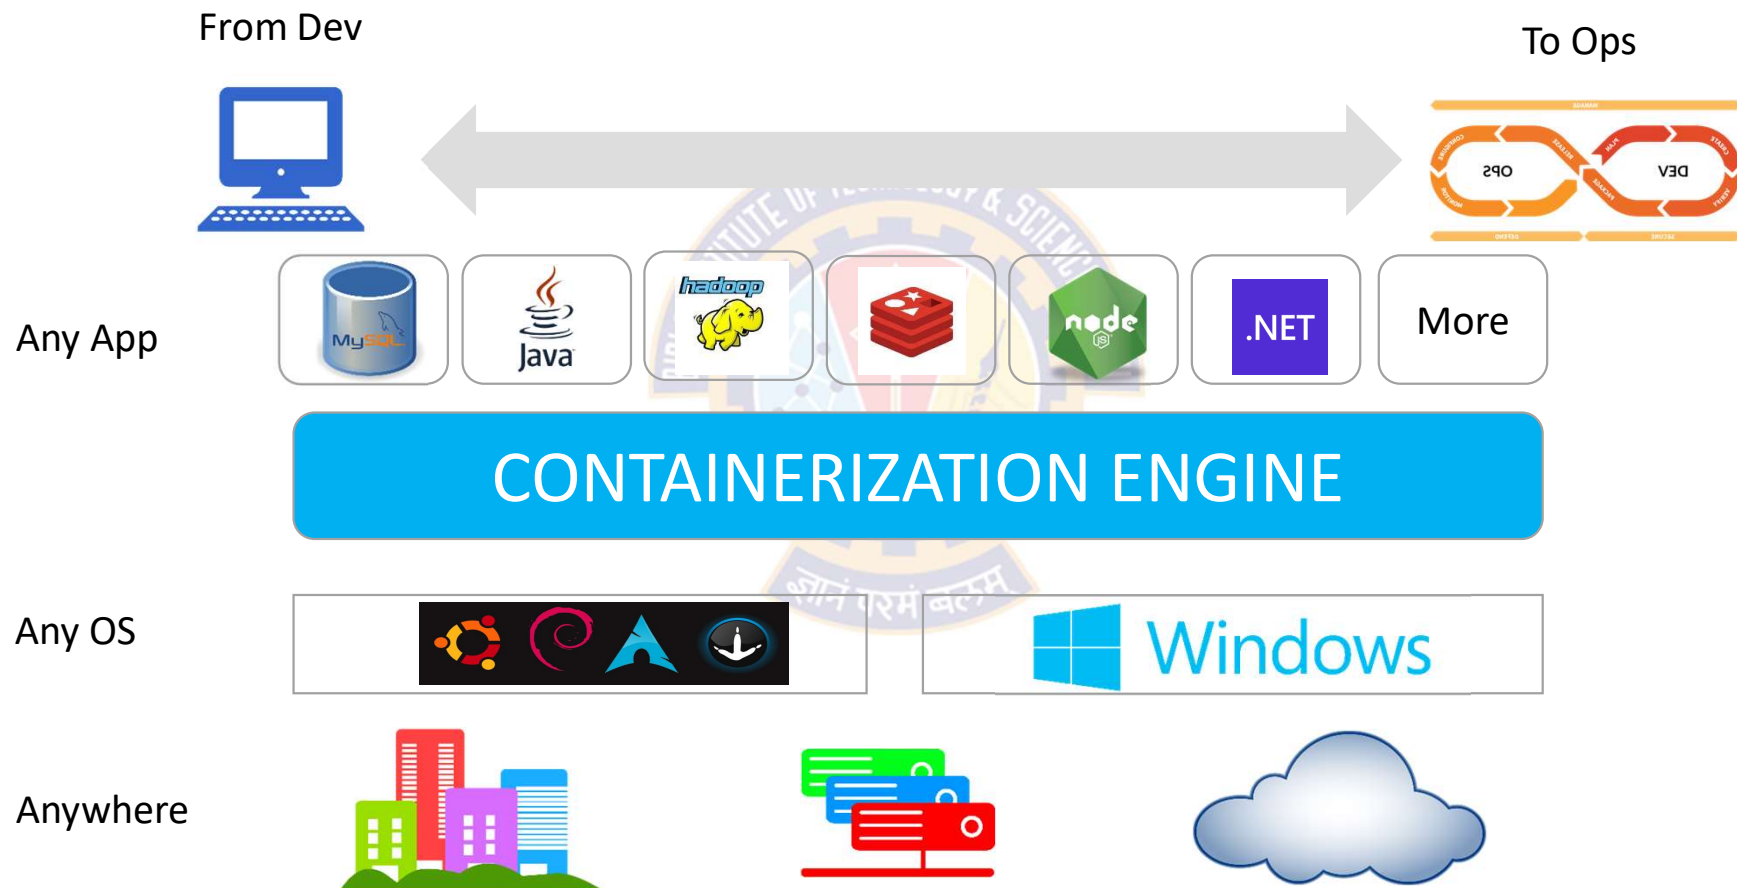
# Namespaces

- Namespaces - kernel mechanism for limiting the visibility that a group of processes has of the rest of a system
- Limit visibility
  - Certain process trees
  - Network interfaces
  - User IDs
  - Filesystem mounts
- Namespace merged into Linux 3.8

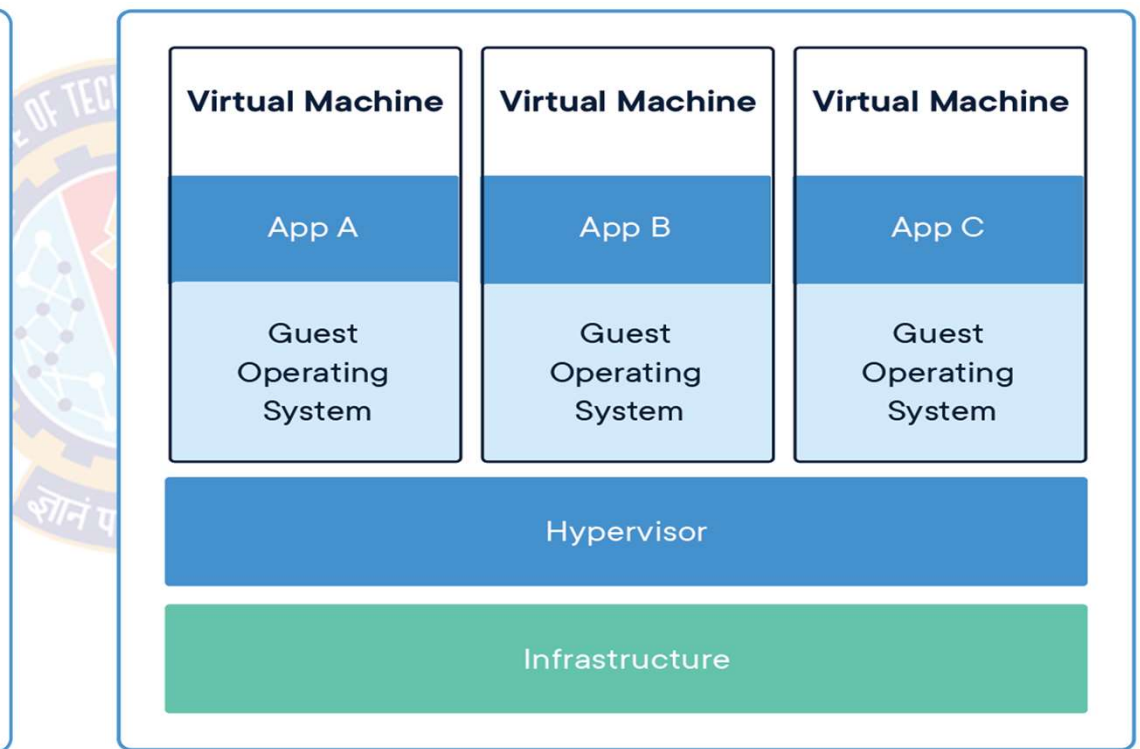
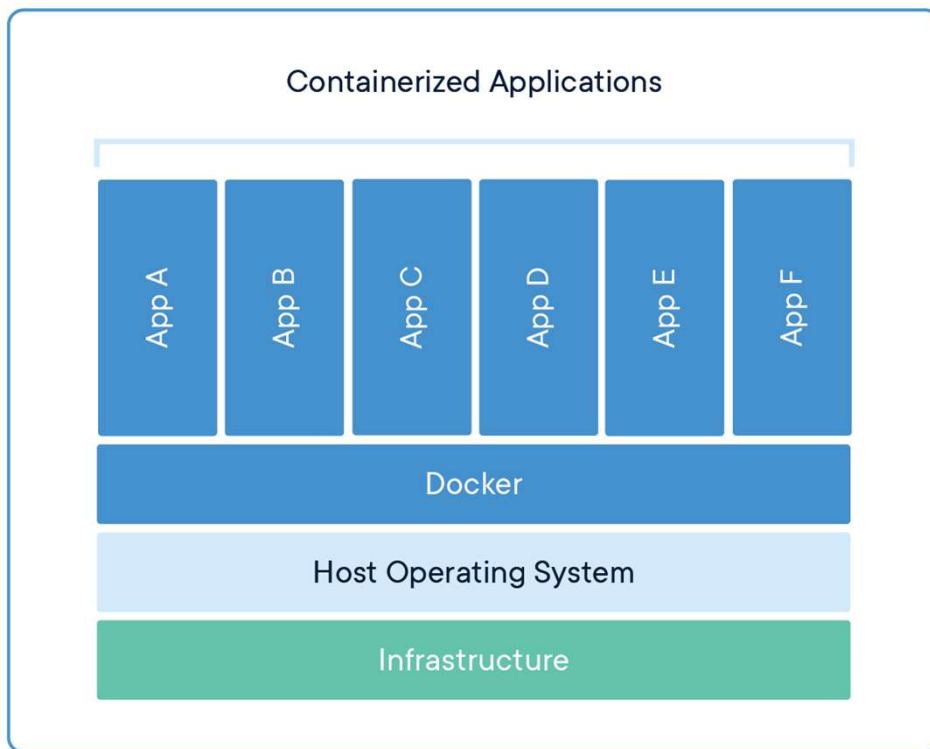




# Build, Ship, Run, Any App Anywhere



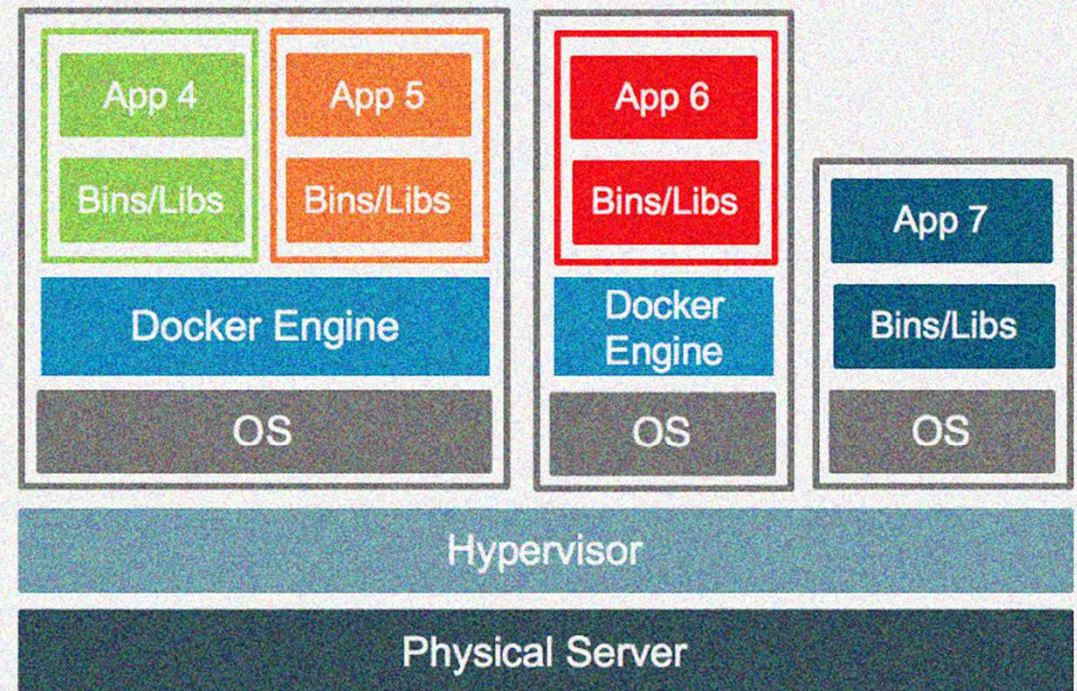
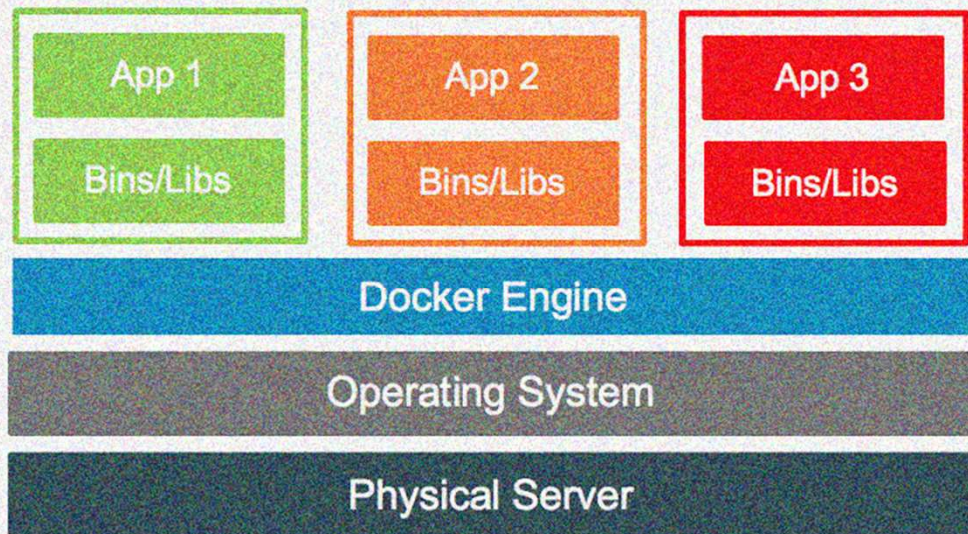
# Virtual Machine vs Containers





# Containers and Virtual Machines

Your Datacenter or VPC



# Types of Containers

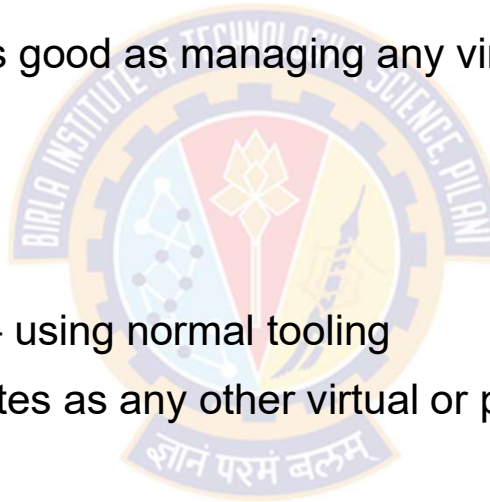
- System Containers
- Application Containers





# System Containers

- Similar to Virtual or physical machine
- They run a full operating system
- Managing a System Container is as good as managing any virtual or physical machine
  - install packages inside them
  - manage services
  - define backup policies
  - Monitoring
- These containers can be updated – using normal tooling
- They get system and security updates as any other virtual or physical machine
- Oldest type of container - 1999



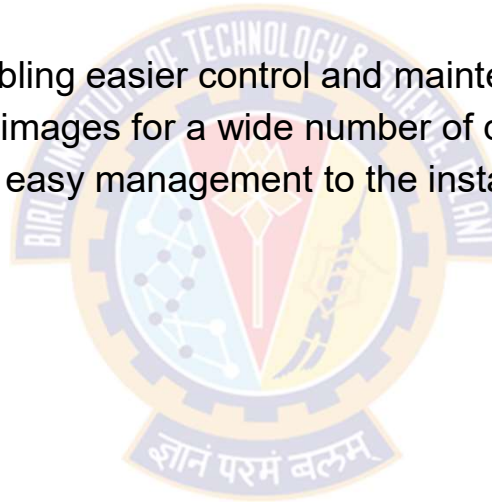


# System Containers - History

- BSD introduced jails, a way of running a second BSD system on the same kernel as the main system
- Linux followed the concept through Linux vServer
- Then Solaris – Zones
- Then OpenVZ project – implemented VPS(Virtual Private Servers) on Linux
- LXC – Linux Containers – mainline Linux implementation
- LXC is a low-level tool that can create both system containers and application containers
- Docker was initially based on LXC
- Goal of LXC: to create an environment as close as possible to a standard Linux installation but without the need for a separate kernel

# System Containers - History

- What is LXD?
  - Is a system container and a virtual machine manager
  - Runs on top of LXC
  - Enhances the experience and enabling easier control and maintenance.
  - LXD is image-based and provides images for a wide number of different Linux distributions.
  - A command-line tool that enables easy management to the instances (Containers and VMs)



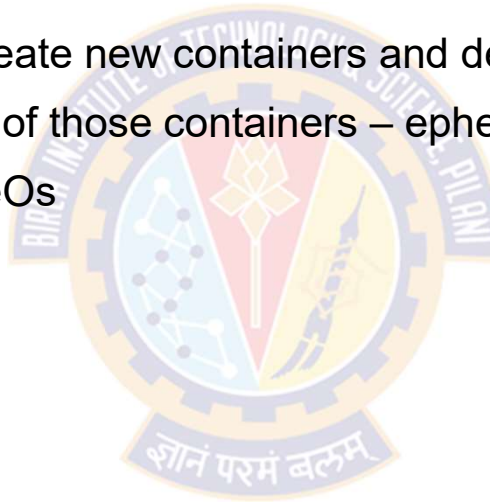
# System Containers - History

- LXD vs LXC


LXC	LXD
Linux container runtime allowing creation of multiple isolated Linux systems (containers) on a control host using a single Linux kernel	System container and virtual machine <b>manager</b> built on top of LXC, enabling easier management, control and integration
Only supports containers	Supports container and VMs
Low-level tool requiring expertise	Better user experience through a simple REST API
Online Demo Tool: <a href="https://linuxcontainers.org/lxd/try-it/">https://linuxcontainers.org/lxd/try-it/</a>	

# Application Containers

- Containers running a single process per container
- They run stateless types of workloads
- Scale up and down as needed – create new containers and delete them at any time
- No need to care about the lifecycle of those containers – ephemeral
- Example: Docker and rkt from CoreOs




# System vs Application Containers



OS containers

- Meant to be used as an OS - run multiple services
- No layered filesystems by default
- Built on cgroups, namespaces, native process resource isolation
- Examples - LXC, OpenVZ, Linux VServer, BSD Jails, Solaris Zones



App containers

- Meant to run for a single service
- Layered filesystems
- Built on top of OS container technologies
- Examples - Docker, Rocket

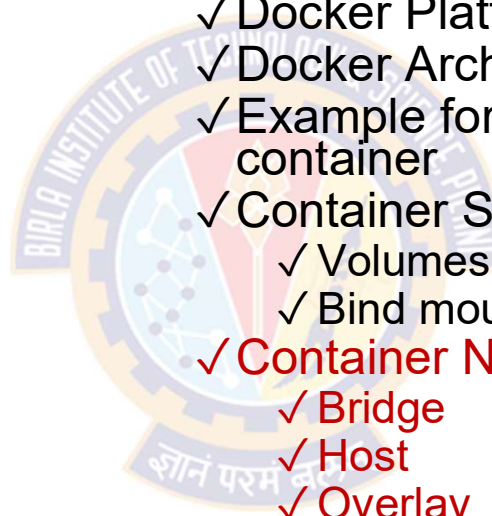


# Agenda



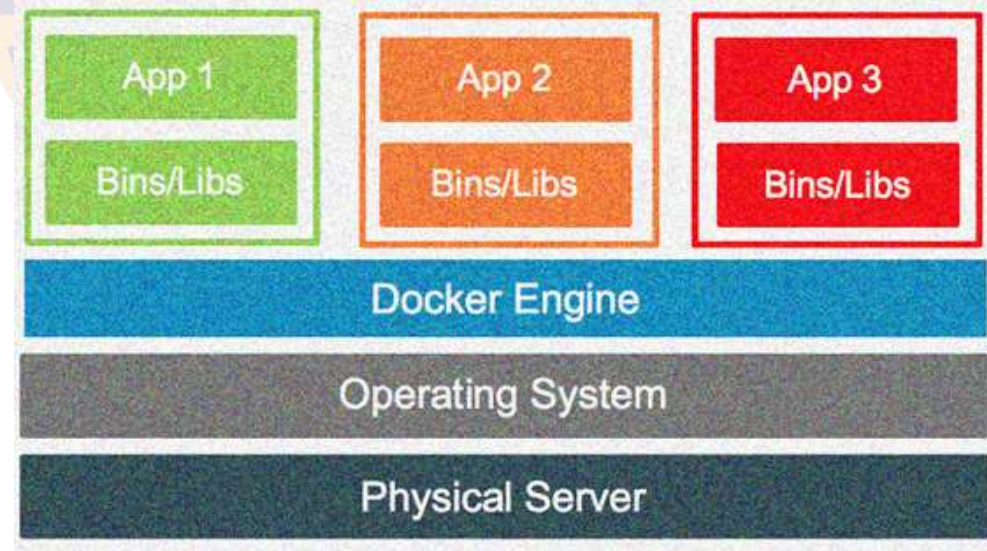
- **Docker Overview**

- ✓ Docker Platform
- ✓ Docker Architecture
- ✓ Example for running a Docker container
- ✓ Container Storage
  - ✓ Volumes
  - ✓ Bind mounts
- ✓ **Container Networking**
  - ✓ Bridge
  - ✓ Host
  - ✓ Overlay
  - ✓ none



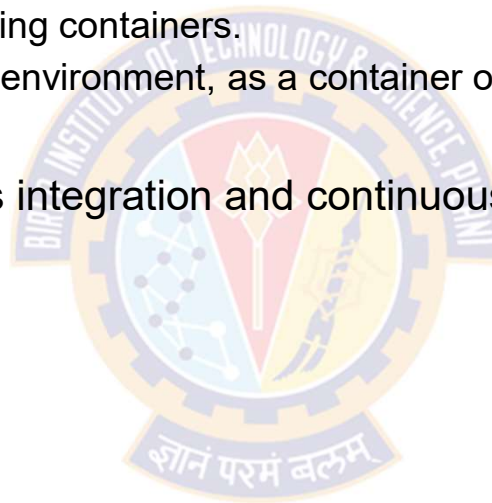
# The Docker platform

- Docker is an open platform
- Docker separates applications from hardware infrastructure
- **Containers are used to package and run an application**
- A single host can run many containers simultaneously
- Containers are lightweight and contain everything needed to run the application



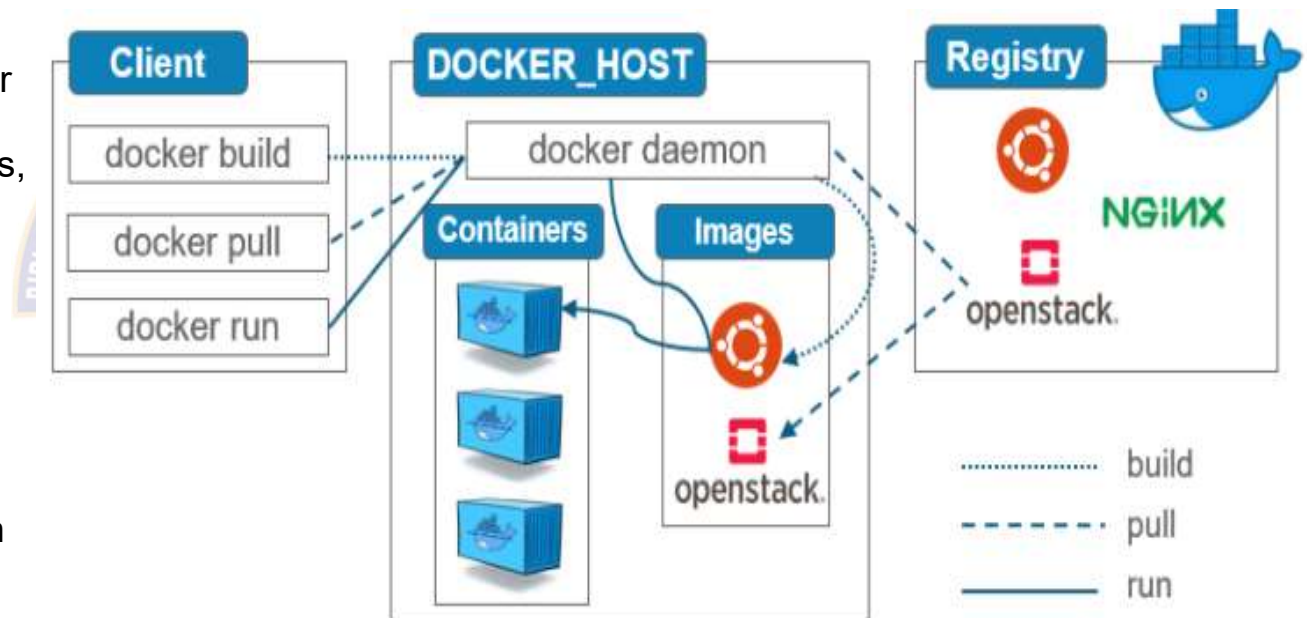
# The Docker platform

- Docker provides tooling and a platform to manage the lifecycle of your containers:
  - Develop application using containers.
  - Distributing and test application using containers.
  - Deploy application into production environment, as a container or an orchestrated service.
- Containers are great for continuous integration and continuous delivery (CI/CD) workflows.



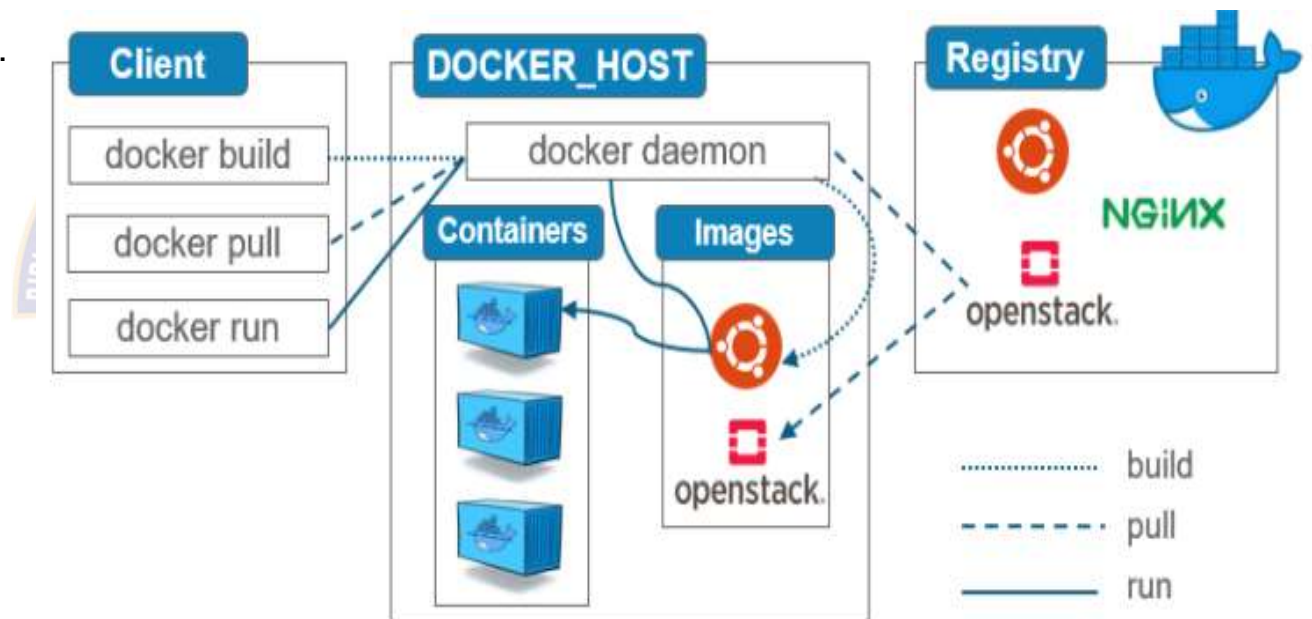
# Docker architecture

- Docker uses a client-server architecture.
- The Docker daemon
  - The Docker daemon (dockerd) listens for Docker API requests
  - Manages Docker objects such as images, containers, networks, and volumes.
  - Builds, runs, and distributes containers
- The Docker client
  - The Docker *client* talks to the Docker *daemon*
  - The Docker client and daemon *can* run on the same system
  - The Docker client can communicate with more than one daemon..
- The Docker client and daemon communicate using a REST API, over UNIX sockets or a network interface.



# Docker architecture

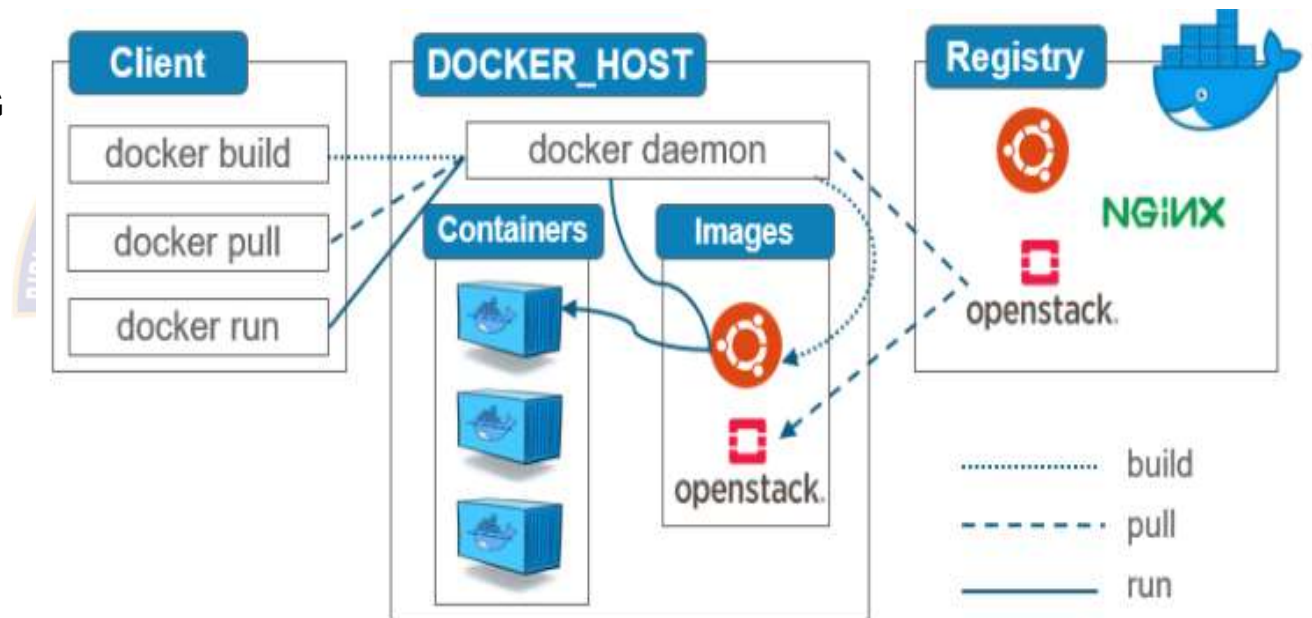
- Docker registries
  - A Docker *registry* stores Docker images.
  - Docker Hub is a public registry that anyone can use.
  - Docker is configured to look for images on Docker Hub by default



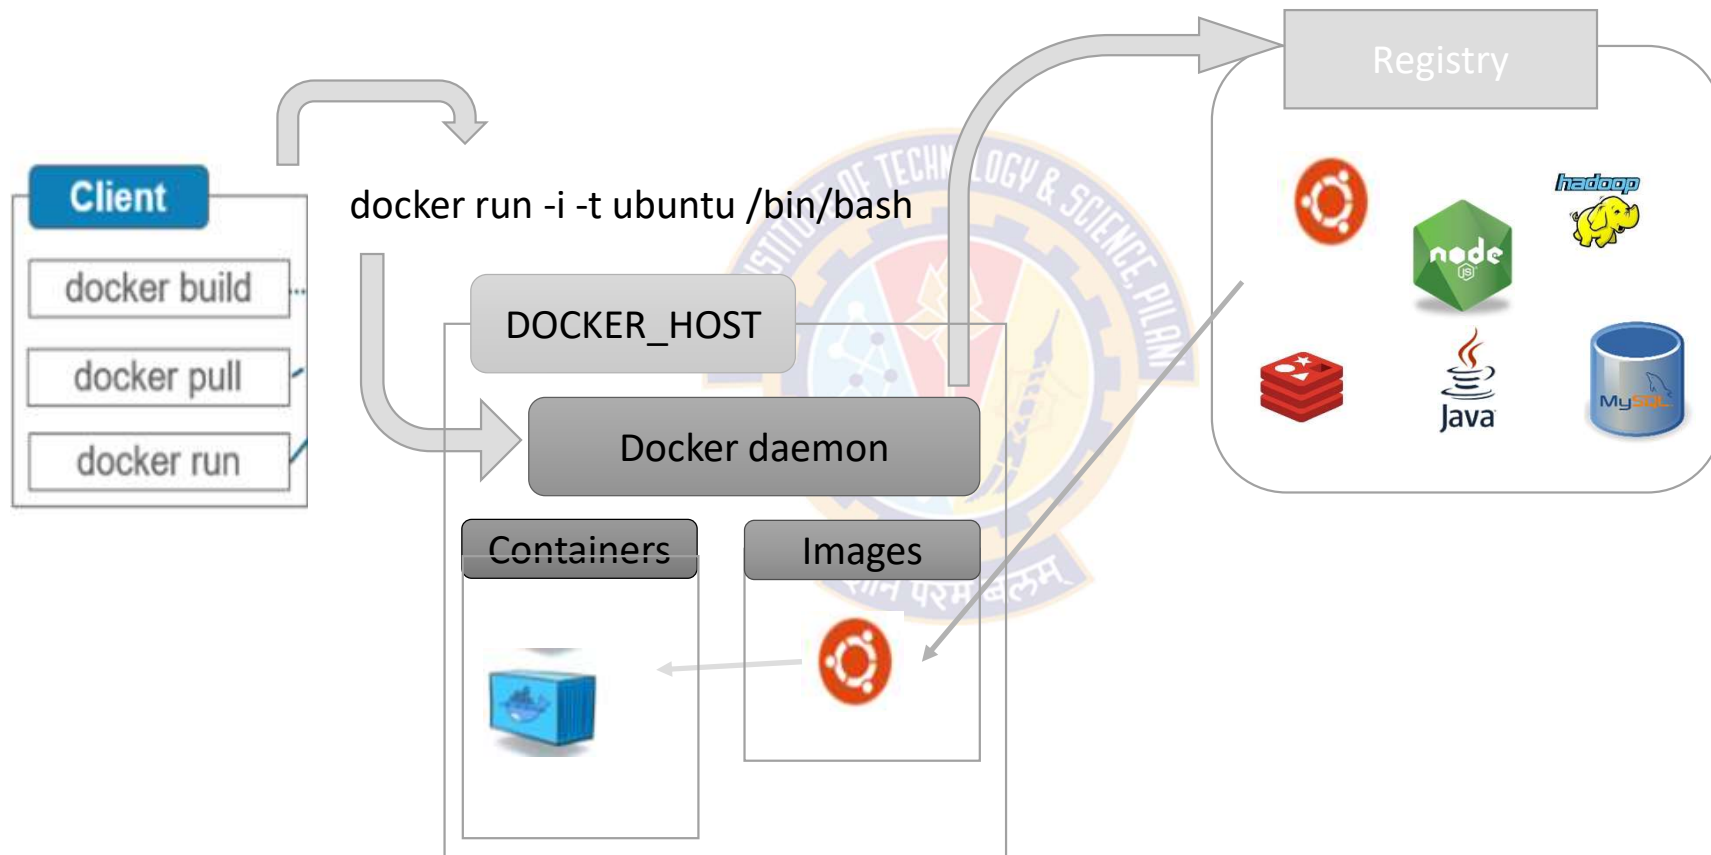


# Docker architecture

- Docker objects
  - IMAGES: AN *IMAGE* IS A READ-ONLY TEMPLATE WITH INSTRUCTIONS FOR CREATING A DOCKER CONTAINER.
  - CONTAINERS: A CONTAINER IS A RUNNABLE INSTANCE OF AN IMAGE



# Example for running a Docker container



# Docker Commands

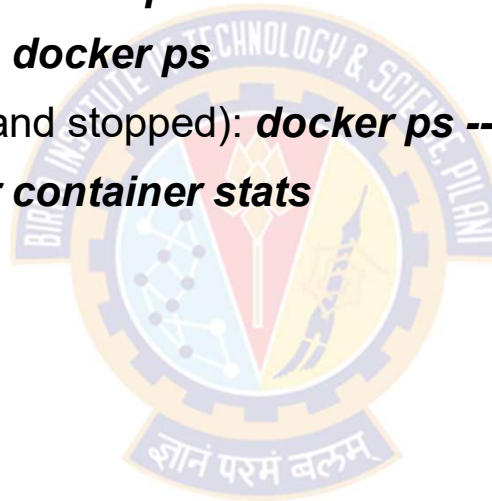
## Containers

- A container is a runtime instance of a docker image
- Create and run a container from an image, with a custom name:
  - `docker run -name <container name> <image name>`  
***docker run -name mylinuxserver Ubuntu***
- Run a container with and publish a container's port(s) to the host.
  - `docker run -p <host port>:<container port> <image name>`  
***docker run -p 8080:80 nginx***
- Run a container in the background
  - `docker run -d <image name>`  
***docker run -d -p 8080:80 nginx***
- Start or stop an existing container:
  - `docker start/stop <container name> (or <container id>)`  
***docker stop 11ed or mynginx***
- Remove a stopped container:
  - `docker rm <container name> (or <container id>)`  
***docker rm -f 11ed or mynginx***
- Open a shell inside a running container:
  - `docker exec -it <container name> sh`  
***docker exec -it myubuntu bash***

# Docker Commands

## Container

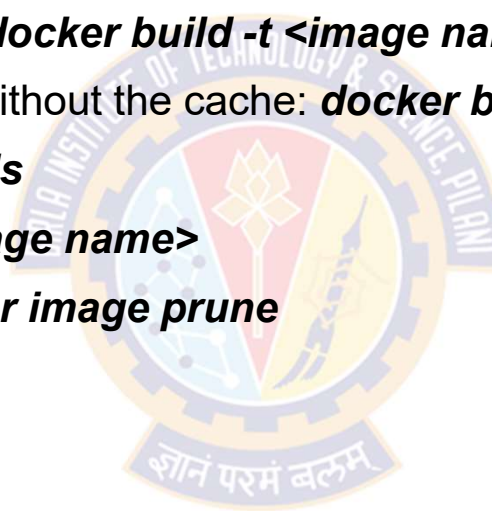
- Fetch and follow the logs of a container: ***docker logs -f <container name>***
- To inspect a running container: ***docker inspect <container id> (or ) <container name>***
- To list currently running containers: ***docker ps***
- List all docker containers (running and stopped): ***docker ps --all***
- View resource usage stats: ***docker container stats***



# Docker Commands

**Images:** Docker images are a lightweight, standalone, executable package of software that includes everything needed to run an application: code, runtime, system tools, system libraries and settings

- Build an Image from a Dockerfile: ***docker build -t <image name>***
- Build an Image from a Dockerfile without the cache: ***docker build -t <image name>. --no-cache***
- List local images: ***docker images ls***
- Delete an Image: ***docker rmi <image name>***
- Remove all unused images: ***docker image prune***





# Build and run customised Docker image

## Dockerfile

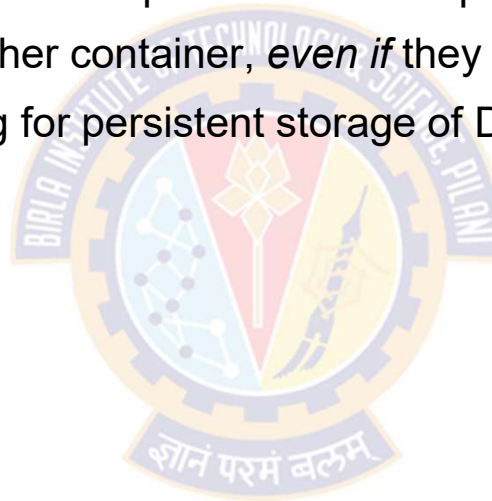
```
FROM ubuntu:latest  
RUN mkdir /app  
RUN apt update  
RUN apt install vim gcc -y  
WORKDIR /app  
ENTRYPOINT ["/bin/bash"]
```

- docker build -t mycustomimage .
- docker run -it mycustomimage



# Container Storage

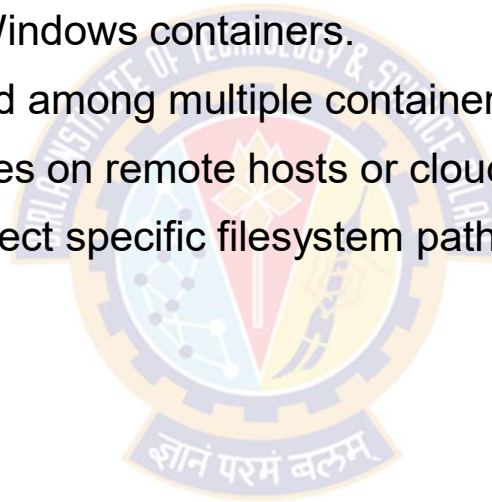
- The container's filesystem
- Each container also gets its own “scratch space” to create/update/remove files.
- Any changes won't be seen in another container, *even if* they are using the same image
- Docker containers use the following for persistent storage of Data
  - Volumes
  - Bind mounts



# Container Storage

## Volumes

- Docker containers use **Volumes** as the *preferred mechanism* to *store persisting data* generated by and used by containers
- Volumes work on both Linux and Windows containers.
- Volumes can be more safely shared among multiple containers.
- Volume drivers let you store volumes on remote hosts or cloud providers
- Volumes provide the ability to connect specific filesystem paths of the container back to the host machine



# Container Storage: Volumes

## Steps to create and mount Volume:

- Create a volume by using the docker volume create command

**`docker volume create mydb`**

- Start the container with mount

**`docker run -dp 3000:3000 --mount type=volume,src=mydb,target=/etc/myappdb getting-started`**

- `docker volume inspect`

**`docker volume inspect mydb`**

**OUTPUT**

```
[
  {
    "CreatedAt": "2019-09-26T02:18:36Z",
    "Driver": "local",
    "Labels": {},
    "Mountpoint": "/var/lib/docker/volumes/mydb/_data",
    "Name": "mydb",
    "Options": {},
    "Scope": "local"
  }
]
```



# Container Storage: bind mounts

- A bind mount - share a directory from the host's filesystem into the container.
- When working on an application, you can use a bind mount to mount source code into the container.
- The container sees the changes you make to the code immediately, as soon as you save a file.
- This means that you can run processes in the container that watch for filesystem changes and respond to them.

***docker run -it --mount type=bind,src="\$(pwd)",target=/src ubuntu bash***

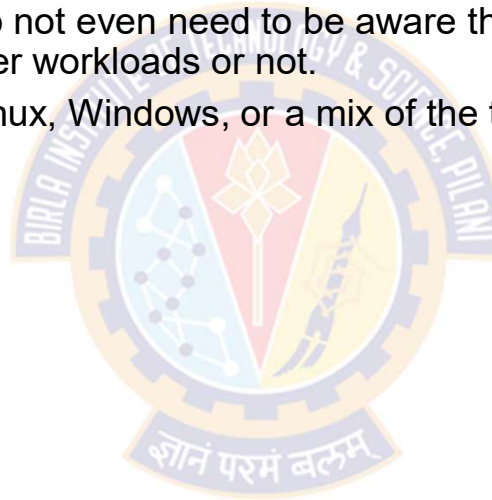
The ***--mount*** option tells Docker to create a bind mount

***src is the current working directory on your host machine (getting-started/app)***

***target is where that directory should appear inside the container (/src)***

# Container Networking


- Docker containers and services are powerful
  - Connect them together, or connect them to non-Docker workloads.
  - Docker containers and services do not even need to be aware that they are deployed on Docker, or whether their peers are also Docker workloads or not.
  - Whether your Docker hosts run Linux, Windows, or a mix of the two, you can use Docker to manage them in a platform-agnostic way.





# Network drivers

- Docker's networking subsystem is pluggable, using drivers.
- Several drivers exist by default, and provide core networking functionality:
  - Bridge
  - Host
  - Overlay
  - Ipvlan
  - Macvlan
  - none



# Network drivers: bridge

- The default network driver.
- If you don't specify a driver, this is the type of network you are creating.
- **Bridge networks are usually used when your applications run in standalone containers that need to communicate.**

`docker network ls`

#command to list all the networks(drivers) in Docker

`docker run -dit --rm --name alpine1 alpine ash`

#command to run a container

`docker run -dit --rm --name alpine2 alpine ash`

`docker container ls`

#list all the containers

`docker network inspect bridge`

#details of bridge network driver

`docker attach alpine1`

#attach to the container (get console)

`$ip addr show`

`$ping -c 2 google.com`

`$ping -c 2 alpine2`

`docker stop alpine1 alpine2`

#to stop the containers

# Network drivers: bridge

- User Defined bridge network

```
docker network create --driver bridge mynet
```

```
docker network ls
```

```
docker network inspect mynet
```

---

```
docker run -dit --rm --name alpine1 --network mynet alpine ash
```

```
docker run -dit --rm --name alpine2 --network mynet alpine ash
```

```
docker run -dit --rm --name alpine3 alpine ash
```

```
docker run -dit --rm --name alpine4 --network mynet alpine ash
```

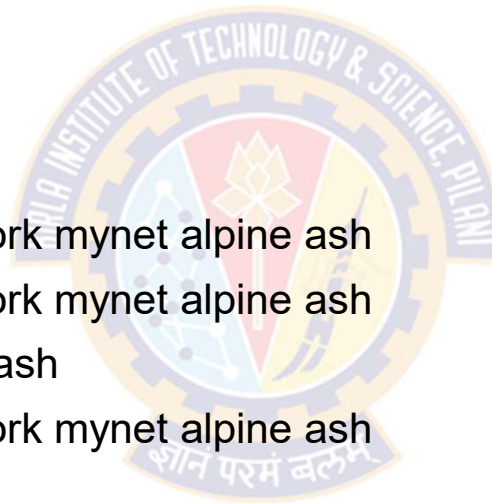
```
docker network connect bridge alpine4
```

---

```
docker network inspect bridge
```

```
docker network inspect mynet
```

---



## Network drivers: bridge

---

docker container attach alpine1

ping -c 2 alpine2

ping -c 2 alpine3 (not pingable)

---

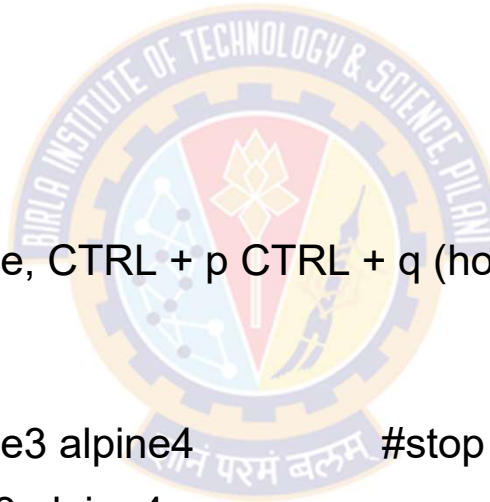
Detach from alpine1 using detach sequence, CTRL + p CTRL + q (hold down CTRL and type p followed by q).

ping -c 2 google.com

docker container stop alpine1 alpine2 alpine3 alpine4 #stop and remove containers

docker container rm alpine1 alpine2 alpine3 alpine4

docker network rm mynet #remove user-defined bridge network



# Network drivers: Host

## Networking using the host network

- The goal of this tutorial is to start a container(nginx) which binds directly to port 80 on the Docker host.
- From a networking point of view, this is the same level of isolation as if the nginx process were running directly on the Docker host and not in a container.
- However, in all other ways, such as storage, process namespace, and user namespace, the nginx process is isolated from the host.
- This procedure requires port 80 to be available on the Docker host.
- The host networking driver only works on Linux hosts, and is not supported on Docker Desktop for Mac, Docker Desktop for Windows, or Docker EE for Windows Server.



# Network drivers: Host

## Networking using the host network

- The goal of this tutorial is to start a container(nginx) which binds directly to port 80 on the Docker host.
- From a networking point of view, this is the same level of isolation as if the nginx process were running directly on the Docker host and not in a container.
- However, in all other ways, such as storage, process namespace, and user namespace, the nginx process is isolated from the host.
- This procedure requires port 80 to be available on the Docker host.
- **The host networking driver only works on Linux hosts**, and is not supported on Docker Desktop for Mac, Docker Desktop for Windows, or Docker EE for Windows Server.

```
docker run --rm -d --network host --name my_nginx nginx
```

```
sudo netstat -tulpn | grep :80
```

#Verify which process is bound to port 80

```
docker container stop my_nginx
```

# Network drivers: Overlay

- Overlay networks connect multiple Docker daemons running in multiple hosts together and enable swarm services to communicate with each other.
- You can also use overlay networks to facilitate communication between a swarm service and a standalone container, or between two standalone containers on different Docker daemons.
- This strategy removes the need to do OS-level routing between these containers.



# Network drivers: Overlay

## Use the default overlay network

- Create the swarm

- On manager. initialize the swarm.

```
docker swarm init
```

```
docker swarm join --token <TOKEN> \
```

```
  --advertise-addr <IP-ADDRESS-OF-WORKER-1> | #optional  
  <IP-ADDRESS-OF-MANAGER>:2377
```

```
docker node ls
```

```
#on manager
```

```
docker network ls
```

The `docker_gwbridge` connects the ingress network to the Docker host's network interface so that traffic can flow to and from swarm managers and workers.

# Network drivers: Overlay

## Use the default overlay network

```
sudo docker network create -d overlay nginx-net
```

```
docker service create \
```

```
--name my-nginx \
```

```
--publish target=80,published=80 \
```

```
--replicas=5 \
```

```
--network nginx-net \
```

```
nginx
```

```
sudo docker service ls
```

```
sudo docker service ps my-nginx
```

```
sudo docker service rm my-nginx
```



# Network drivers: Overlay

## Use an overlay network for standalone containers

- Set up the swarm

```
docker swarm init
```

```
docker swarm join --token <your_token> <your_ip_address>:2377
```

On host1, create an attachable overlay network called test-net:

```
docker network create --driver=overlay --attachable test-net
```

```
docker run -it --name alpine1 --network test-net alpine
```

```
docker network ls          #on host2
```

```
docker run -dit --name alpine2 --network test-net alpine
```

```
docker network ls
```

```
ping -c 2 alpine2          #from host1
```

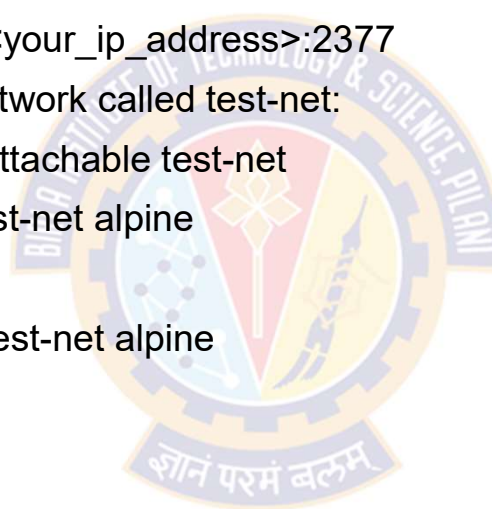
```
docker container stop alpine2
```

```
docker network ls
```

```
docker container rm alpine2
```

```
docker container rm alpine1
```

```
docker network rm test-net
```



# Network drivers

- Ipvlan
  - Ipvlan networks give users total control over both IPv4 and IPv6 addressing.
  - The VLAN driver builds on top of that in giving operators complete control of layer 2 VLAN tagging and even Ipvlan L3 routing for users interested in underlay network integration.
- Macvlan
  - Macvlan networks allow you to assign a MAC address to a container, making it appear as a physical device on your network.
  - The Docker daemon routes traffic to containers by their MAC addresses
- None
  - For this container, disable all networking.
  - Usually used in conjunction with a custom network driver.