



**BITS Pilani**  
Pilani Campus

# Popular scaling approaches

Akanksha Bharadwaj  
CSIS Department



**BITS Pilani**  
Pilani Campus



# **SE ZG583, Scalable Services**

## **Lecture No. 2**



**BITS Pilani**  
Pilani Campus



# Partitioning and Sharding

# Introduction



- In many large-scale solutions, data is divided into *partitions* that can be managed and accessed separately.

## Why partition data?

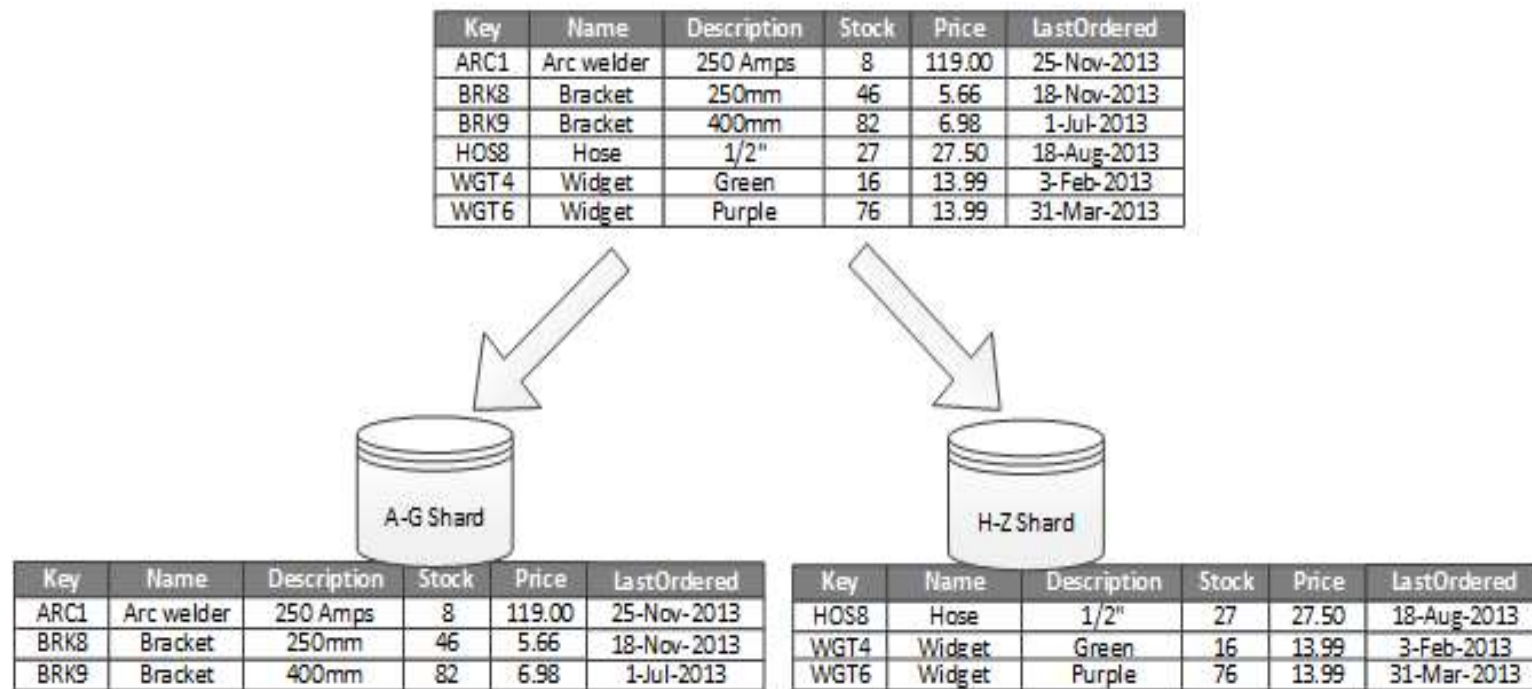
- Improve scalability
- Improve performance
- Improve security
- Provide operational flexibility
- Improve availability

# Types of Partitioning

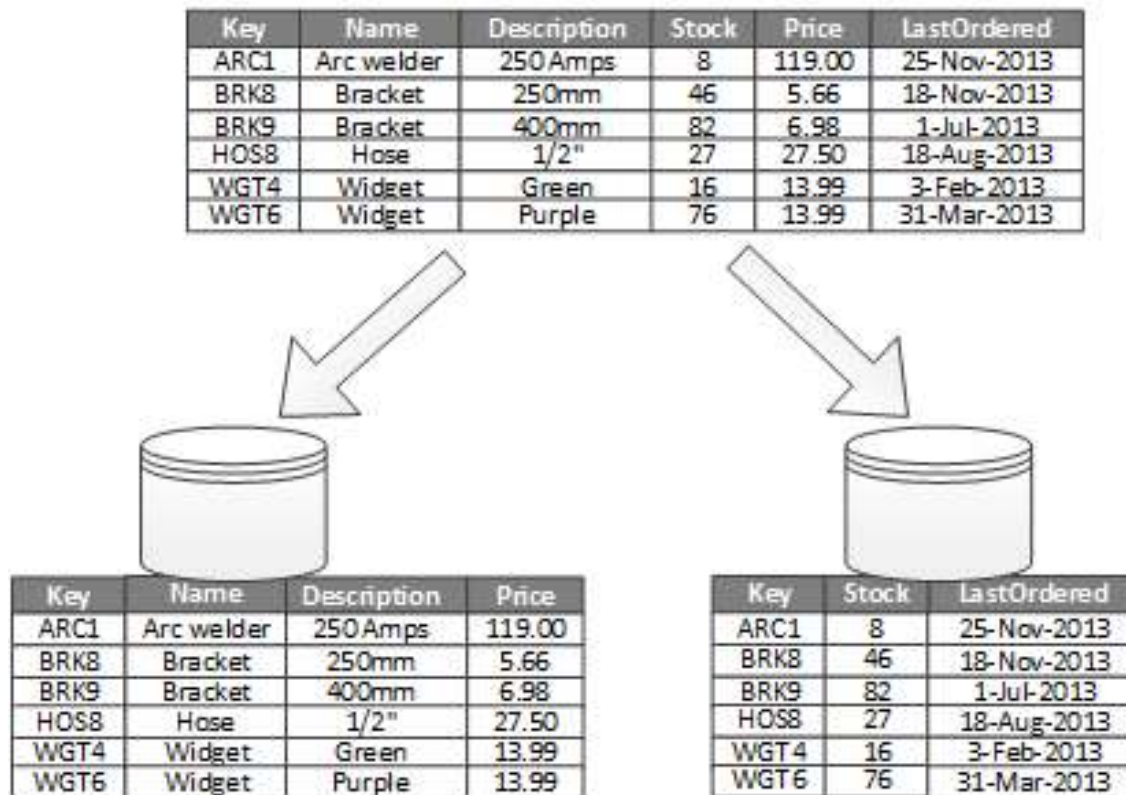


- Horizontal partitioning
- Vertical partitioning
- Functional partitioning

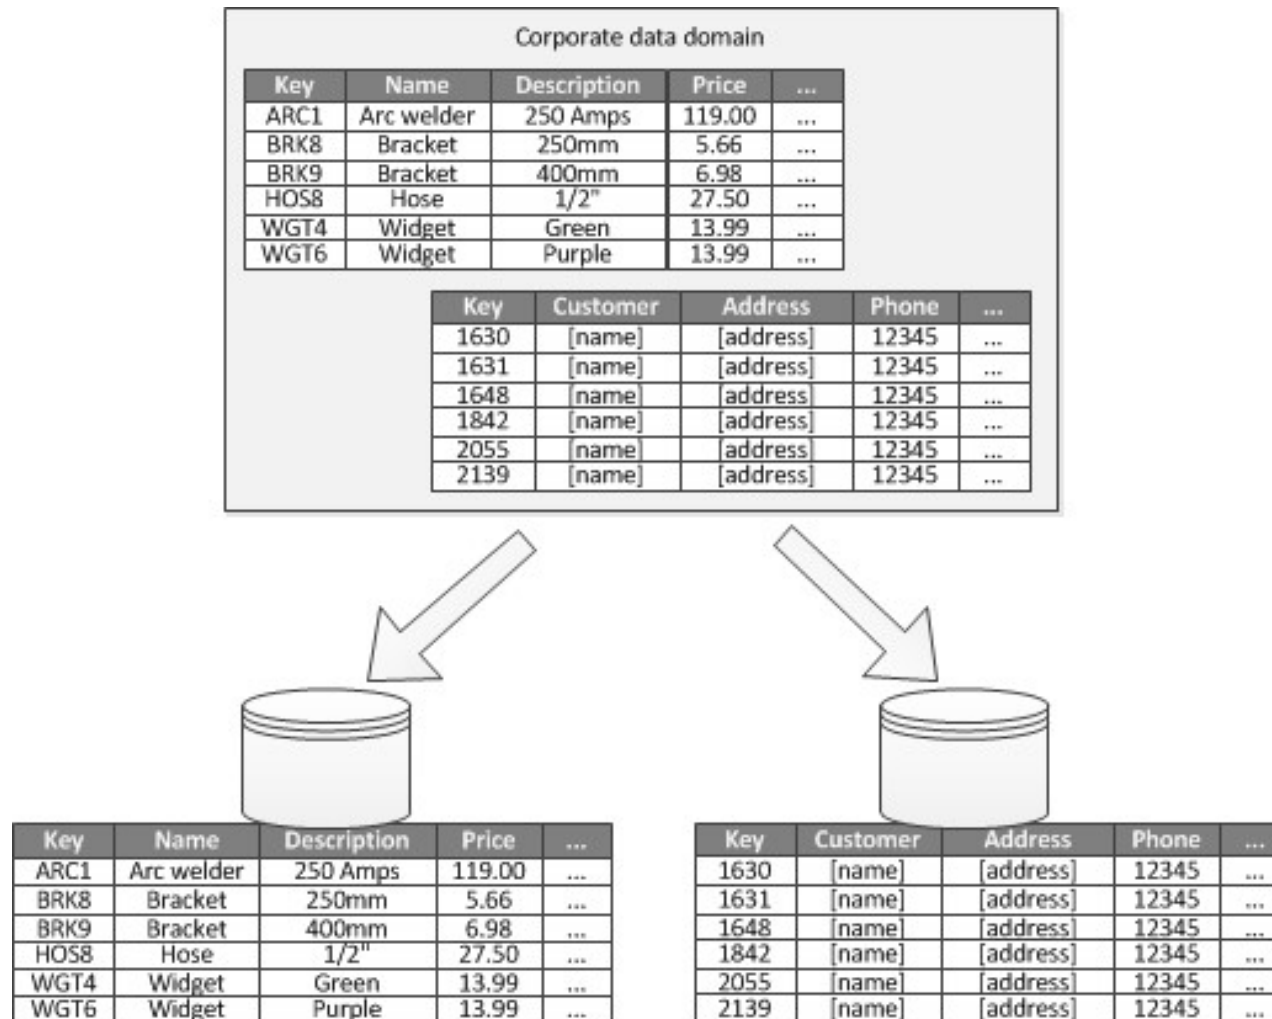
# Horizontal partitioning (Sharding)



# Vertical partitioning



# Functional partitioning






# NoSQL



NoSQL databases are non tabular, and store data differently than relational tables



# Document model



- These NoSQL databases replace the familiar rows and columns structure with a document storage model.
- Document-Oriented NoSQL DB stores and retrieves data as a key value pair

# Graph model



- It is database that uses graph structures for semantic queries with nodes and edges
- The entity is stored as a node with the relationship as edges.
- Every node and edge has a unique identifier.

# Key-value model



- In this NoSQL database model, a key is required to retrieve and update data.

# Column-based

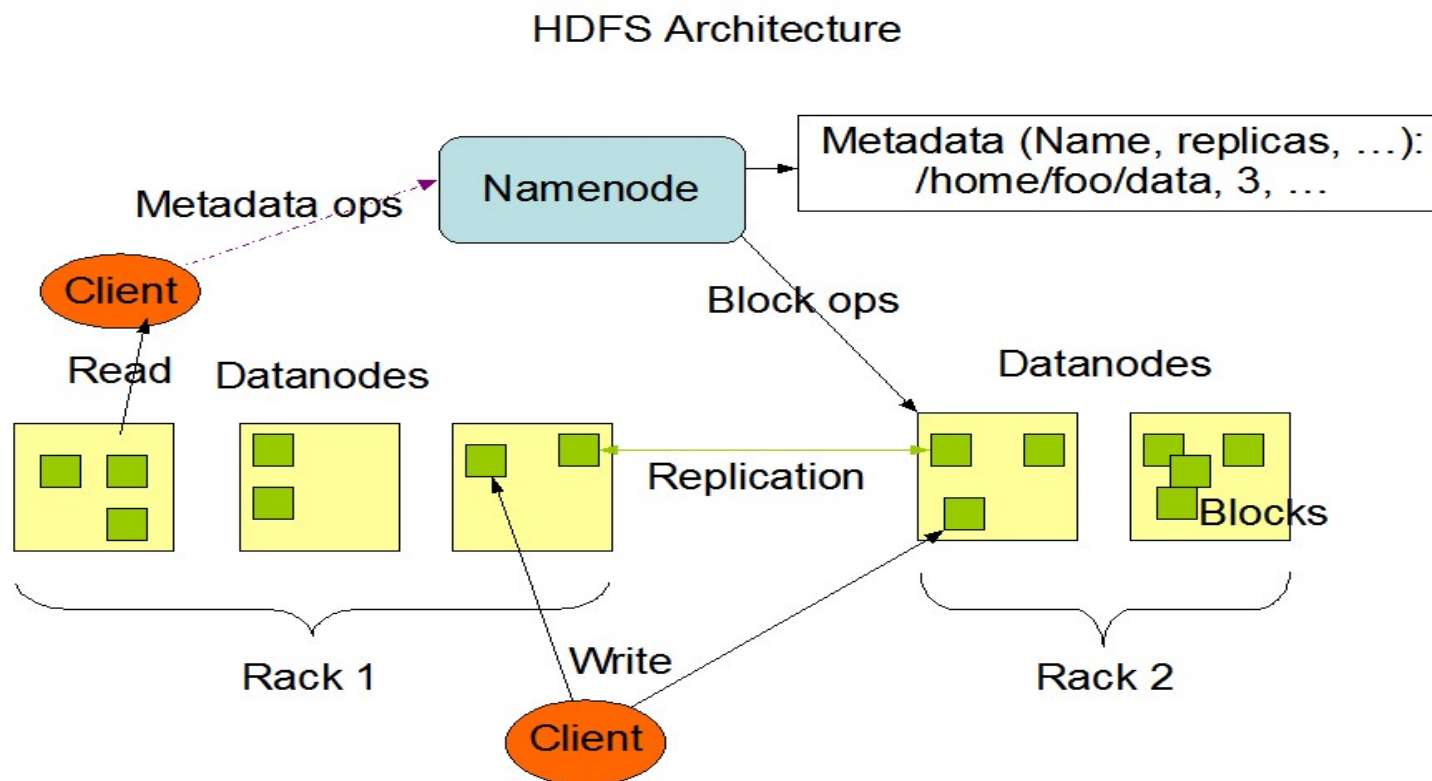


- Column-oriented databases work on columns and are based on BigTable paper by Google.
- Every column is treated separately. Values of single column databases are stored contiguously.

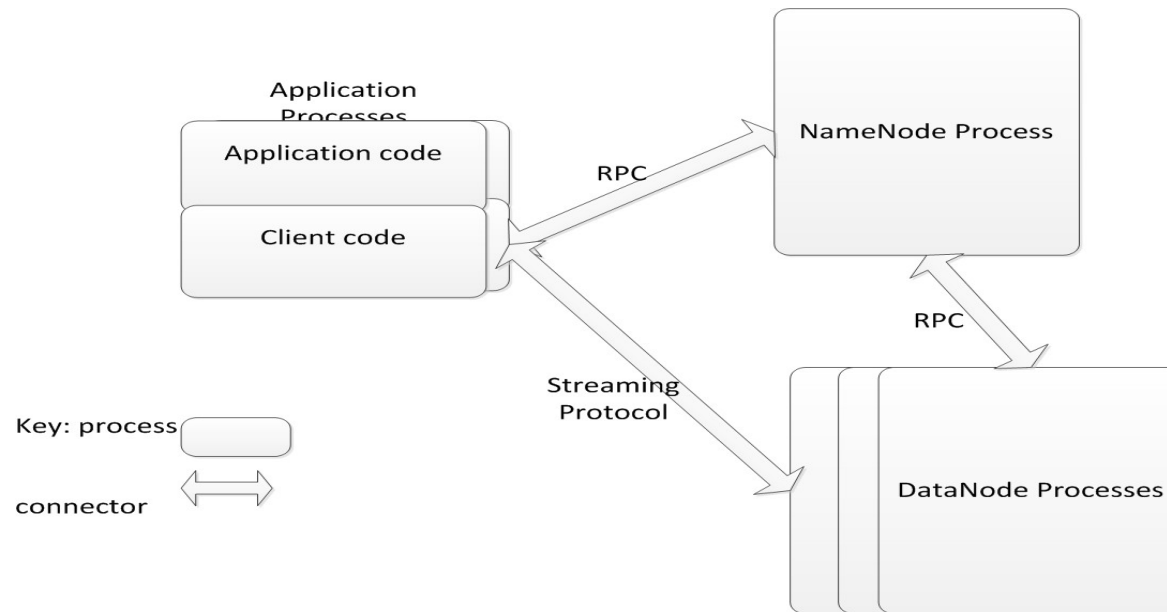
# HDFS



- The Hadoop Distributed File System (HDFS) is a distributed file system designed to run on commodity hardware.



# HDFS Components





# HDFS Write

---

- Application writes as to any file system
- Client buffers until it gets 64K block
- Client informs NameNode it wishes to write a new block
- NameNode returns list of three DataNodes to hold block
- Client sends block to first DataNode and informs DataNode of other two replicas.
- First DataNode writes block and sends it to second DataNode. Second DataNode writes block and sends it to last DataNode.
- Each DataNode reports to client when it has completed its write
- Client commits write to NameNode when it has heard from all three DataNodes.



# HDFS Write – Failure Cases



- Client fails
  - Application detects and retries
  - Write is not complete until committed by Client
- NameNode fails
  - Backup NameNode takes over
  - Log file maintained to avoid losing information
  - DataNodes maintain true list of which blocks they each have
  - Client detects and retries
- DataNode fails
  - Client (or earlier DataNode in pipeline) detects and asks NameNode for different DataNode.
  - Since each block is replicated three times, a failure in a DataNode does not lose any data.

# Goals of HDFS



- Fast recovery from hardware failures
- Access to streaming data
- Accommodation of large data sets
- Portability



# How MapReduce Works

---

MapReduce are two functions: Map and Reduce. They are sequenced one after the other.

- The **Map** function takes input from the disk as  $\langle \text{key}, \text{value} \rangle$  pairs, processes them, and produces another set of intermediate  $\langle \text{key}, \text{value} \rangle$  pairs as output.
- The **Reduce** function also takes inputs as  $\langle \text{key}, \text{value} \rangle$  pairs, and produces  $\langle \text{key}, \text{value} \rangle$  pairs as output.



# Combine and Partition

---

There are two intermediate steps between Map and Reduce.


- **Combine** is an optional process. The combiner is a reducer that runs individually on each mapper server. It reduces the data on each mapper further to a simplified form before passing it downstream.
- **Partition** is the process that translates the <key, value> pairs resulting from mappers to another set of <key, value> pairs to feed into the reducer. It decides how the data has to be presented to the reducer and also assigns it to a particular reducer.

# MapReduce Pattern

Input-Map-Reduce-Output




Input-Multiple Maps-Reduce-Output



# MapReduce Pattern



Input-Map-Combiner-Reduce-Output



# Example: Word count problem

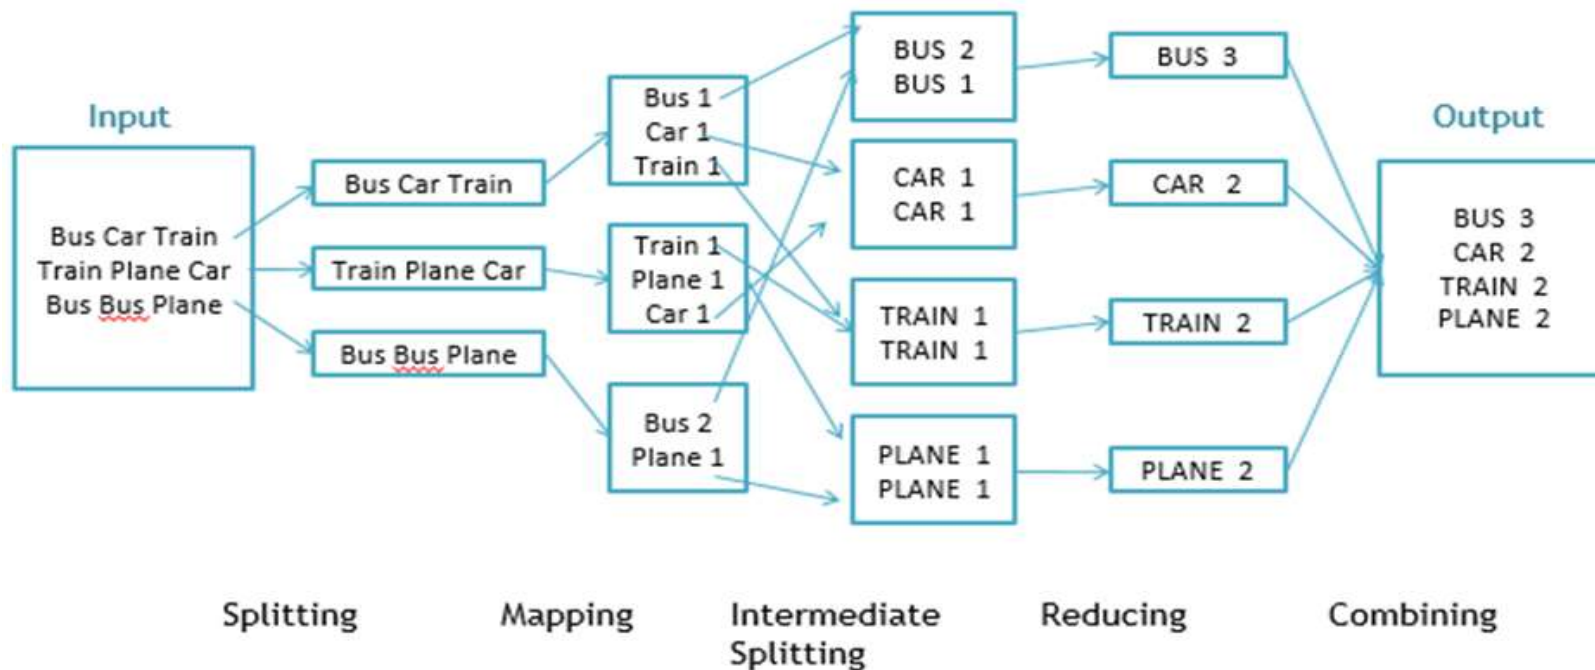


Image: dzone



**BITS Pilani**  
Pilani Campus



# Managing high velocity data streams

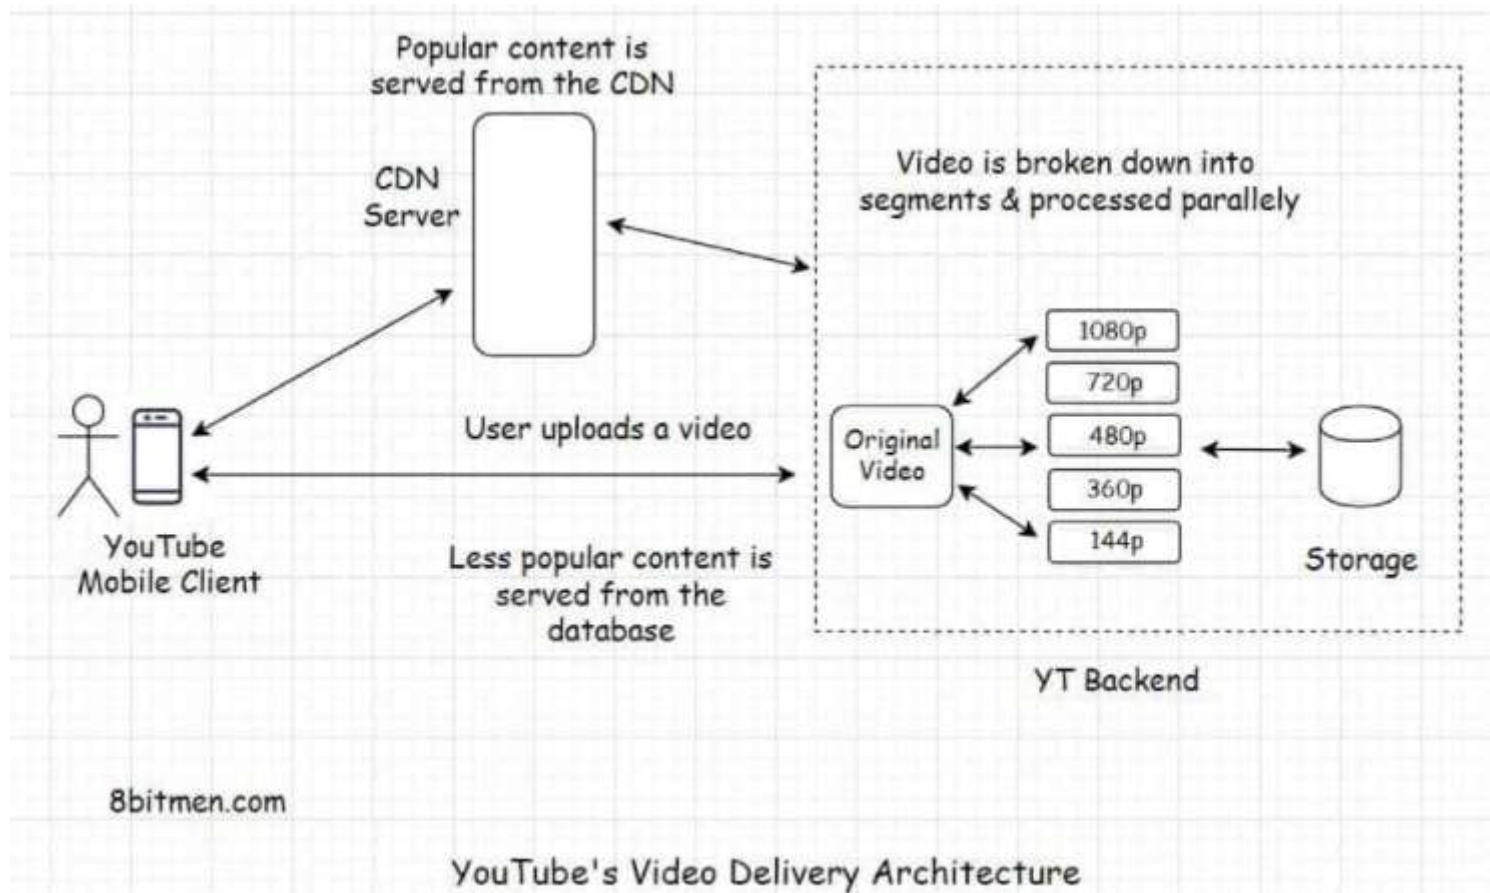


# Content Delivery Network



- While back in 2005 about 1 billion people used the internet on a daily basis, today there are 3.5 billion internet users that share 4 Exabytes (4,000,000,000 Gigabytes) of data every single day.
- Basically a CDN is nothing more than a bunch of globally distributed computers that are directly connected and move data from one end to another.
- A good example of this is YouTube.

# YouTube working

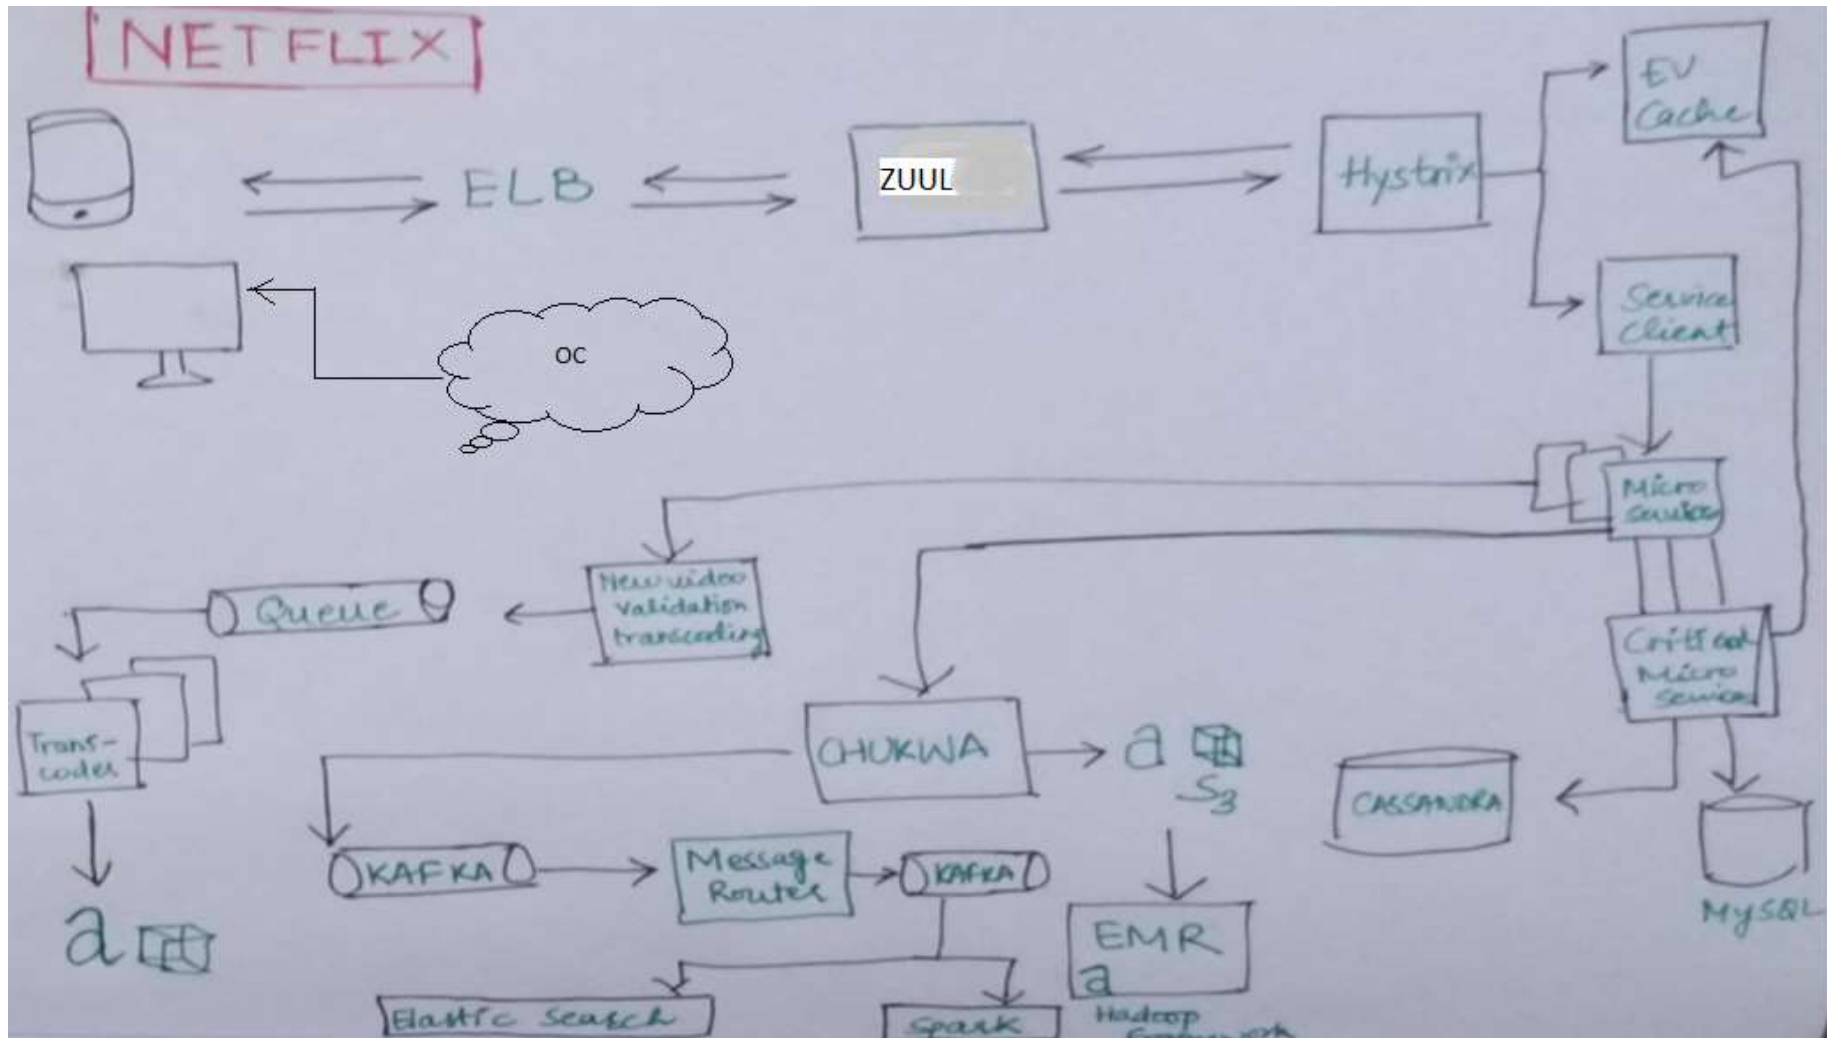


# Video streaming: Netflix



- Netflix launched in 1998. At first they rented DVDs through the US Postal Service. But Netflix saw the future was on-demand streaming video
- In 2007 Netflix introduced their streaming video-on-demand service
- It starts when you hit 'Play.'
- When Netflix hands off your video to your ISP, they must carry it through their network to your home.

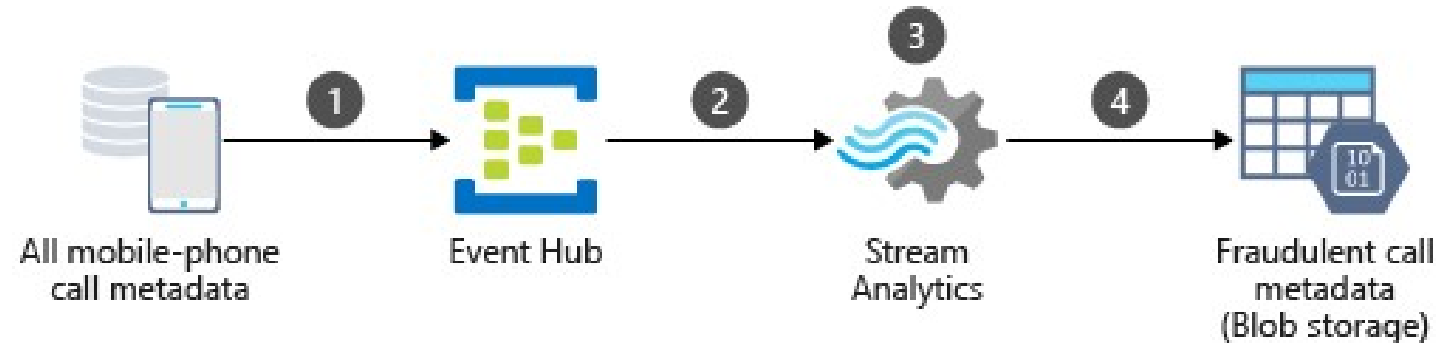
# Netflix Architecture



# Real Time Fraud Detection



## Architecture



- Mobile phone call metadata is sent from the source system to an Azure Event Hubs instance.
- A Stream Analytics job is started, which receives data via the event hub source.
- The Stream Analytics job runs a predefined query to transform the input stream and analyze it based on a fraudulent-transaction algorithm.
- The Stream Analytics job writes the transformed stream representing detected fraudulent calls to an output sink in Azure Blob storage.

# Web conferencing: Zoom



- Zoom customers with Business subscriptions can enjoy three times as many video participants in their meetings at no additional cost — and without doing a thing.
- such an increase is made possible in the way the Zoom platform is engineered
- From the very beginning, Zoom was engineered to be cloud-native and optimized for video.

# Web conferencing: Zoom



There are two important aspects of Zoom's technology stack:

- Cloud network
- Video architecture
  - Distributed architecture
  - Multimedia routing
  - Multi-bitrate encoding
  - Application layer quality of service



# What is Kafka?

---


- Apache Kafka is a **publish-subscribe based durable messaging system**.
- A messaging system sends messages between processes, applications, and servers.
- Apache Kafka is a software where topics can be defined (think of a topic as a category), applications can add, process and reprocess records.

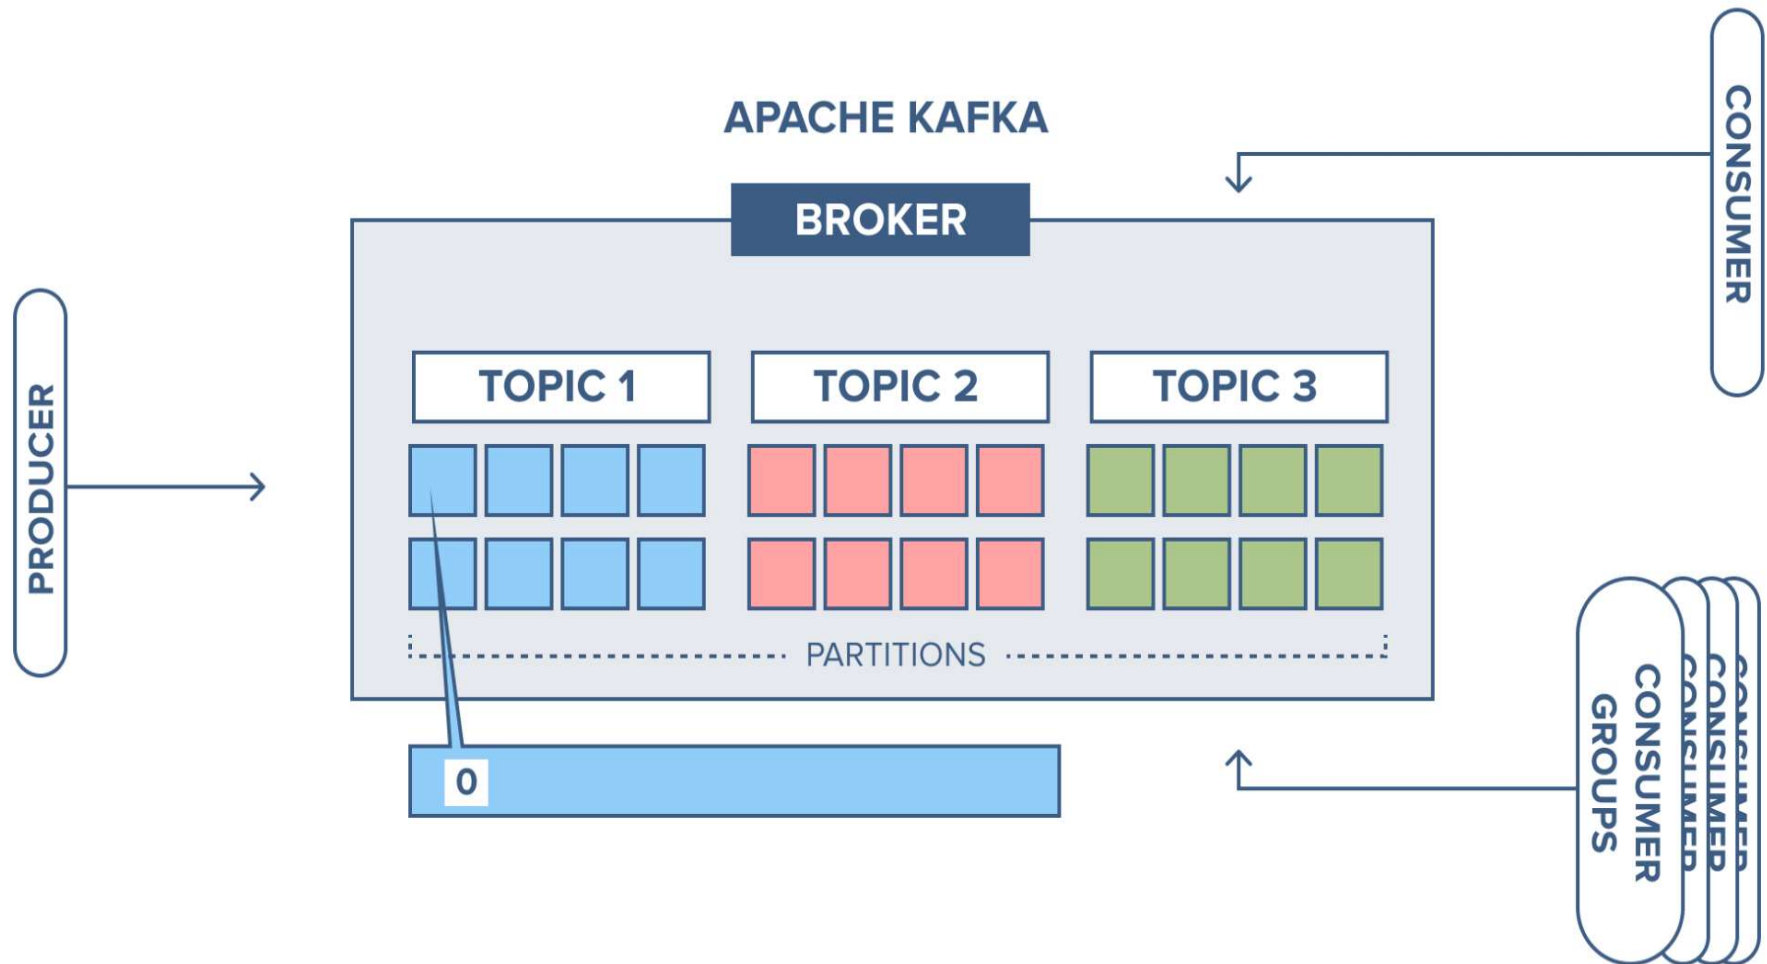


# Kafka



- Kafka Streams simplifies application development by building on the Kafka producer and consumer libraries







# Kafka Related Concepts

---

- Kafka Topics
- Partitioning
- Kafka brokers
- Replication
- Kafka Producers
- Kafka Consumers
- Kafka Connect
- Kafka Streams

# Kafka



There are close links between Kafka Streams and Kafka in the context of parallelism:

- Each **stream partition** is a totally ordered sequence of data records and maps to a Kafka **topic partition**.
- A **data record** in the stream maps to a Kafka **message** from that topic.
- The **keys** of data records determine the partitioning of data in both Kafka and Kafka Streams, i.e., how data is routed to specific partitions within topics.

# What is Edge Computing?



- Edge computing is a distributed information technology (IT) architecture in which client data is processed at the periphery of the network, as close to the originating source as possible.
- It helps to provide server resources, data analysis, and artificial intelligence to data collection sources and cyber-physical sources like smart sensors and actuators

# Edge computing: IoT systems



- Rapidly increasing numbers of IoT devices and resultant data, mean that new techniques are needed to meet customer requirements and ensure effective management need to be explored

# Key Benefits of Edge for the IoT

---



- Low latency
- Longer battery life for IoT devices
- Access to data analytics and AI
- Resilience
- Scalability
- More efficient data management

# Example: IoT Image and Audio Processing

---



- IoT edge introduces new ways of analysing data without having to backhaul the entire image or audio stream
- An edge cloudlet can be used to process the image, video or audio data to determine key information, such as licence plate numbers or the number of people in an area.



# Self Study



- <https://developer.cisco.com/docs/webex-meetings/#!architecture/overview>
- Kafka

# References



- <https://docs.microsoft.com/en-us/azure/architecture/best-practices/data-partitioning>
- <https://www.mongodb.com/nosql-explained>
- [https://hadoop.apache.org/docs/r1.2.1/hdfs\\_design.html](https://hadoop.apache.org/docs/r1.2.1/hdfs_design.html)
- <https://netflixtechblog.com/>
- <http://highscalability.com/youtube-architecture>
- <https://docs.microsoft.com/en-us/azure/architecture/example-scenario/data/fraud-detection>
- <https://blog.zoom.us/>
- <https://kafka.apache.org/11/documentation/streams/architecture>
- <https://www.gsma.com/iot/wp-content/uploads/2018/11/loT-Edge-Opportunities-c.pdf>
- <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7696529/>
- Reading material available on Confluent